

Autotuning OpenCL Workgroup Sizes

Tuning GPU Stencils with machine learning outperforms human experts

Chris Cummins Pavlos Petoumenos Michel Steuwer Hugh Leather
c.cummins@ed.ac.uk

**3.79x
speedup!**

Predicting OpenCL workgroup sizes
of 429 stencil programs, execution
devices, and datasets.




Hand tuning programs
is **expensive** and time
consuming

We **automate** this tuning
using collaborative
machine learning

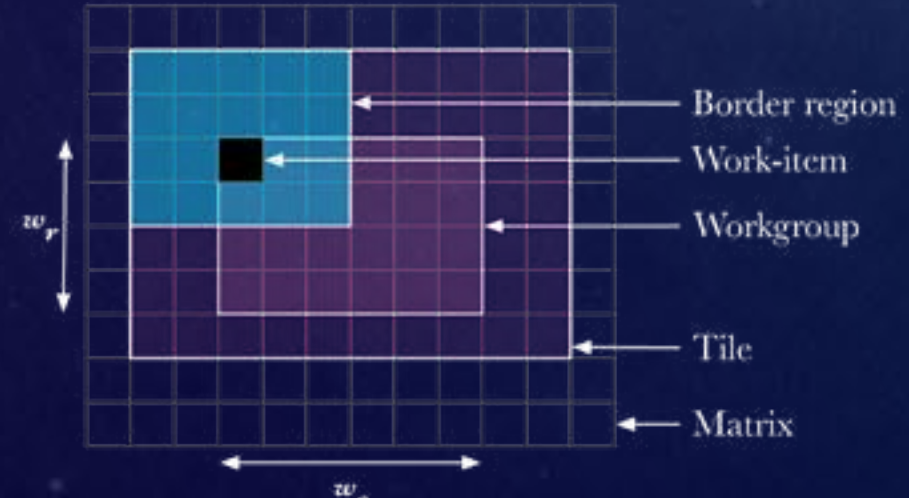
Stencil Skeletons

Stencil Skeletons are a common data parallel pattern with a range of applications from image processing to partial differential equations and cellular automata.



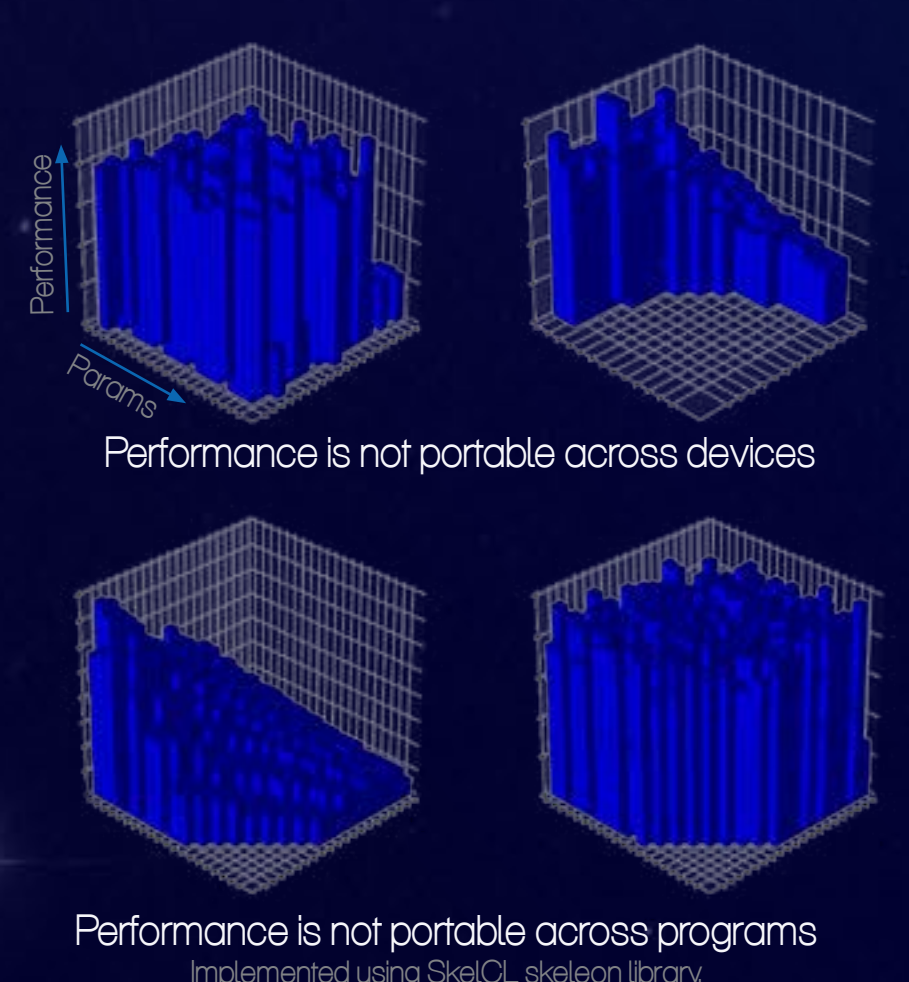
OpenCL Workgroup Size

OpenCL stencil skeletons are parameterised with a workgroup size, which controls grouping of hardware threads and local memory utilisation.



Optimization Space

Choosing the right OpenCL workgroup size for stencil kernels depends on the program, device and dataset



Introducing OmniTune ...



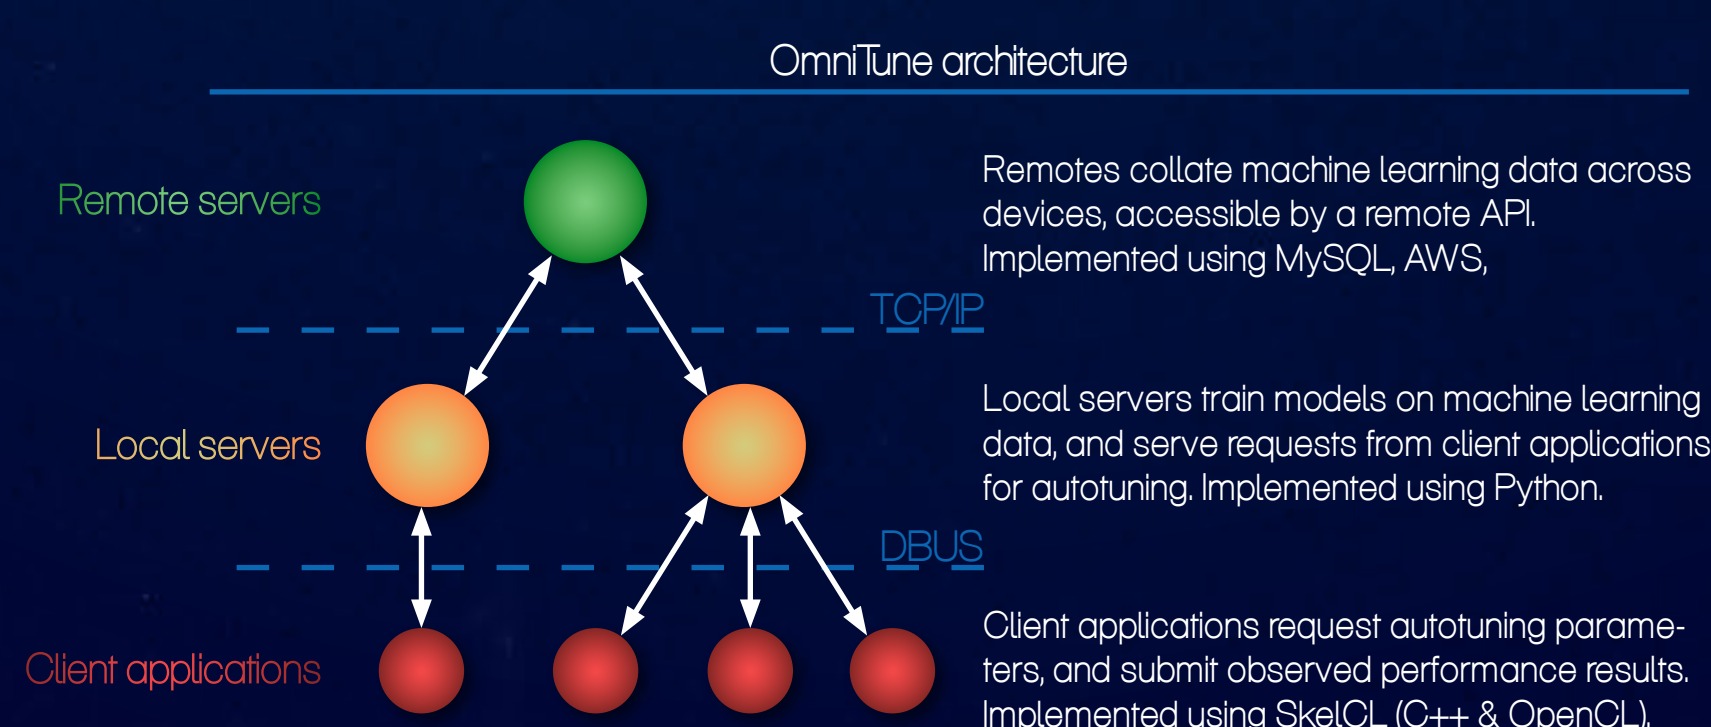
OmniTune generates synthetic benchmark programs to use for empirical testing



OmniTune collaboratively gathers performance data by testing different parameter values



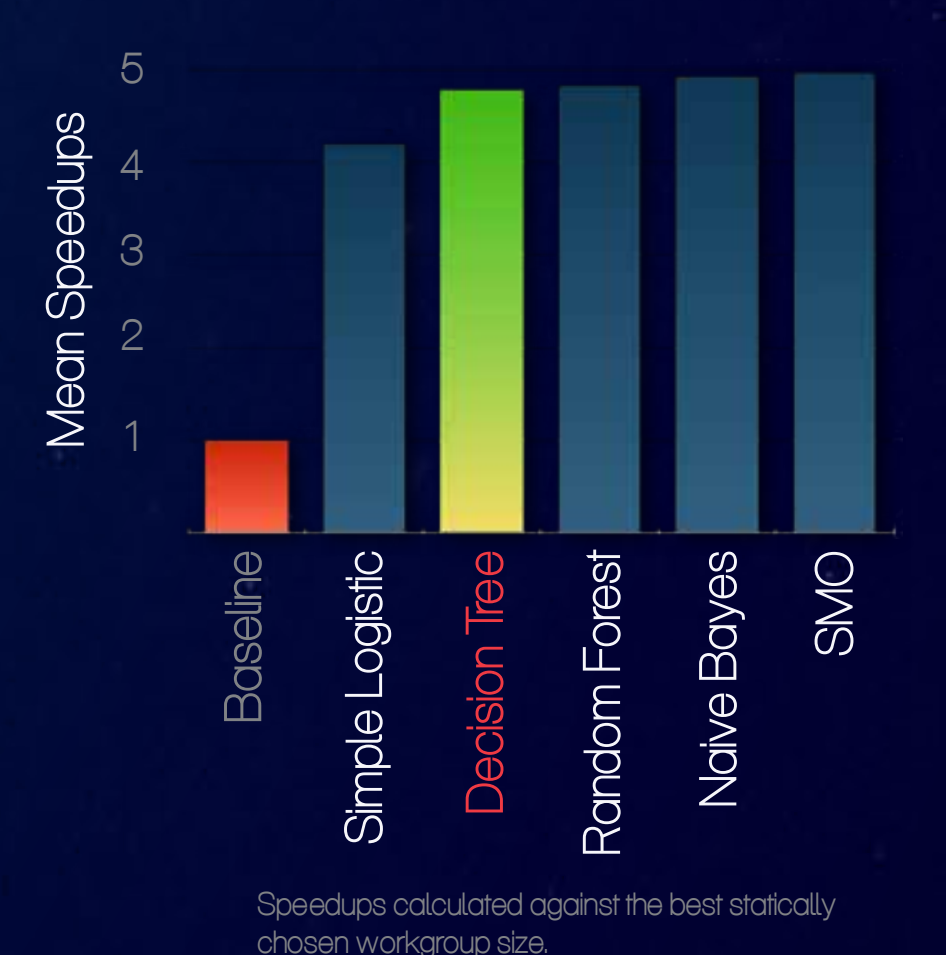
OmniTune uses machine learning to predict parameters for unseen programs at runtime



Machine Learning features



Results



Read more ...

C. Cummins, P. Petoumenos, M. Steuwer, H. Leather
"Autotuning OpenCL Workgroup Size for Stencil Patterns" ADAPT 2016.

C. Cummins, P. Petoumenos, M. Steuwer, H. Leather
"Towards Collaborative Performance Tuning of Algorithmic Skeletons" HLPGPU 2016.

<http://chriscummins.cc>

