# Dynamic Autotuning of Algorithmic Skeletons

Chris Cummins

THE UNIVERSITY of EDINBURGH informatics

EPSRC Centre for Doctoral Training in Pervasive Parallelism

EPSRC Engineering and Physical Sciences Research Council

# What are algorithmic skeletons?

**Parallel** implementations of common **patterns** of computation.

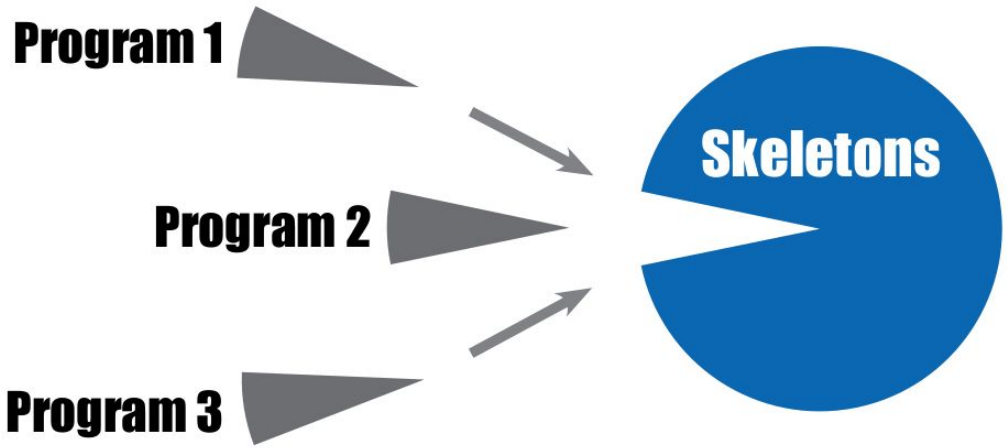

# Why do we need them?

Heterogeneous parallelism offers massive **performance**.

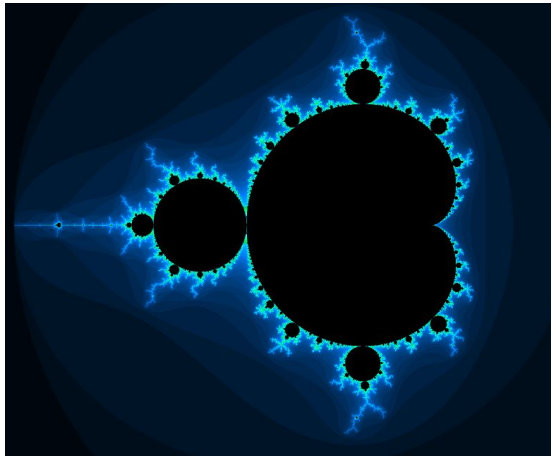

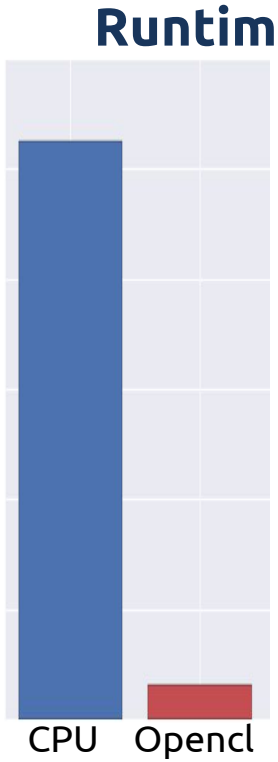


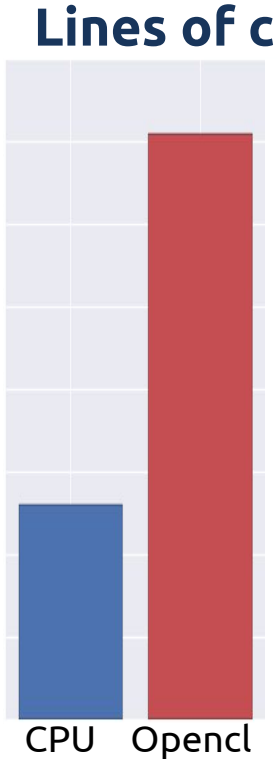

# Why do we need them?

Heterogeneous parallelism offers massive **performance**.

Algorithmic Skeletons offer **ease of use**.

For both performance **and** ease of use, we need autotuning.

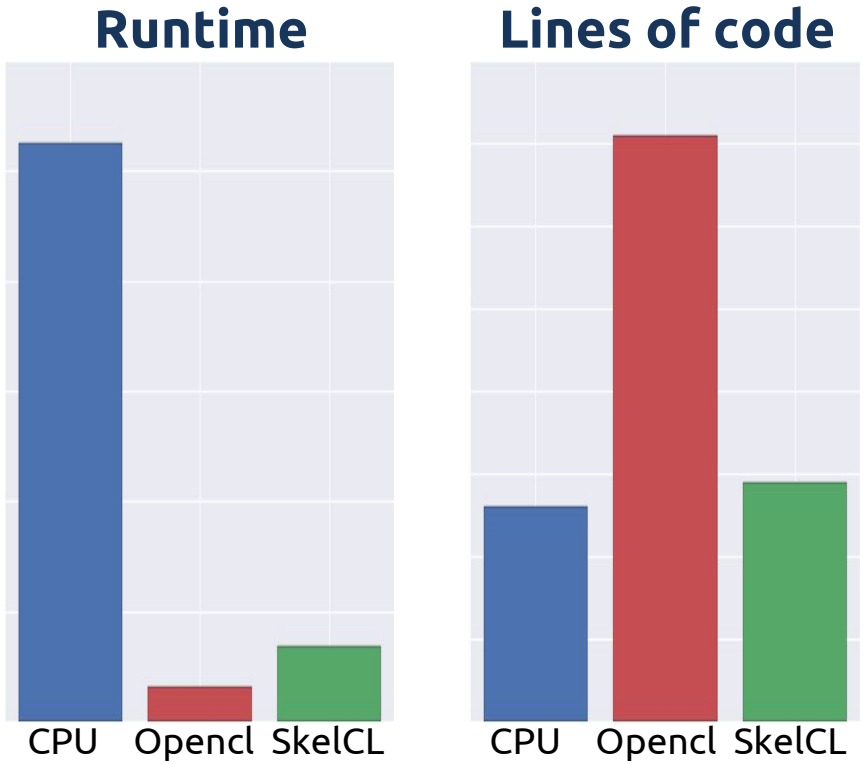

# My Project

Demonstrate **dynamic autotuning** of algorithmic skeletons.

Using the **SkelCL** data-parallel skeleton library.

Targeting **Stencils** applications on GPUs and CPUs.

# Stencil codes

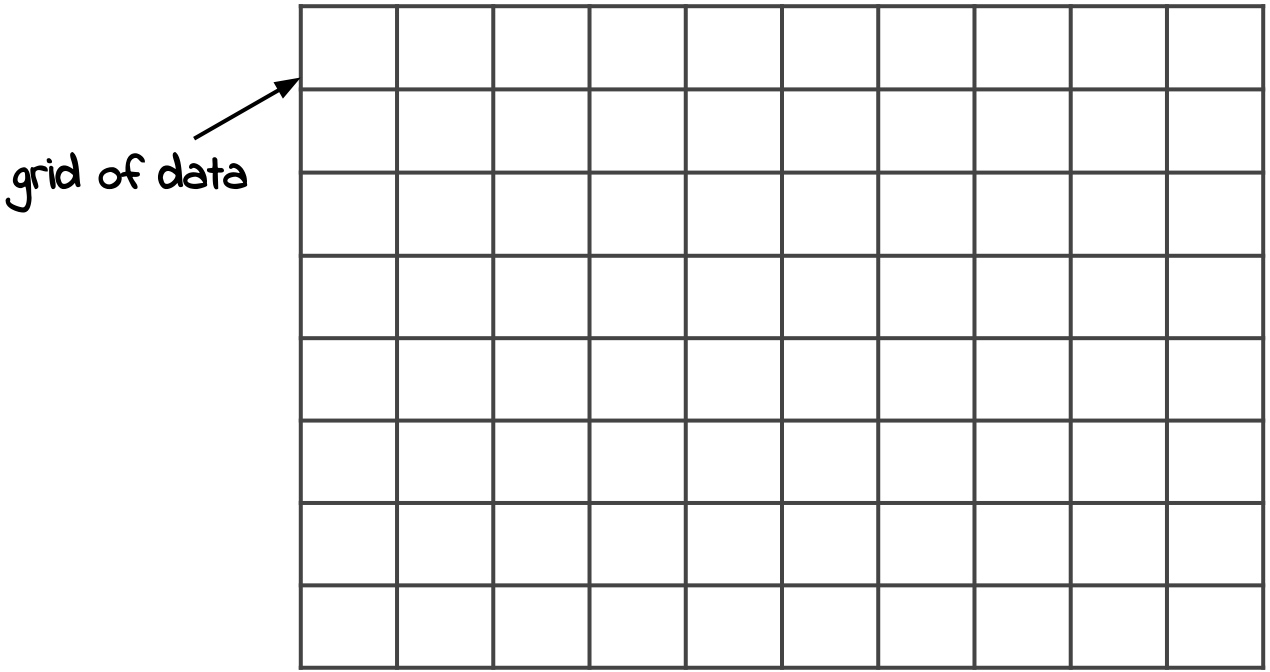

# Stencil codes

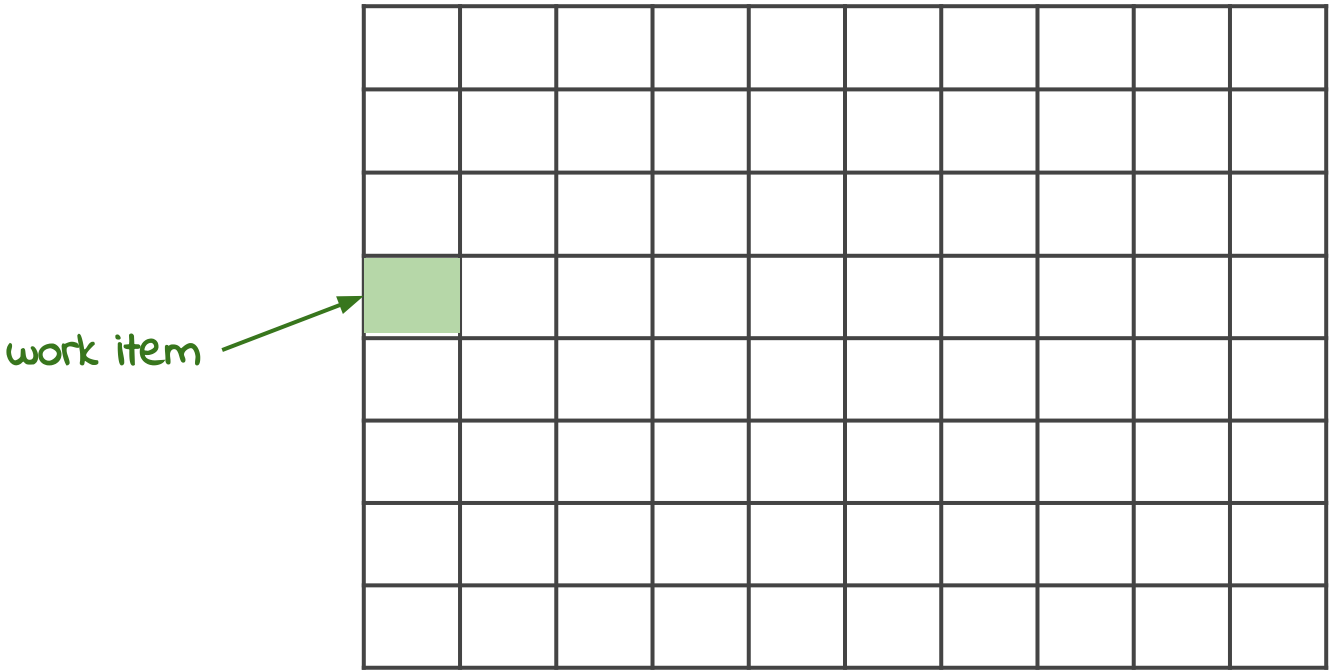

# Stencil codes

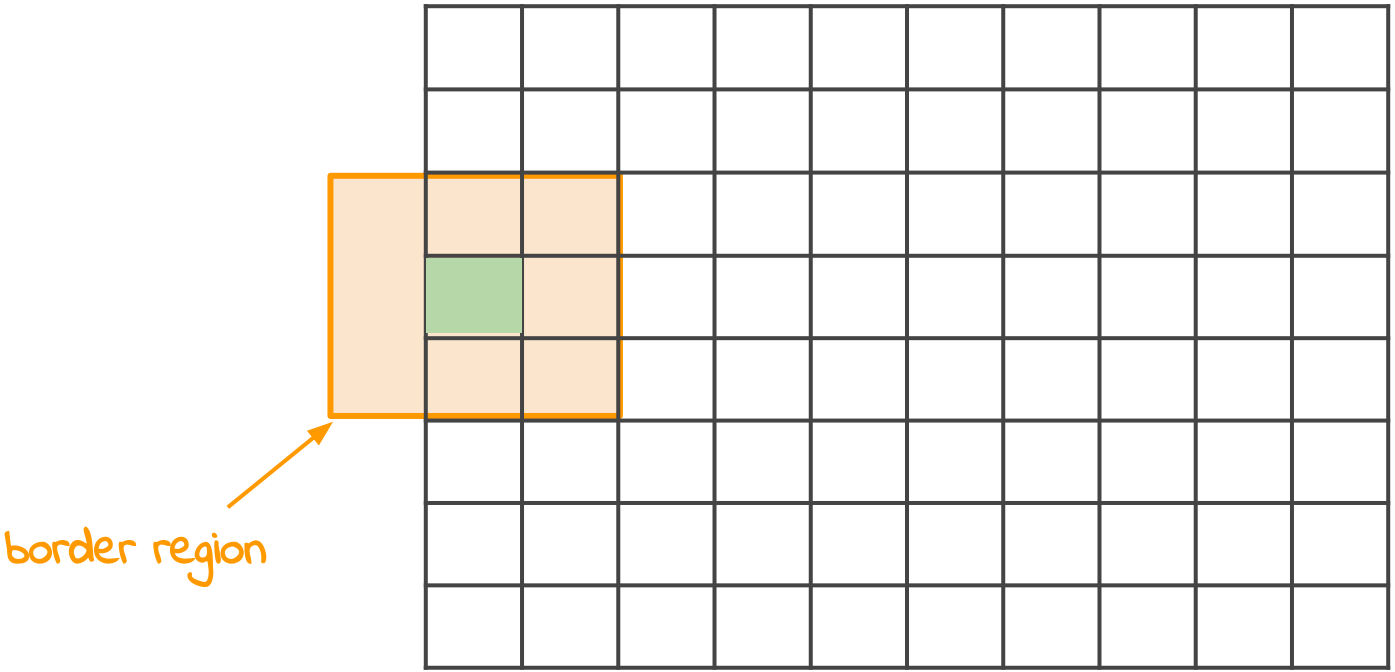

# Stencil codes

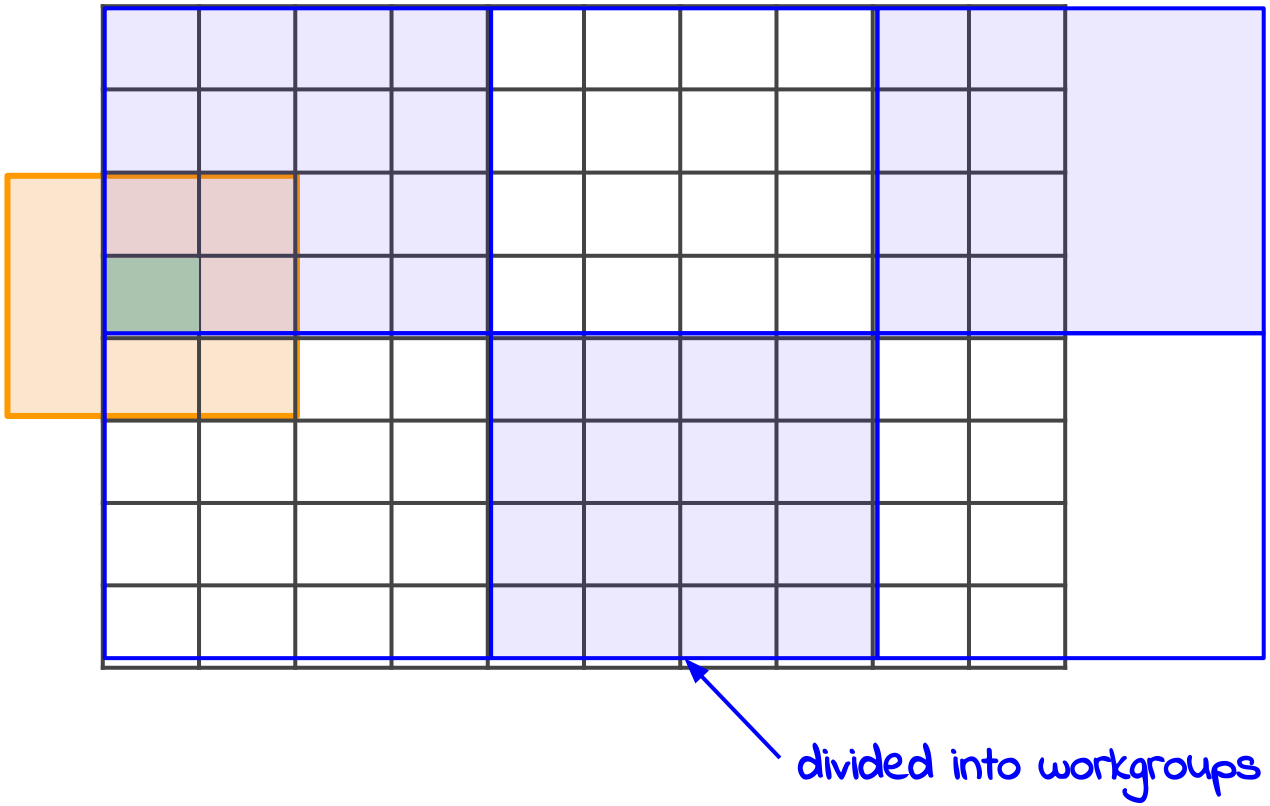

# Stencil codes

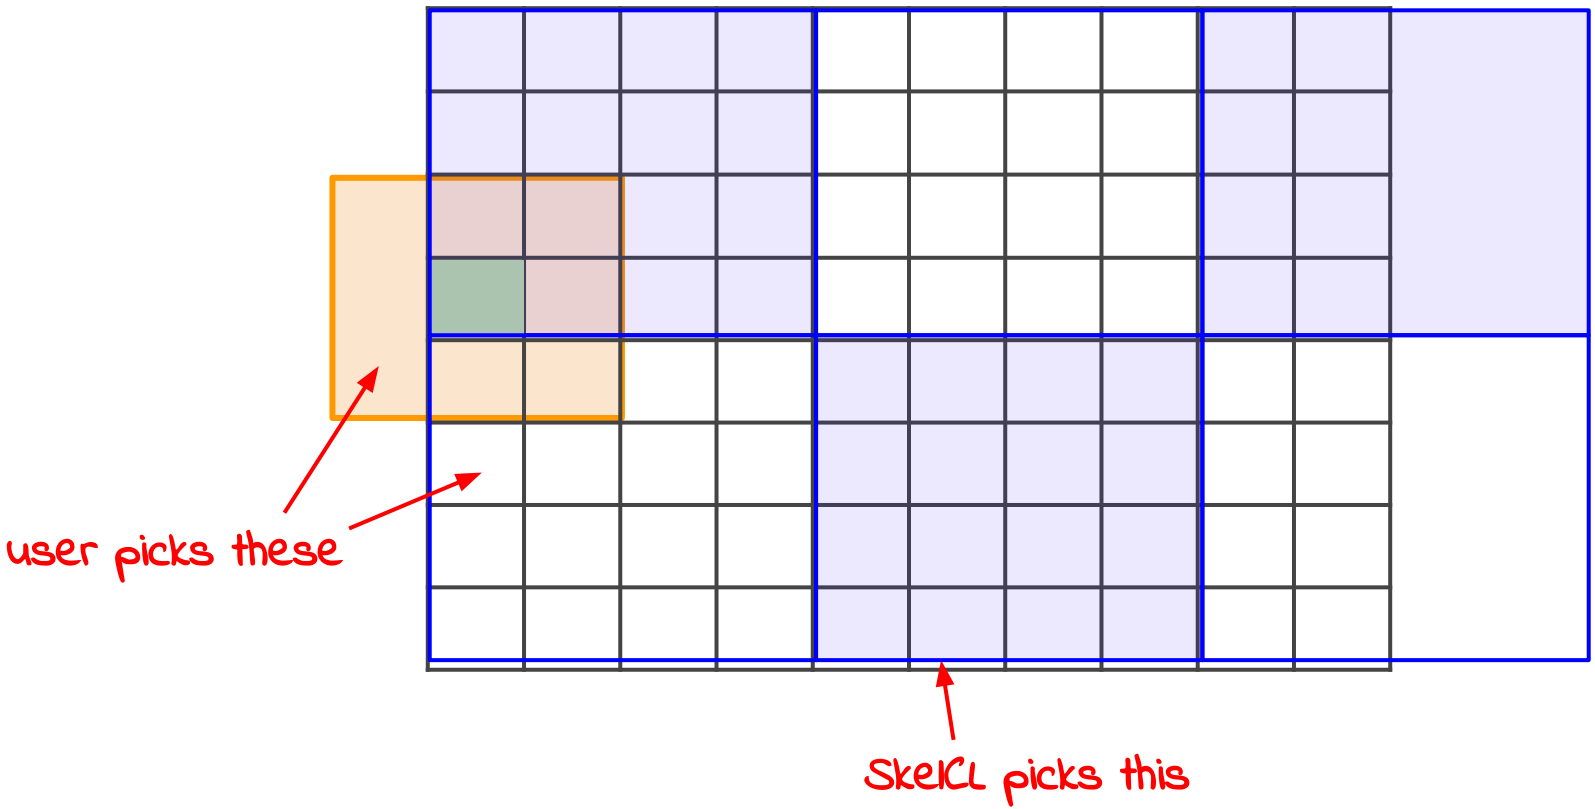

# Exploring optimisation space

Performance of workgroup size depends on:

| | |
|---|---|
| Program | Shape of border region, static instruction counts, … |
| Hardware | Local memory capacity, num processors, … |
| Dataset | Number of elements, data types, … |

**How can we test this?**

Try a bunch of **synthetic** workloads.

Measure runtime of different workgroup sizes, compare **performance**.

# Distribution of best values

No silver bullet!

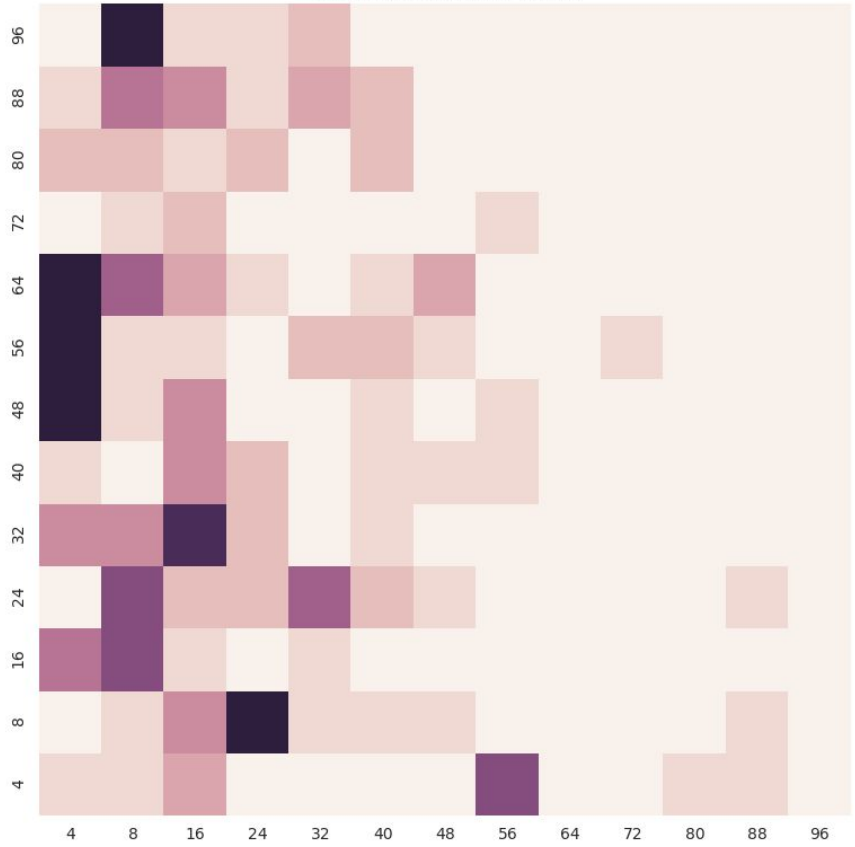

# How many values do we need?

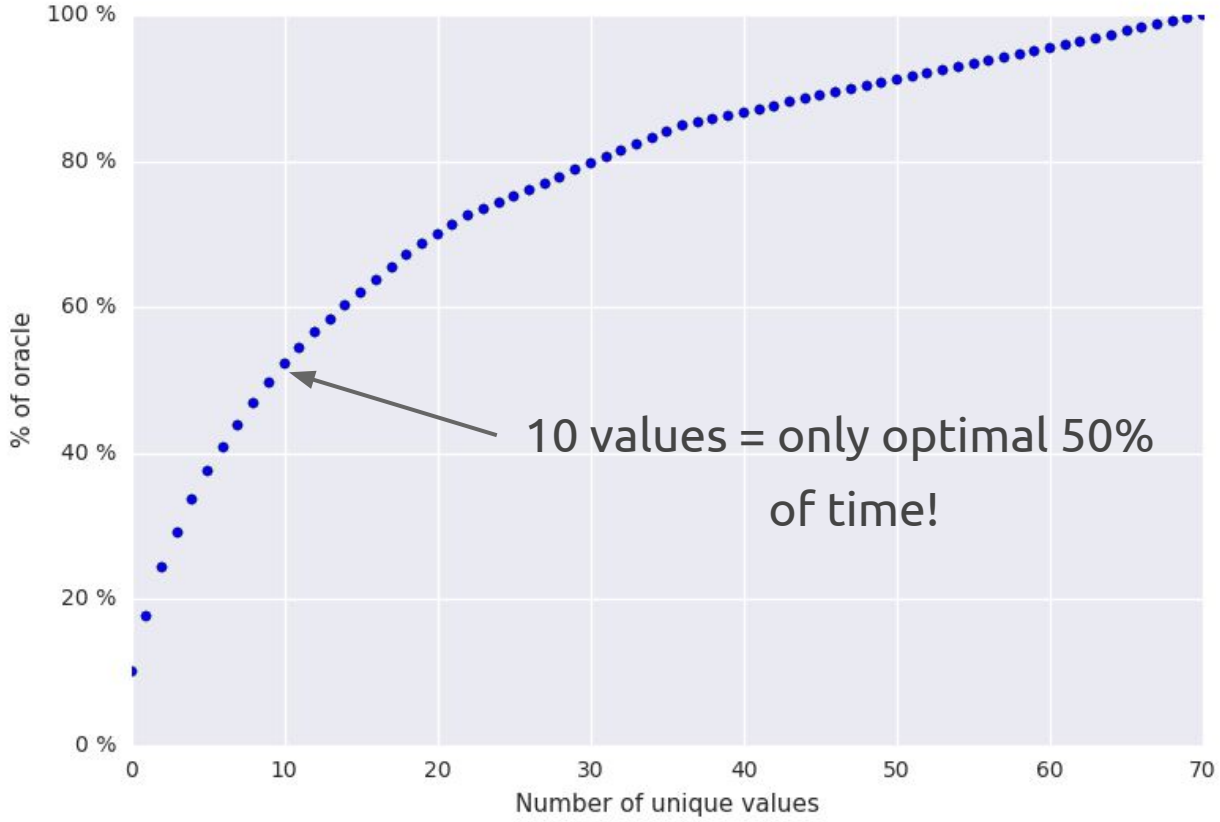

# What's the best we can do *statically*?

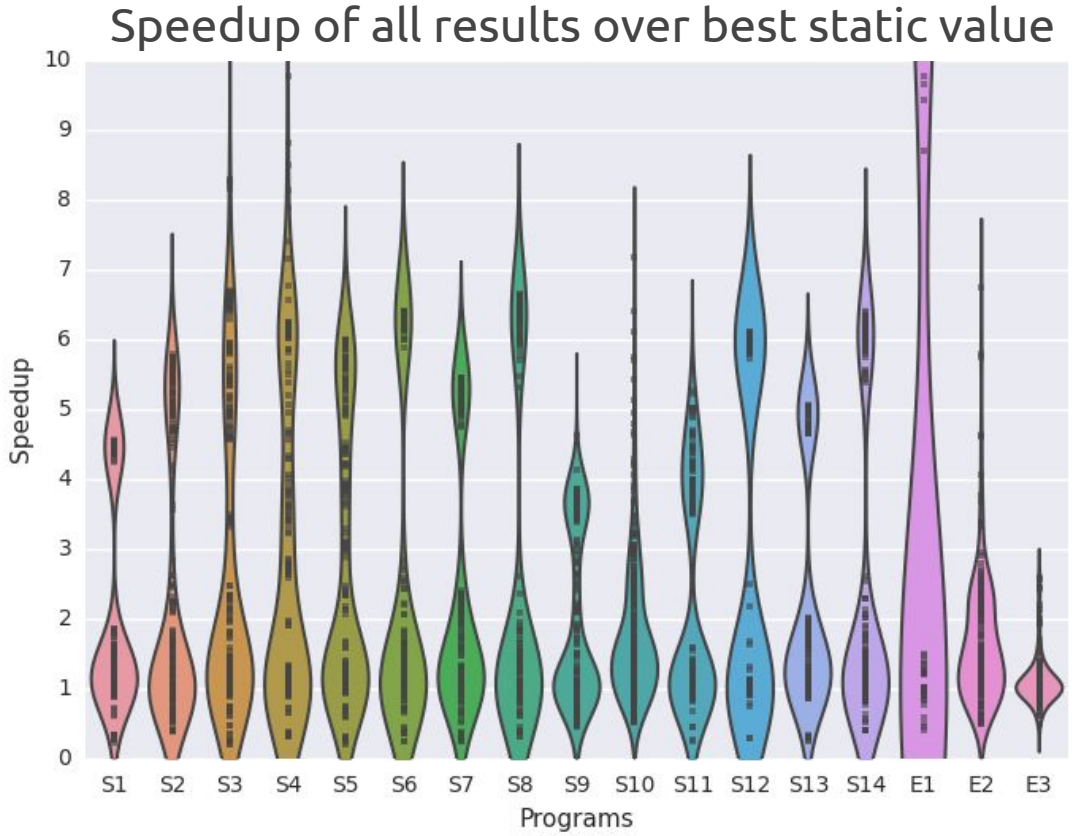

# Autotuner design

Extract **features** from hardware, program, and dataset.

Use best workgroup sizes as **training data**.

Use **machine learning** to predict:

*f(features) -> (workgroup size)*

# Autotuner design

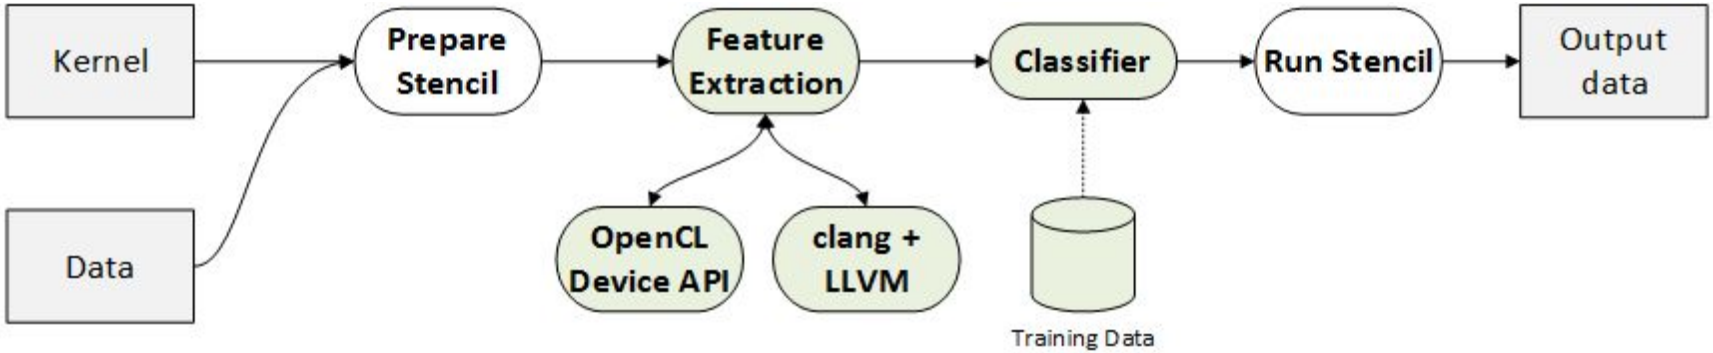

# Autotuner performance

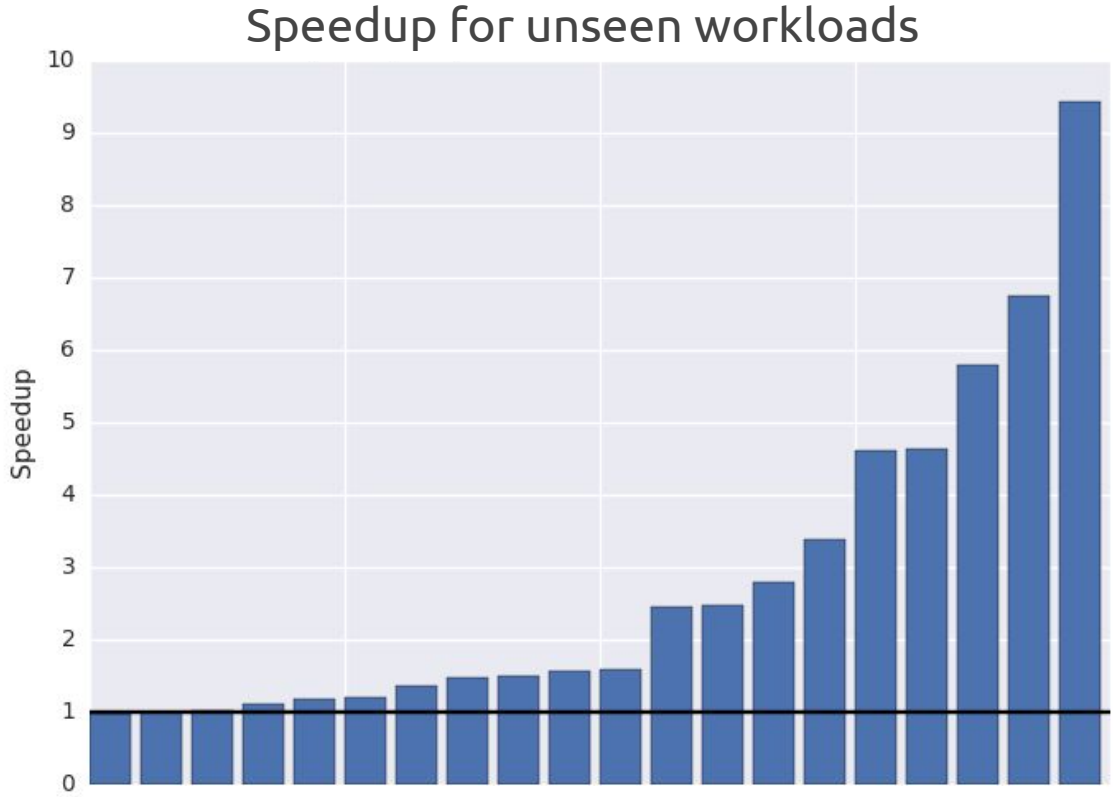


Mean: **2.82x**

# Moving Forward

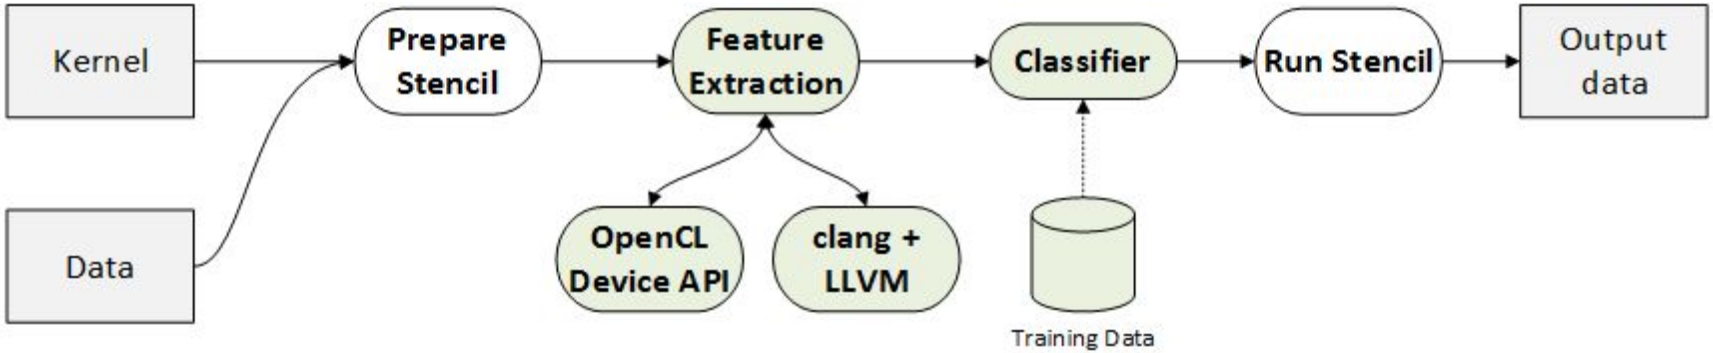

# Moving Forward

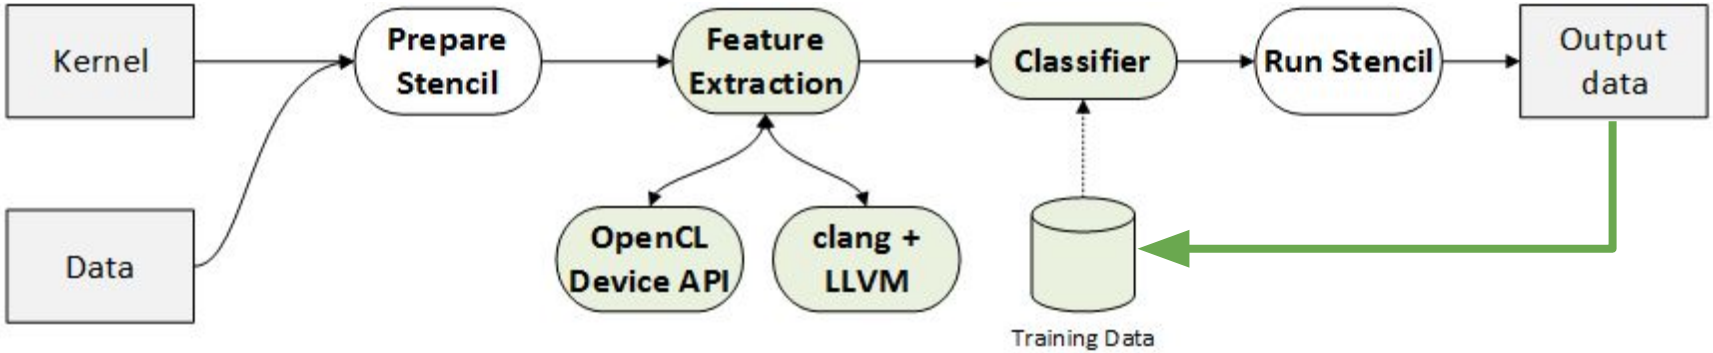

# Conclusions

High level skeletons needed for complexity of GPU programming.

Values used to parameterise these skeletons offer 10x performance margin.

Synthetic benchmarks + runtime features + machine learning
= **2.8x** performance improvement of real programs.