

Progression Review

Chris Cummins

Review

Thesis Autotuning Stencil Codes with Algorithmic Skeletons


Thesis Autotuning

Stencil Codes

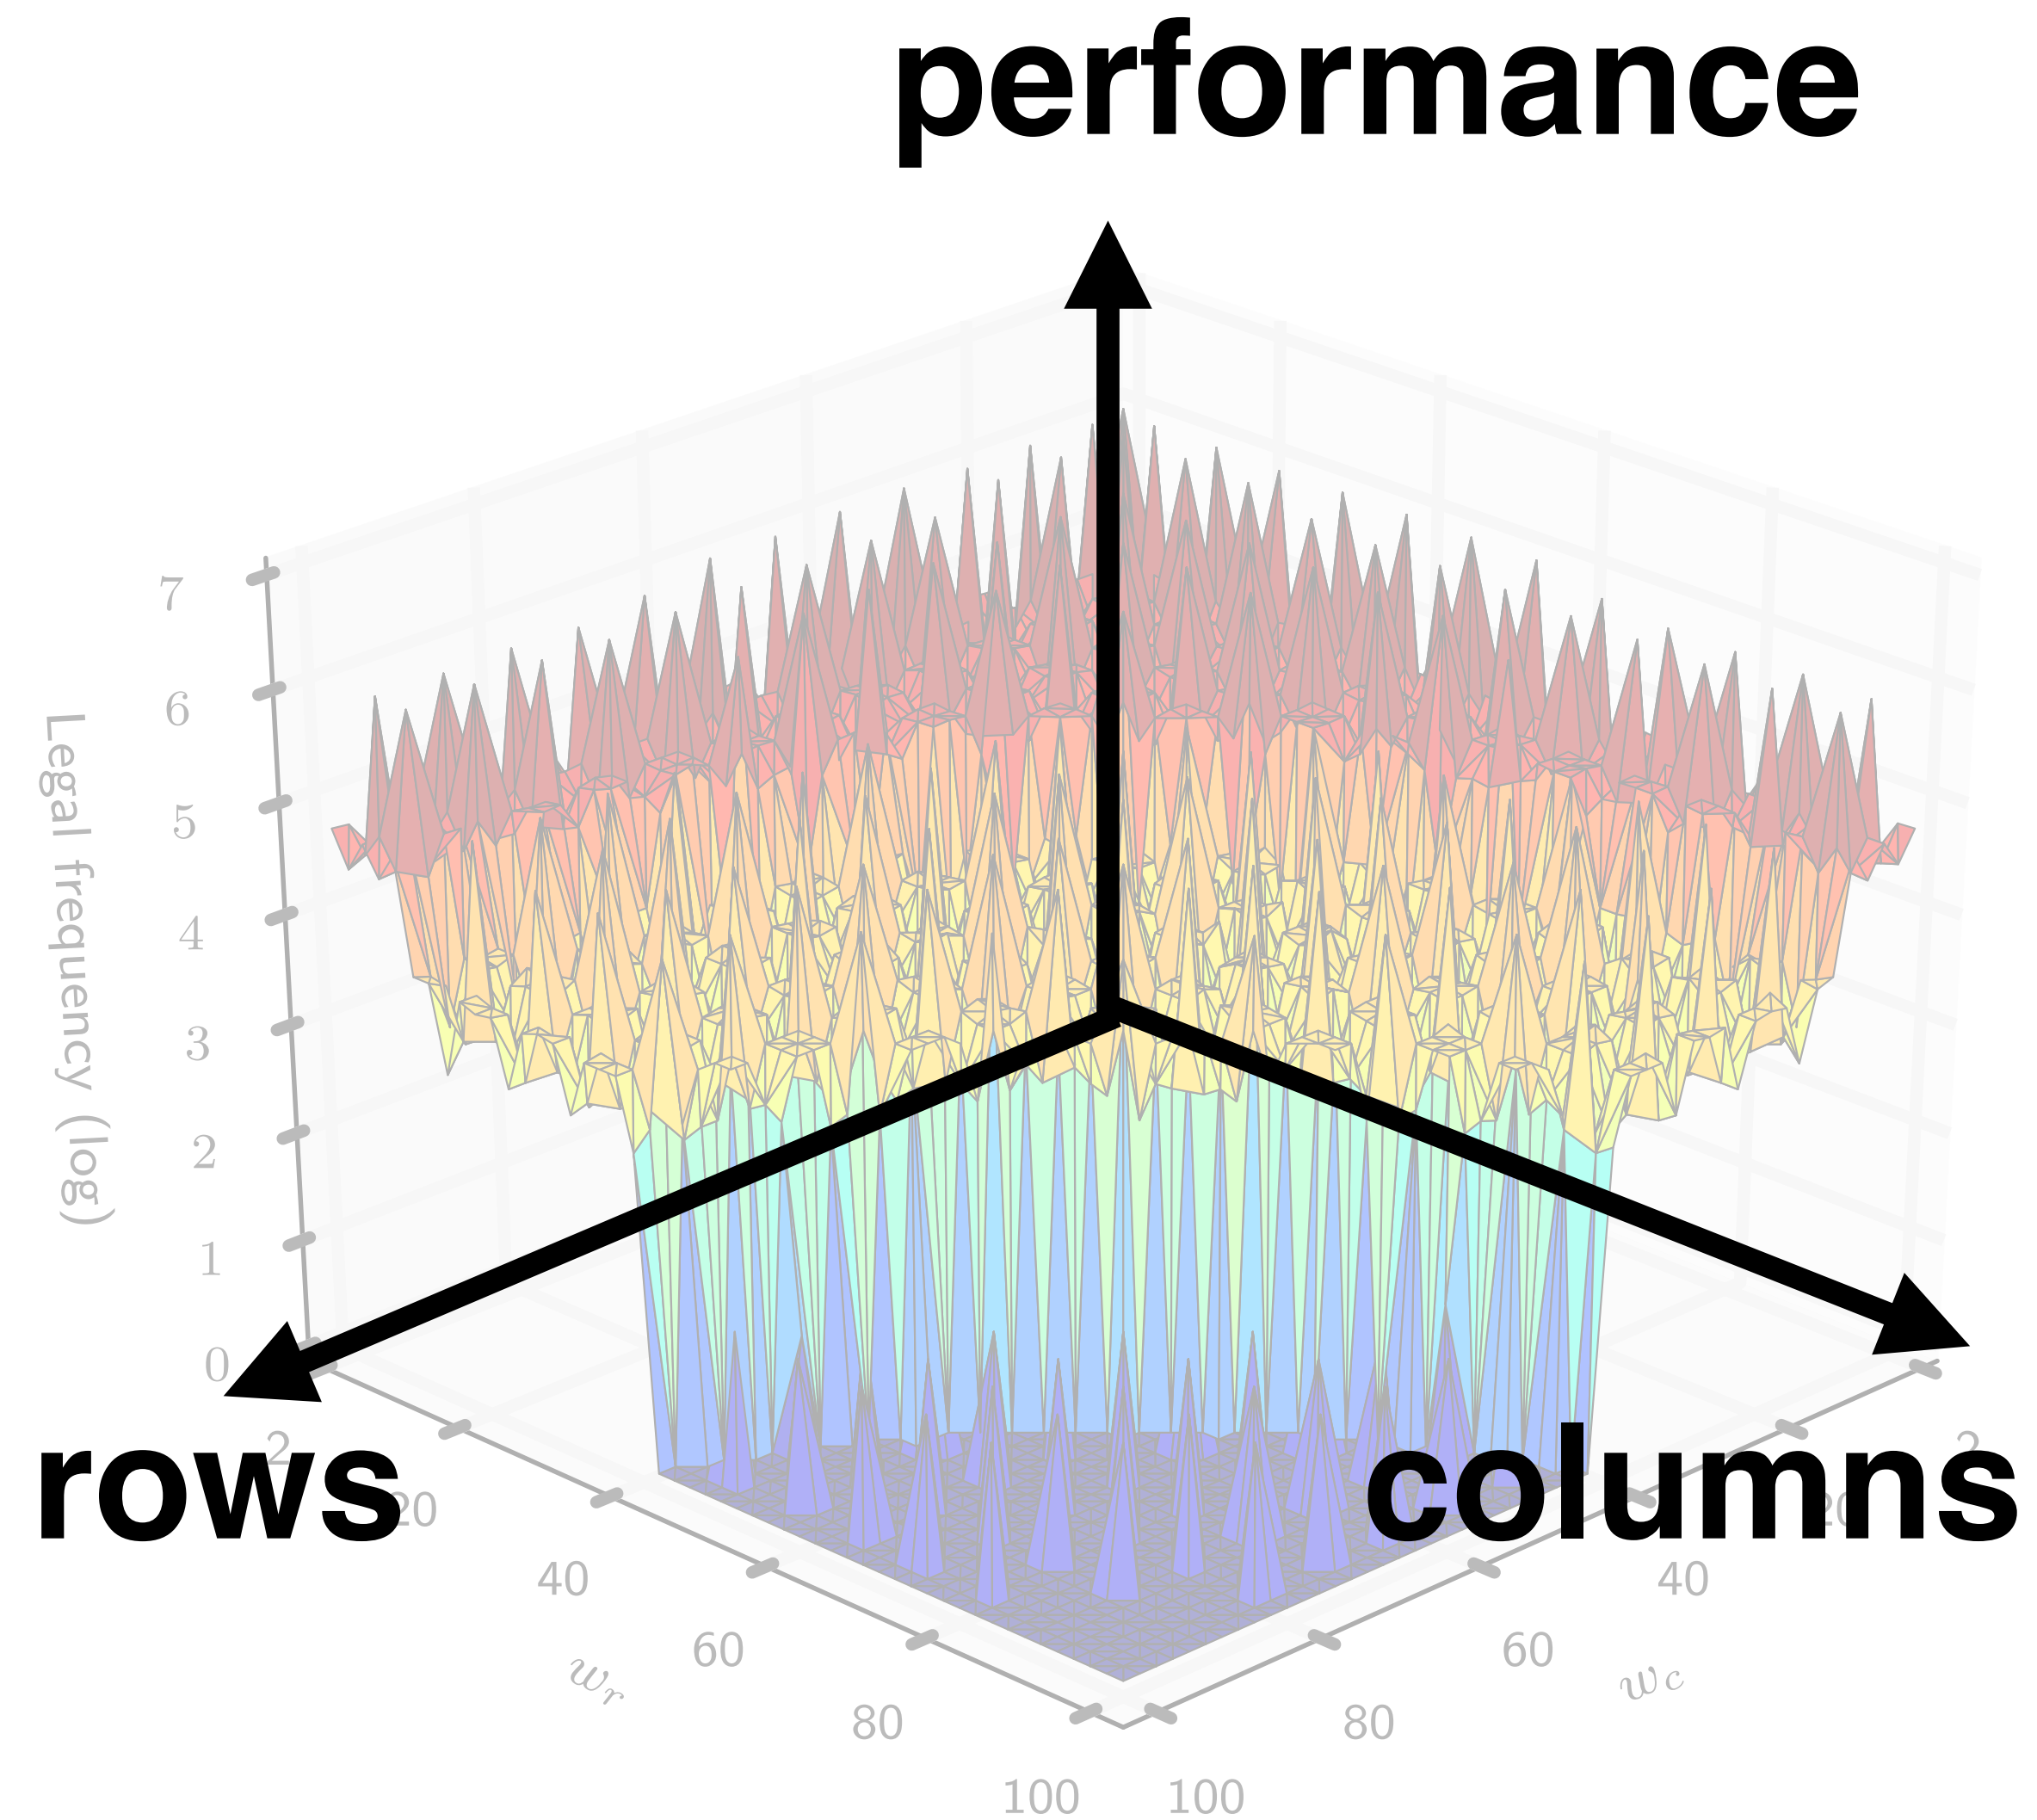
with Algorithmic

Skeletons

2D
parameter
space




2D parameter space

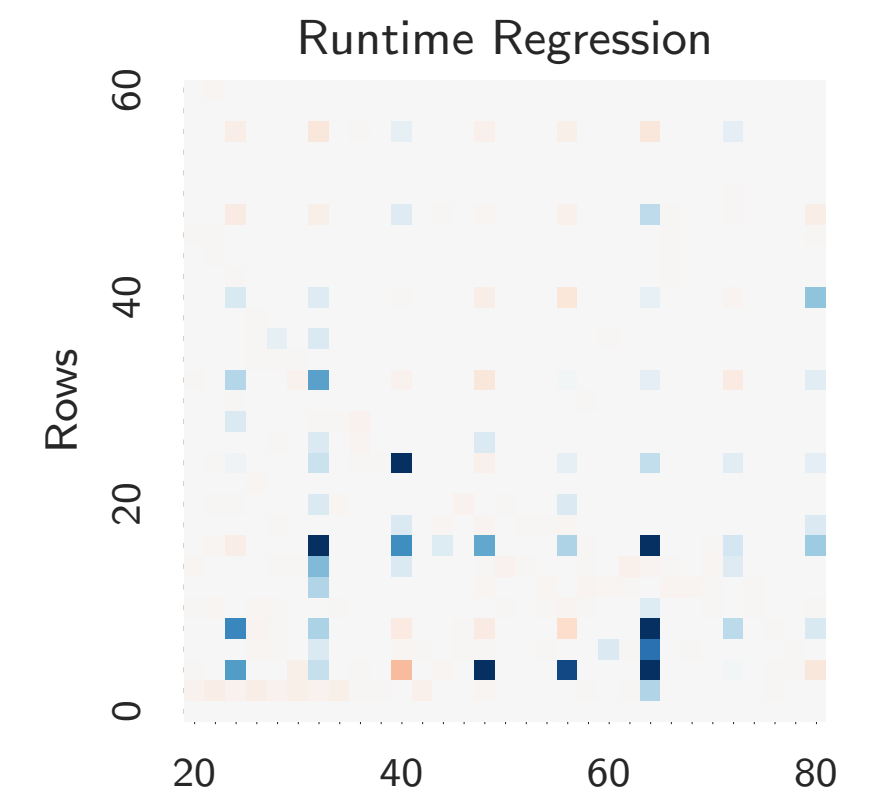
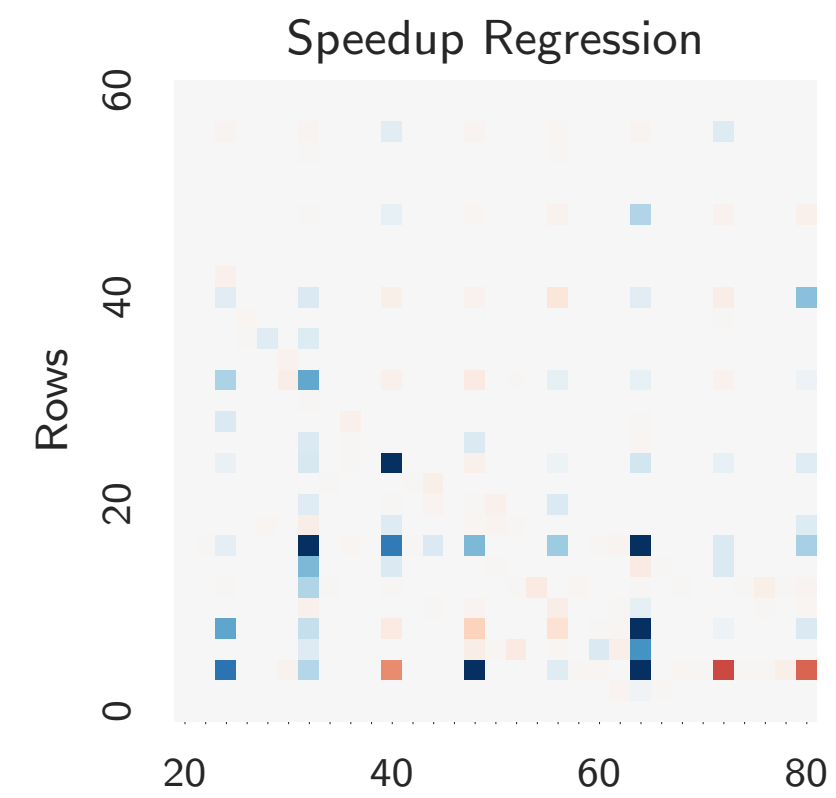

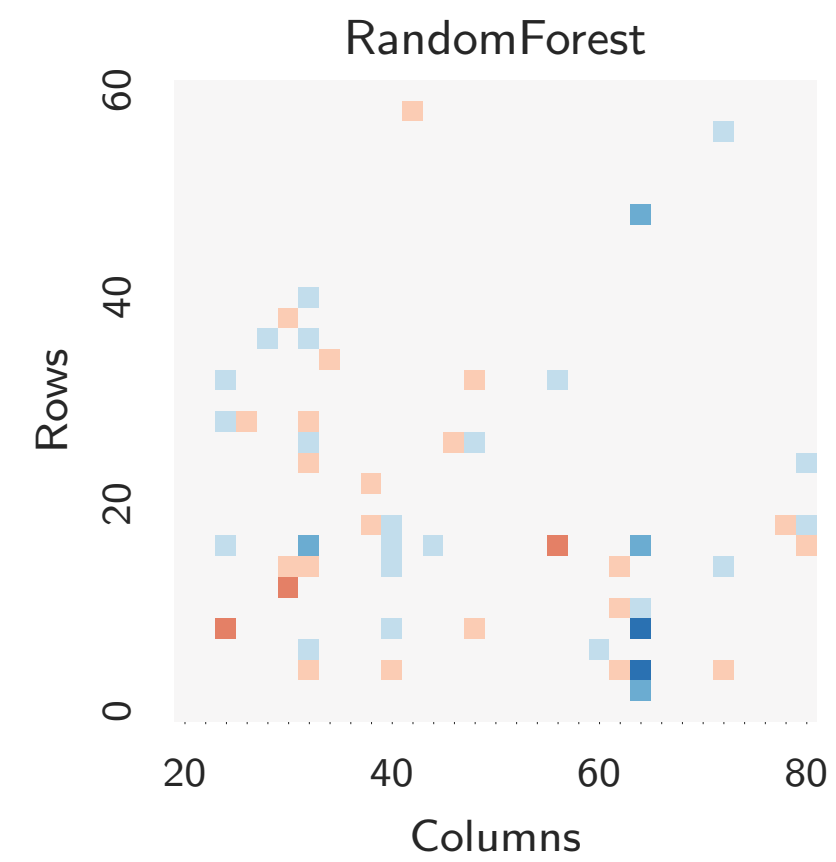

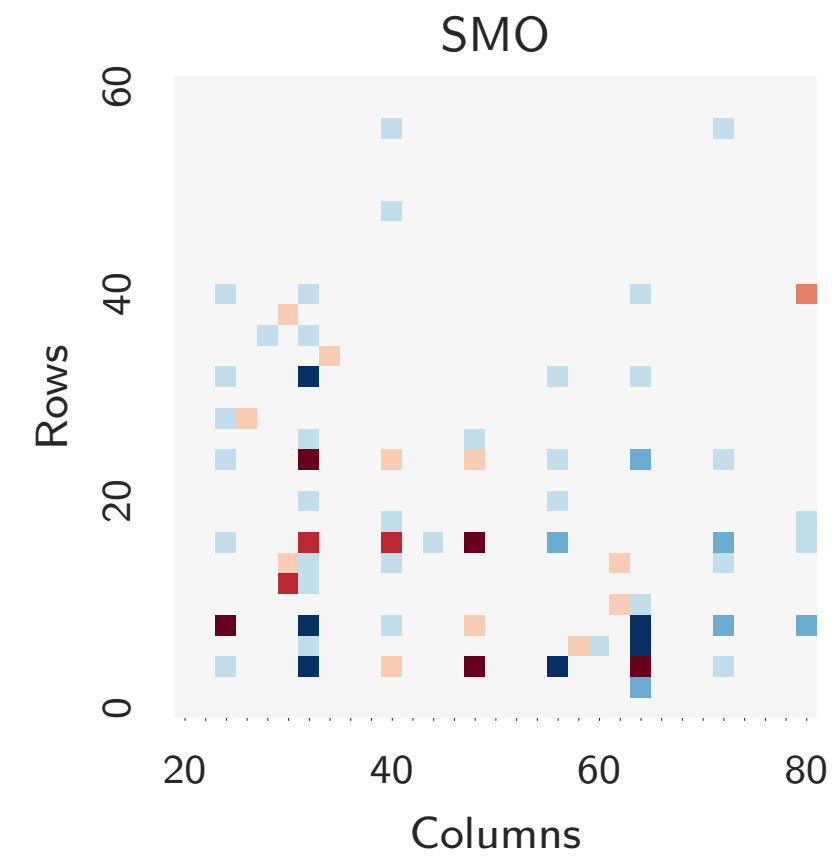



**3x Supervised
machine learning**

Predict:
**oracle
runtime
speedup**



Compare Predictions





Compare Speedups

Thesis Autotuning

Stencil Codes

with Algorithmic

Skeletons

Thesis Autotuning **Stencil Codes** **with Algorithmic** **Skeletons**

2D, iterative operations



```
// "Complex" kernel. Performs lots of trigonometric heavy lifting.
DATA_T func(input_matrix_t *img) {
    DATA_T sum = 0;

    // Iterate over all except outer neighbouring elements.
    for (int y = -SCL_NORTH + 1; y < SCL_SOUTH; y++) {
        for (int x = -SCL_WEST + 1; x < SCL_EAST; x++) {
            // Do *some* computation on values.
            DATA_T a = sin((float)getData(img, -1, 0));
            DATA_T b = native_sin((float)getData(img, 0,
1) * a);

            sum += getData(img, y, x) * a * (b / b);
        }
    }

    DATA_T out = 0;
    // Loop over horizontal region.
    for (int i = SCL_EAST; i >= -SCL_WEST; i--) {
        // DO *some* computation on values.
        * (DATA_T *) (i) = (float)getData(i, 0, i) % i;
    }
}
```

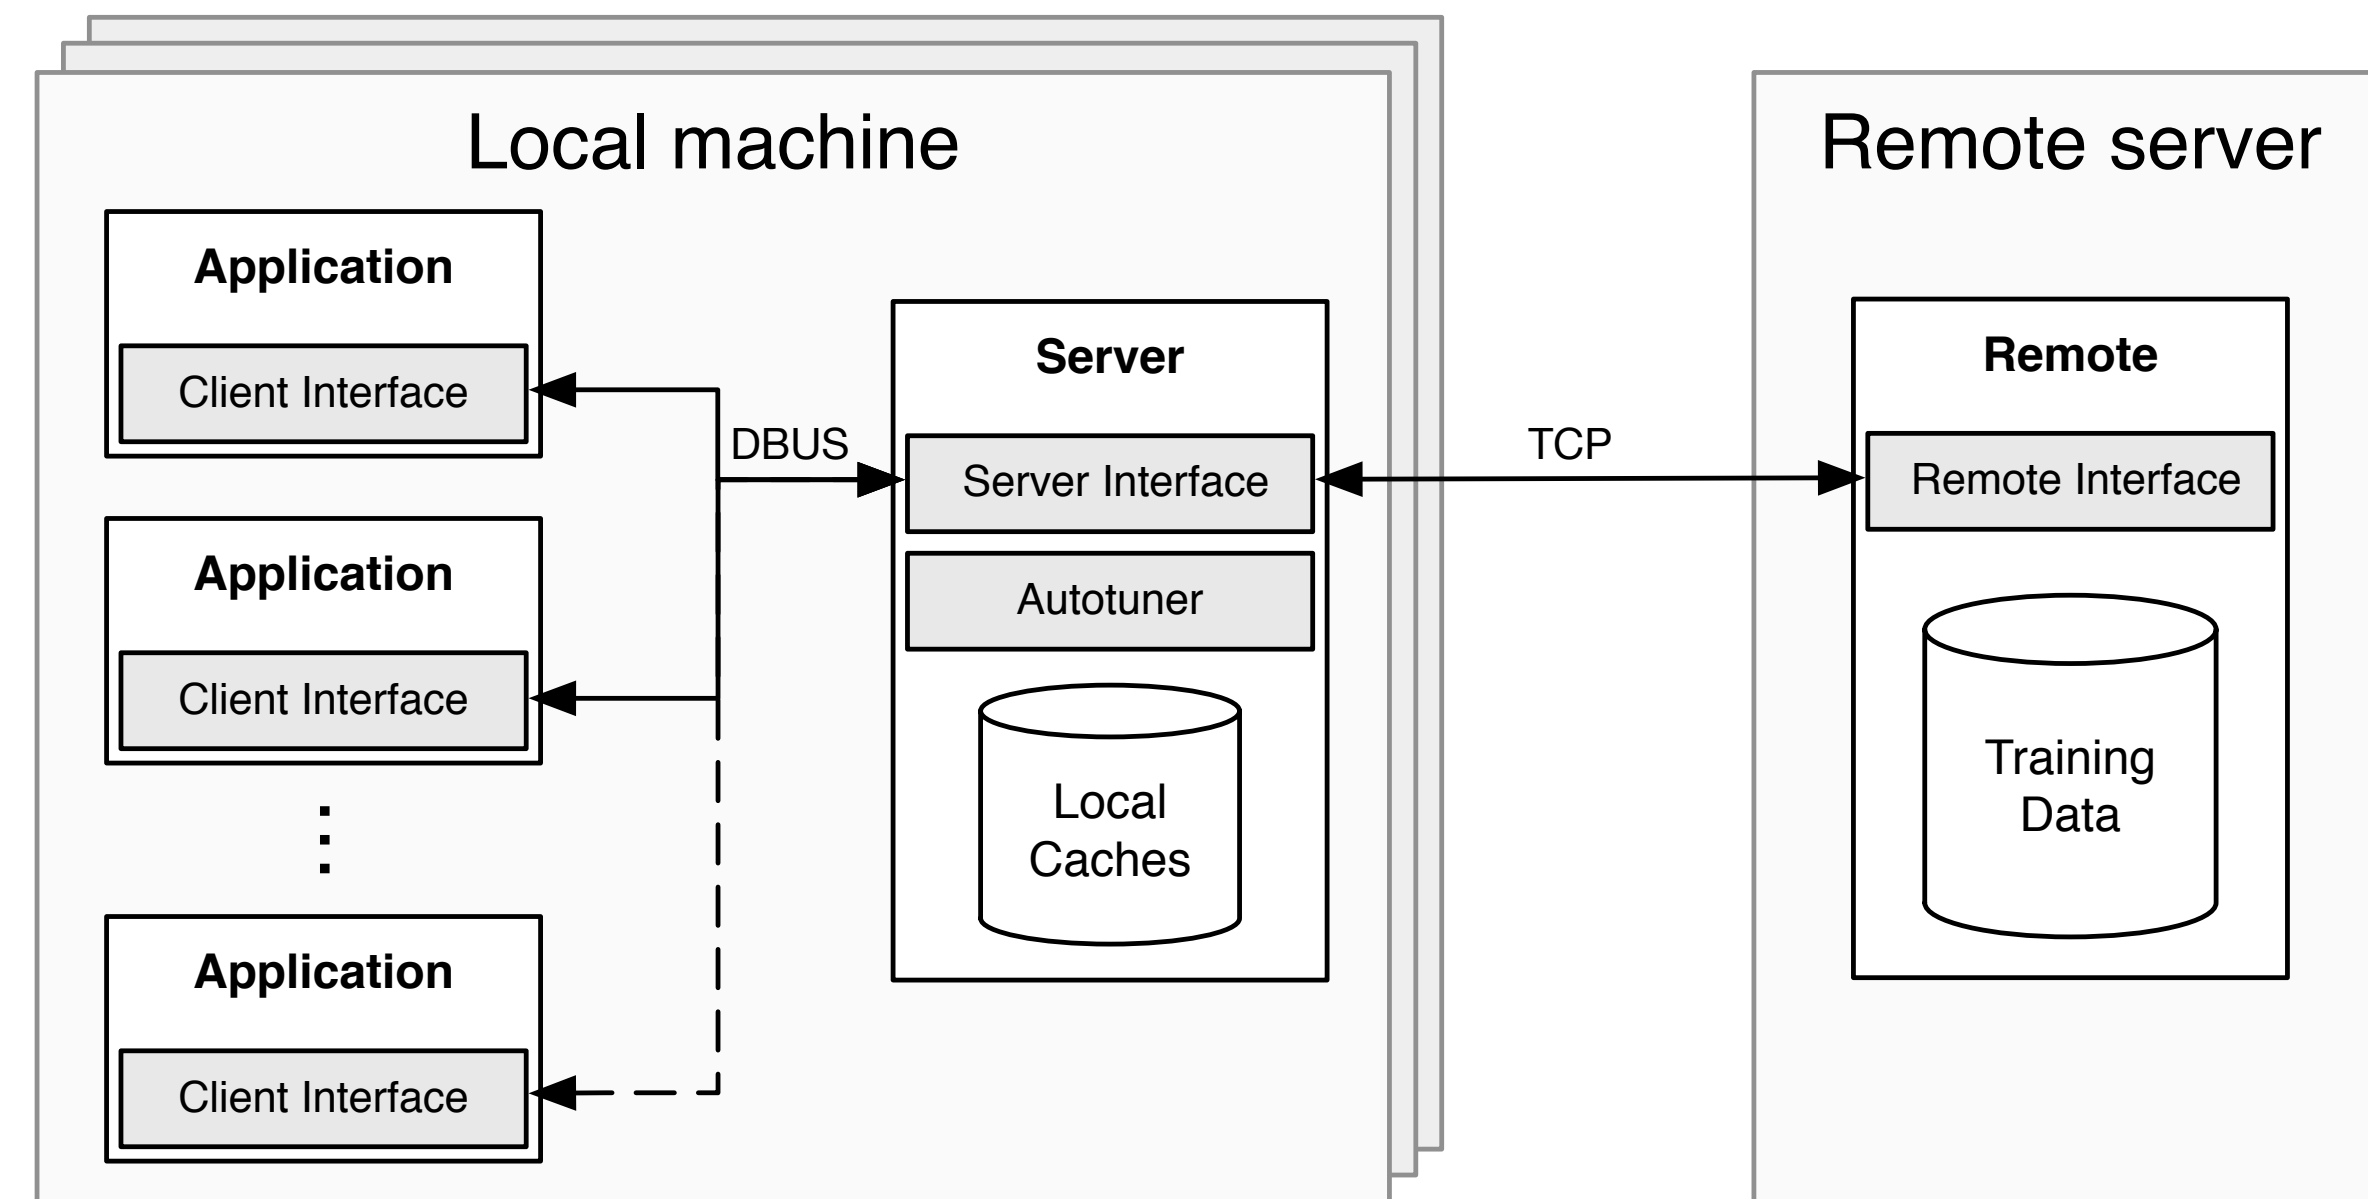
Synthetic benchmark generator

Thesis Autotuning **Stencil Codes** **with Algorithmic** **Skeletons**

Thesis Autotuning Stencil Codes with Algorithmic Skeletons

SkelCL
OpenCL
C++ Templates
GPUs

OmniTune



Thesis Autotuning Stencil Codes with Algorithmic Skeletons

Thesis Autotuning Stencil Codes with Algorithmic Skeletons

**How did
it go?**


Lessons learned

Scalable experiments

Write up often

Plan for planning

The Plan



Design



Implement

Proposal

**Skeleton-aware
compilation**

There is only so much optimisation that can be performed at the library level.

Can't optimise

**Across muscle
functions**

Across skeletons


SKELETON “COARSENING”.

LOAD BALANCING.

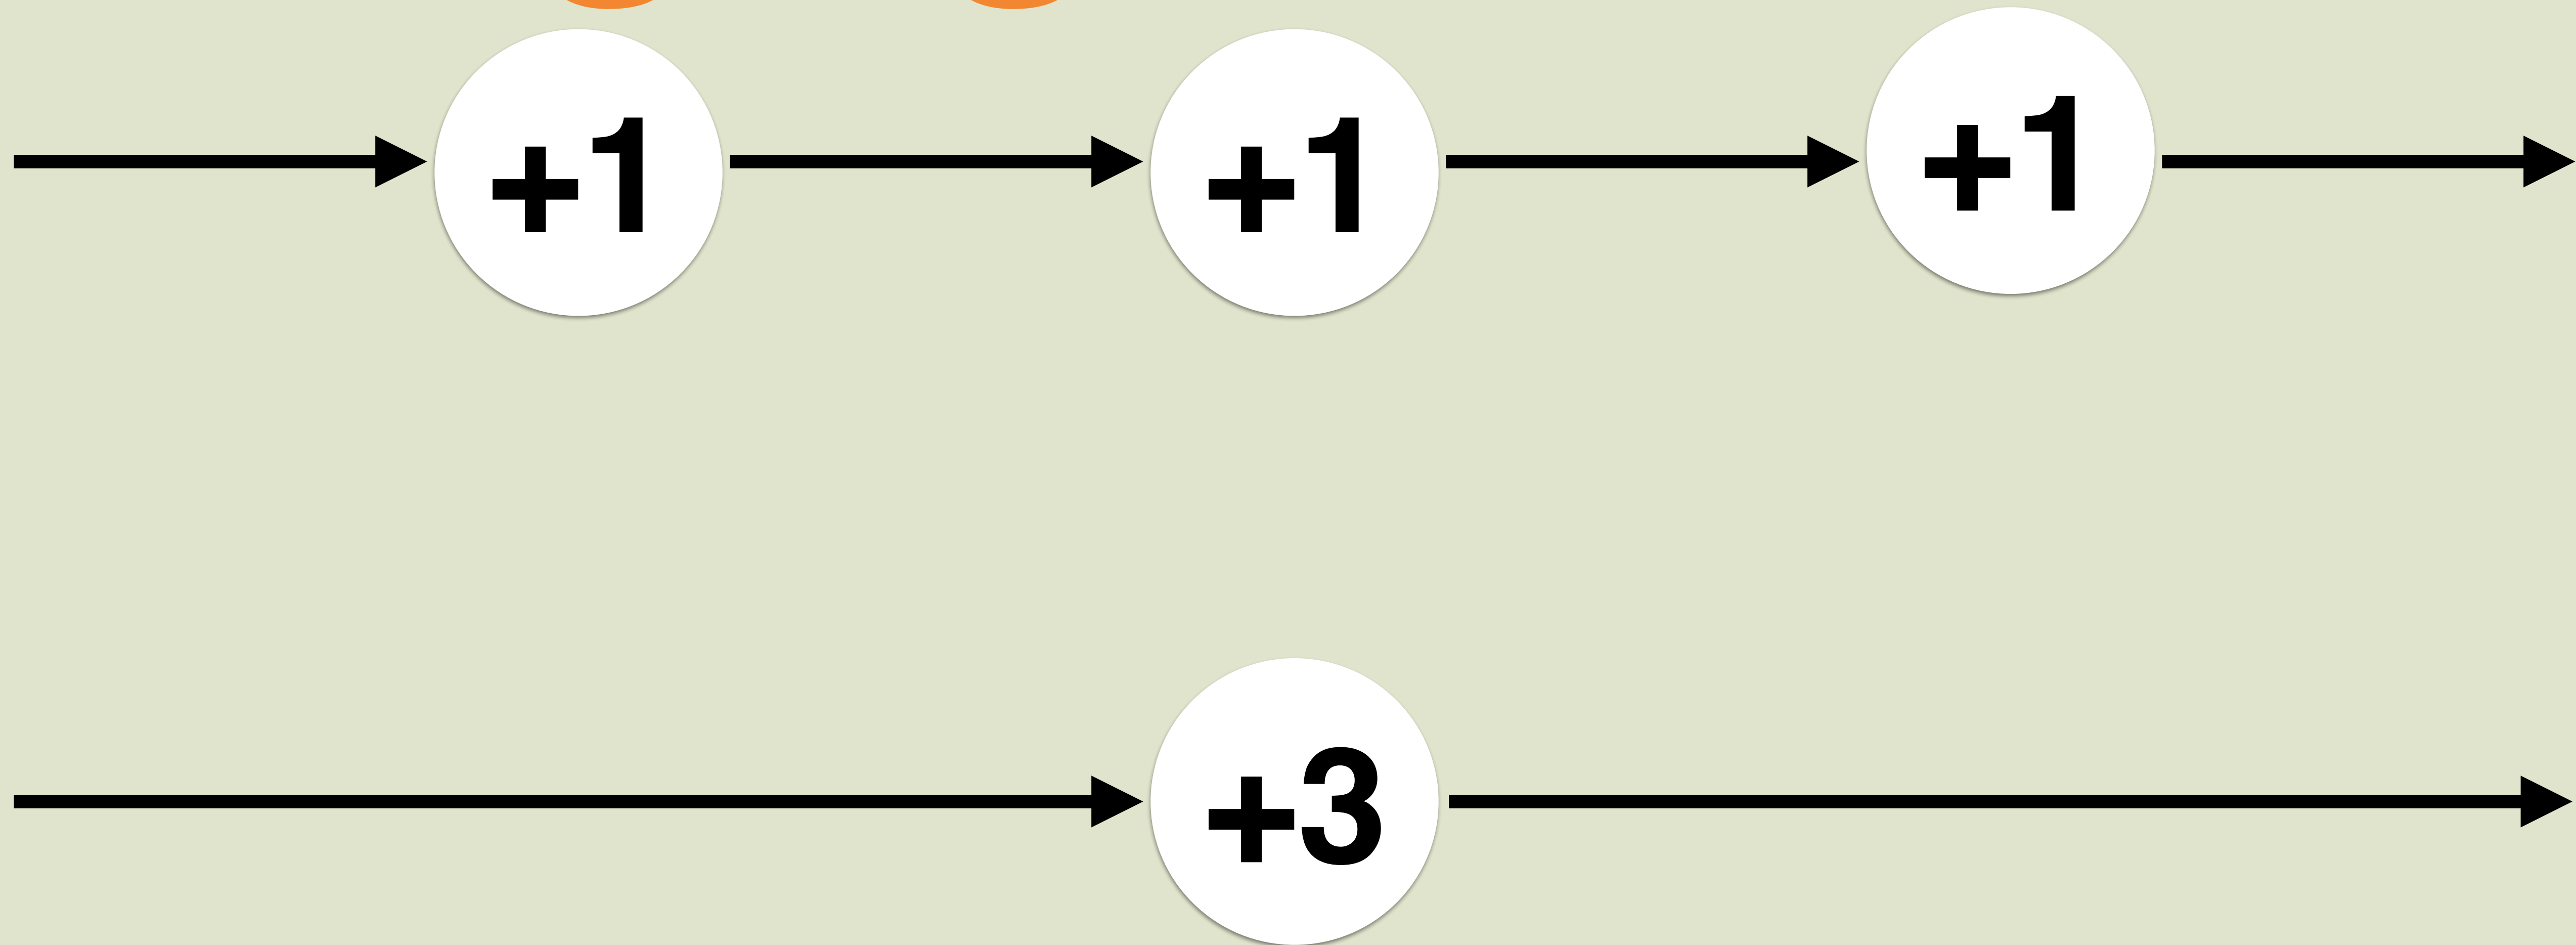
NESTED SKELETON RE-ORDERING.

CONSTANT PROPAGATION.



Merging



Merging



Load balancing



**What is the
simplest first step
that would need
to be taken?**

**Enumerate range of what
can be done.**

**Evaluate range of what
has been done.**

To test an optimisation

Implement by hand.

Measure performance.

Implement compiler pass.

Measure success.

Composition

```
(pmap f4  
  (pmap f3  
    (pmap f2  
      (pmap f1 x))))
```

Composition

```
(pmap (comp f4 f3 f2 f1) x)
```


Composition

(pmap (comp f4 f3 f2 f1) x)

1.81x faster
(1e6 floats, 30 samples)

Summary

My TODO list



1 review the literature.

2 hand-test optimisations

3 publish a workshop paper

(ADAPT or
HLGP GPU)

**Skeleton-aware
compilation**

