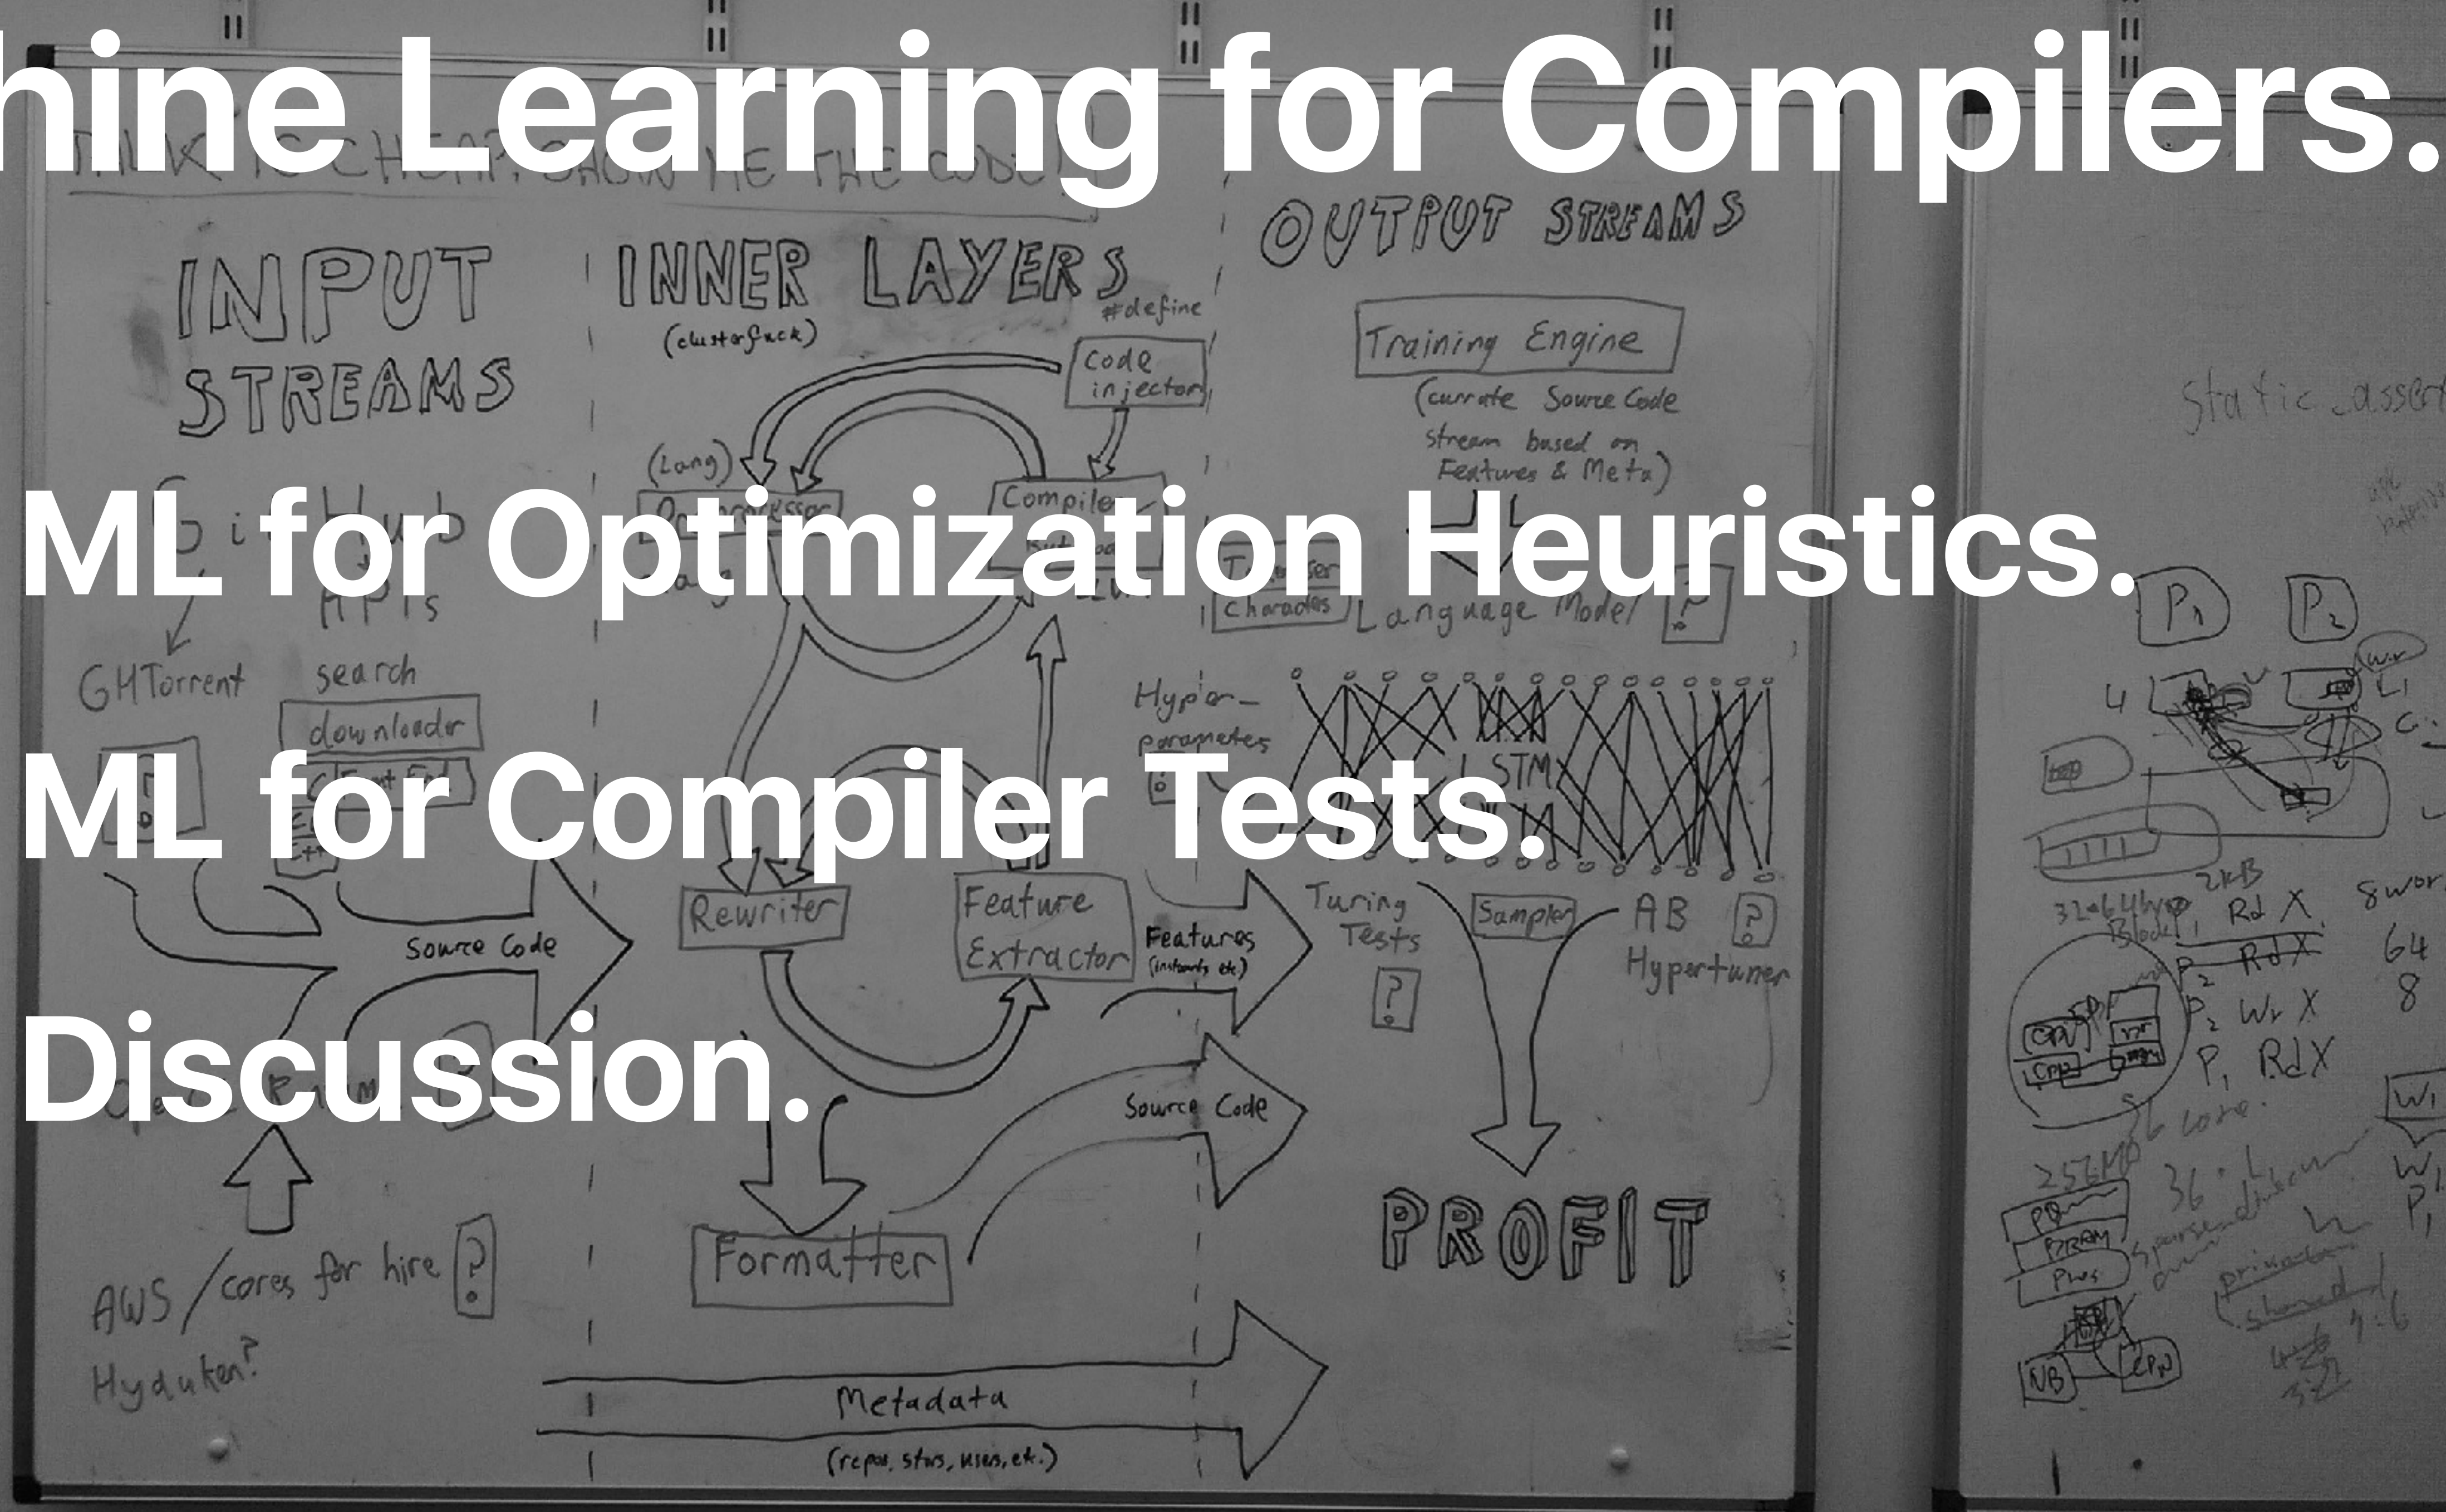


Machine Learning for Compilers.

Chris Cummins

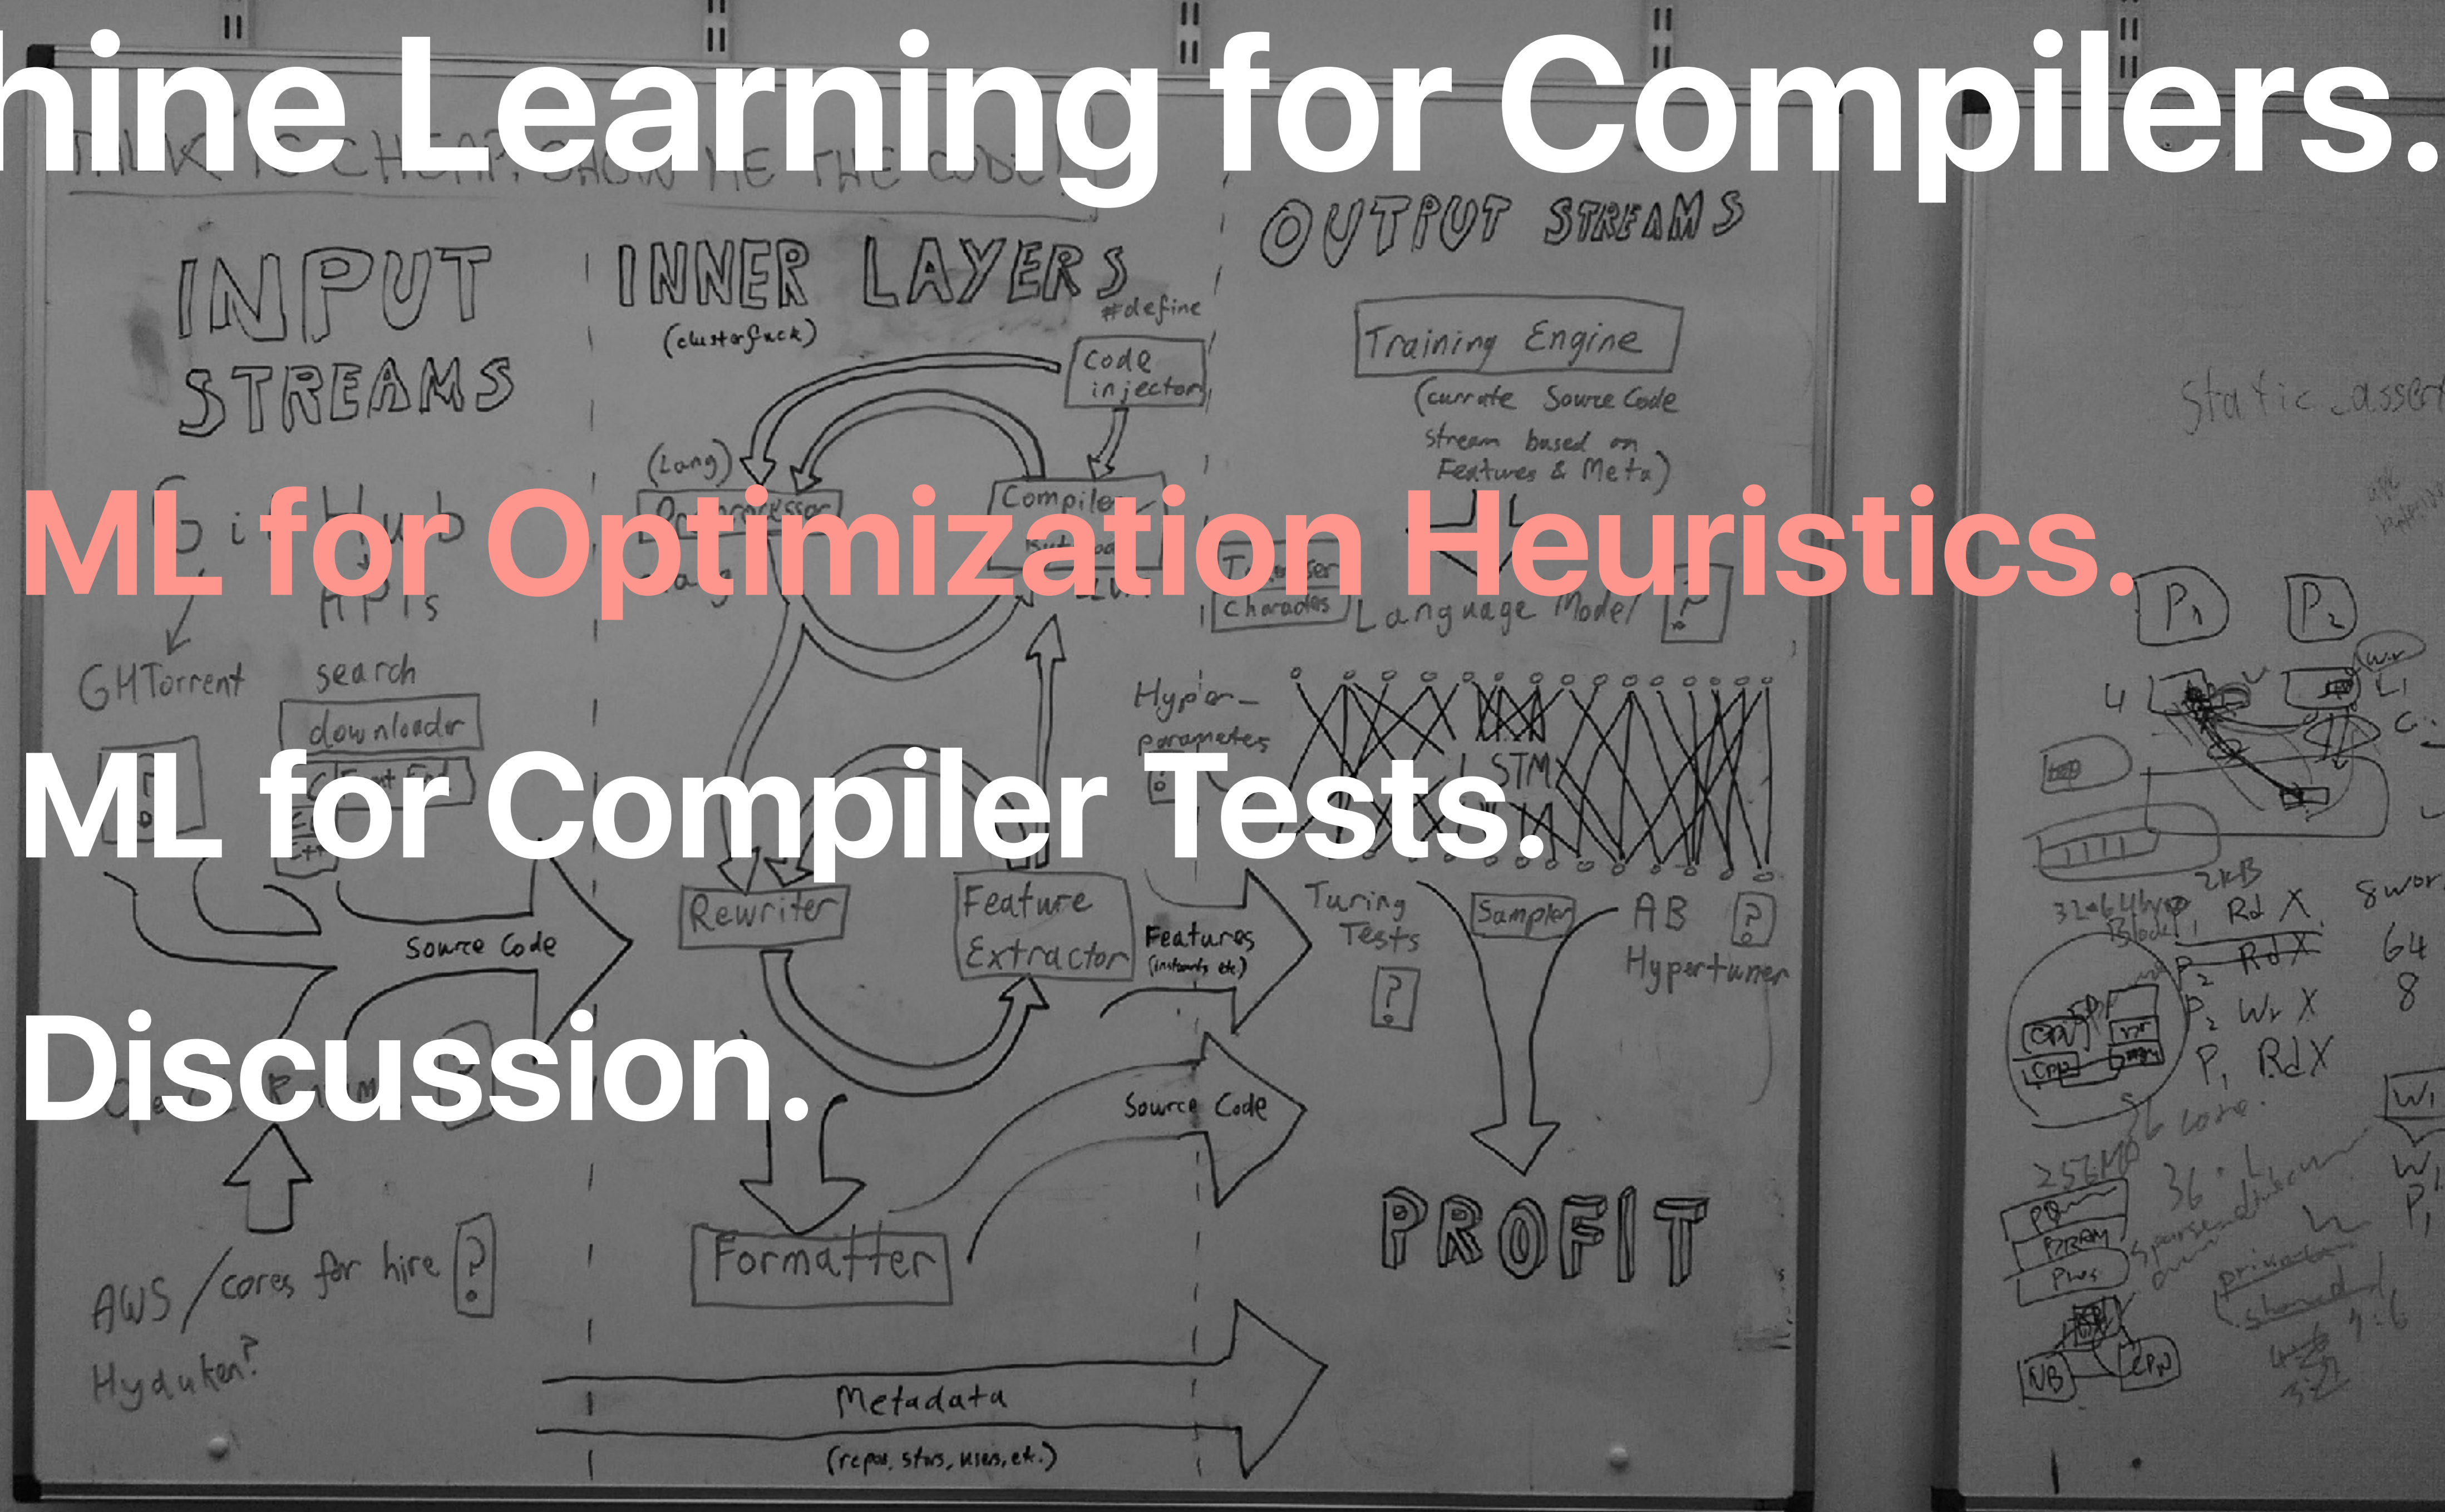
Machine Learning for Compilers.

1. ML for Optimization Heuristics.
2. ML for Compiler Tests.
3. Discussion.



Machine Learning for Compilers.

1. ML for Optimization Heuristics.
2. ML for Compiler Tests.
3. Discussion.




Problem:

Optimization heuristics are too coarse.

(aka. it doesn't take much to do better than -O3)

compilers are Very complex



Hand-coded heuristics:

Analytical models difficult

Huge number of variables

Architectures change

Other bits of the compiler change

Execution is non-deterministic (OoO exec, caches)

Simple components are complex together

Better than -O3

```
#!/bin/sh
```

```
while true; do  
    sort --random-sort < "cflags.txt" | head -n 20 | xargs gcc -O1 app.c  
    time ./a.out  
done
```

after 200 attempts, ~5% speedup

Iterative Compilation

Very simple idea

Don't get best heuristic analytically


Try many versions of heuristic and choose best

Platform blind

Adapts to every change

Always based on evidence not belief

Example: Stencil Workgroup size



Workgroup size affects





mapping to SIMD hardware.



device occupancy.

local memory utilisation.



Choosing workgroup size depends on



- 1. Device**
- 2. Program**
- 3. Dataset**







Avg 15x max speedup











Same program,
different device



same device,
different program





324 attempts sounds like a lot...

($324 \times 10 \times 30 \text{ s} \approx 27 \text{ hrs}$)


Drop in the ocean

250 GCC flags $\geq 2^{250}$ options $\approx 10^{75}$

Atoms in the universe $\approx 10^{80}$


Machine Learning

Estimate $y = f(x)$



Machine Learning

Estimate $y = f(x)$



Machine Learning

Estimate $y = f(x)$

Optimisations

Cflags

Workgroup size

CPU or GPU

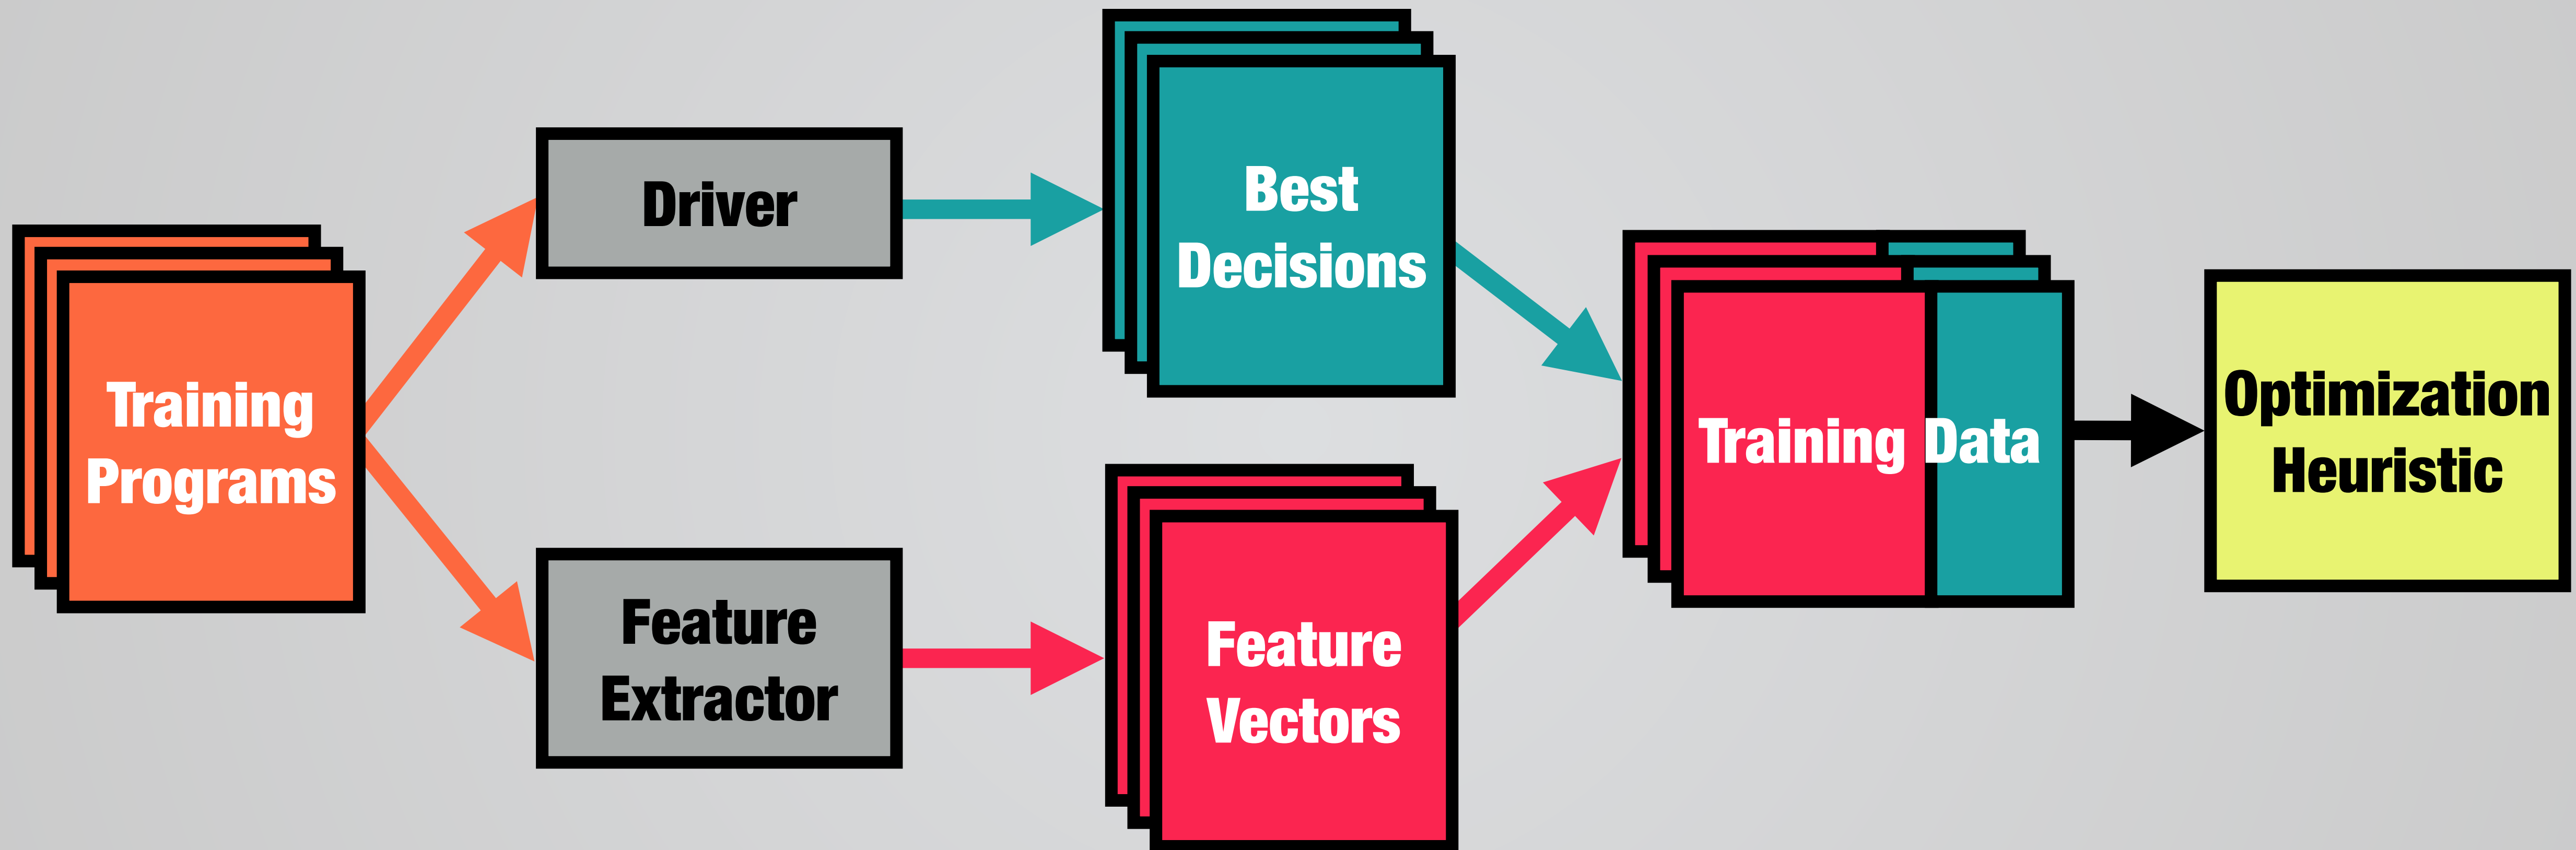
Features

instructions

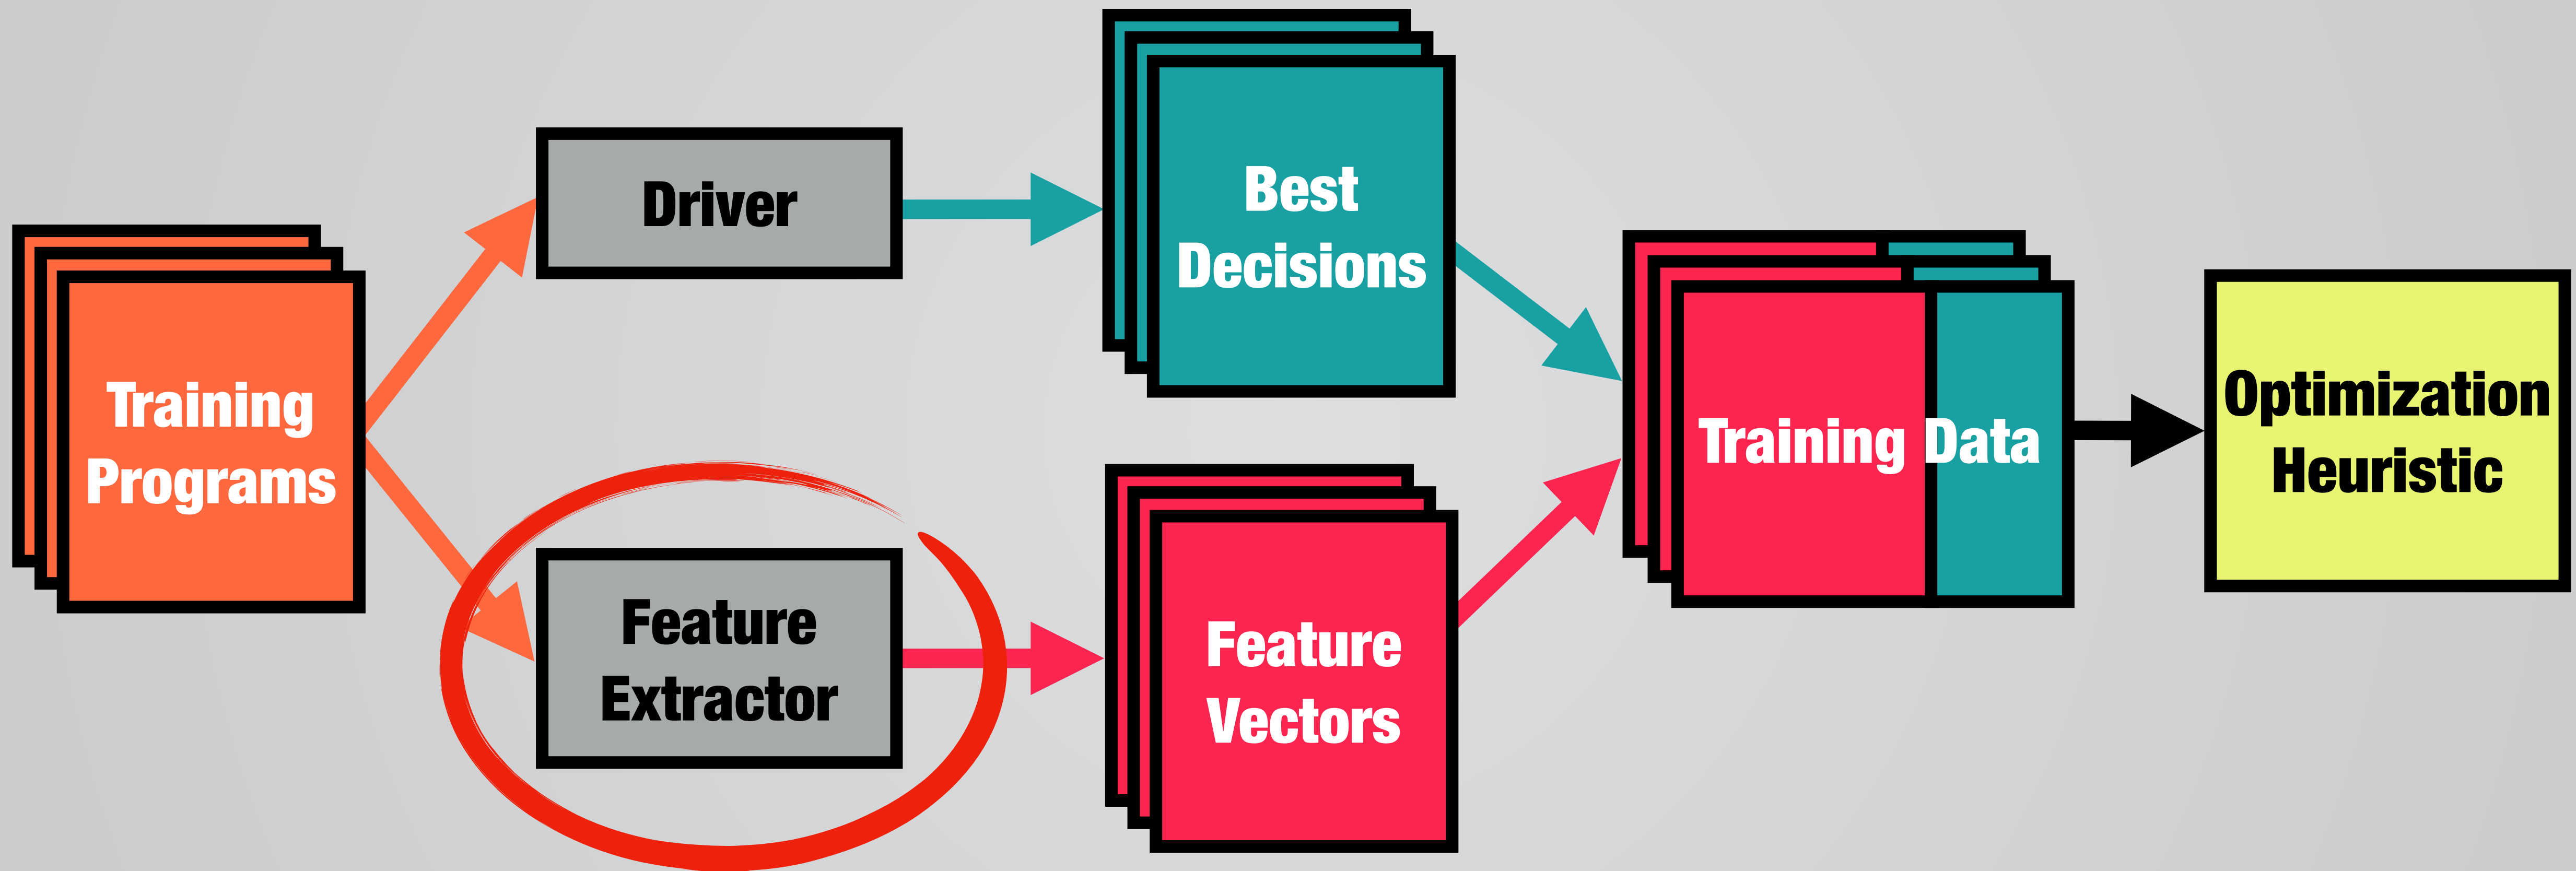
Arithmetic density

Dataset size

Learning an Optimization Heuristic



Learning an Optimization Heuristic




Features for Stencil Workgroup size

- 1. Device**
- 2. Program**
- 3. Dataset**

Features for Stencil Workgroup size

1. **Device**
2. **Program**
3. **Dataset**



or



How many compute units?

How much memory?

Cache size?

etc.



Features for Stencil Workgroup size

1. **Device**
2. **Program**
3. **Dataset**

Features for Stencil Workgroup size

1. Device
2. Program
3. Dataset


How big is border region?

What shape is it?

How many instructions?

What type of instructions?

etc.



Features for Stencil Workgroup size

1. Device
2. Program
3. Dataset

Features for Stencil Workgroup size


1. Device
2. Program
3. Dataset

$$\mathbf{A} = \begin{bmatrix} A_{11} & A_{12} & \cdots & A_{1n} \\ A_{21} & & & A_{2n} \\ \vdots & & & \vdots \\ A_{n1} & A_{n2} & \cdots & A_{nn} \end{bmatrix}$$

How big is the data?

What type is the input?

What type is the output?



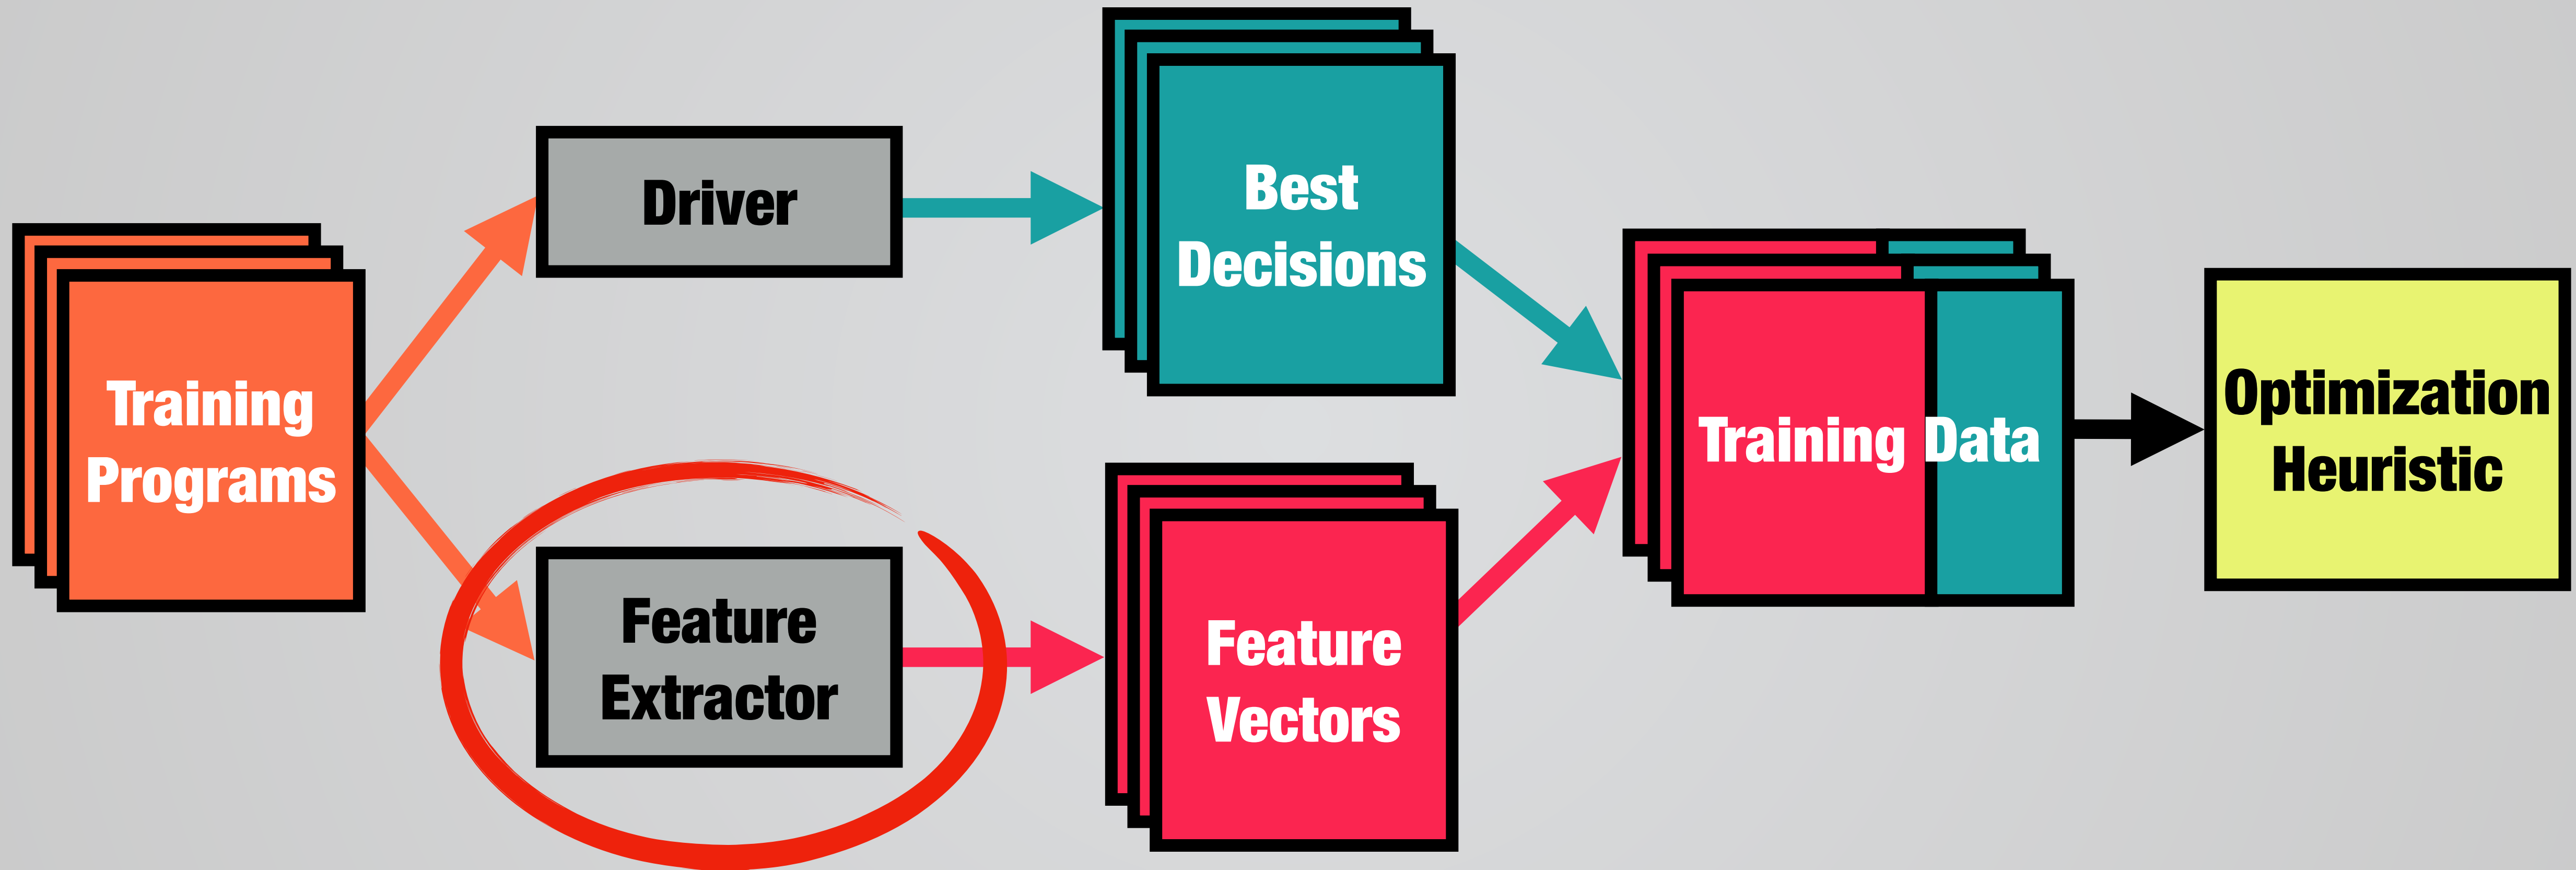
Features for Stencil Workgroup size

1. Device
2. Program
3. Dataset

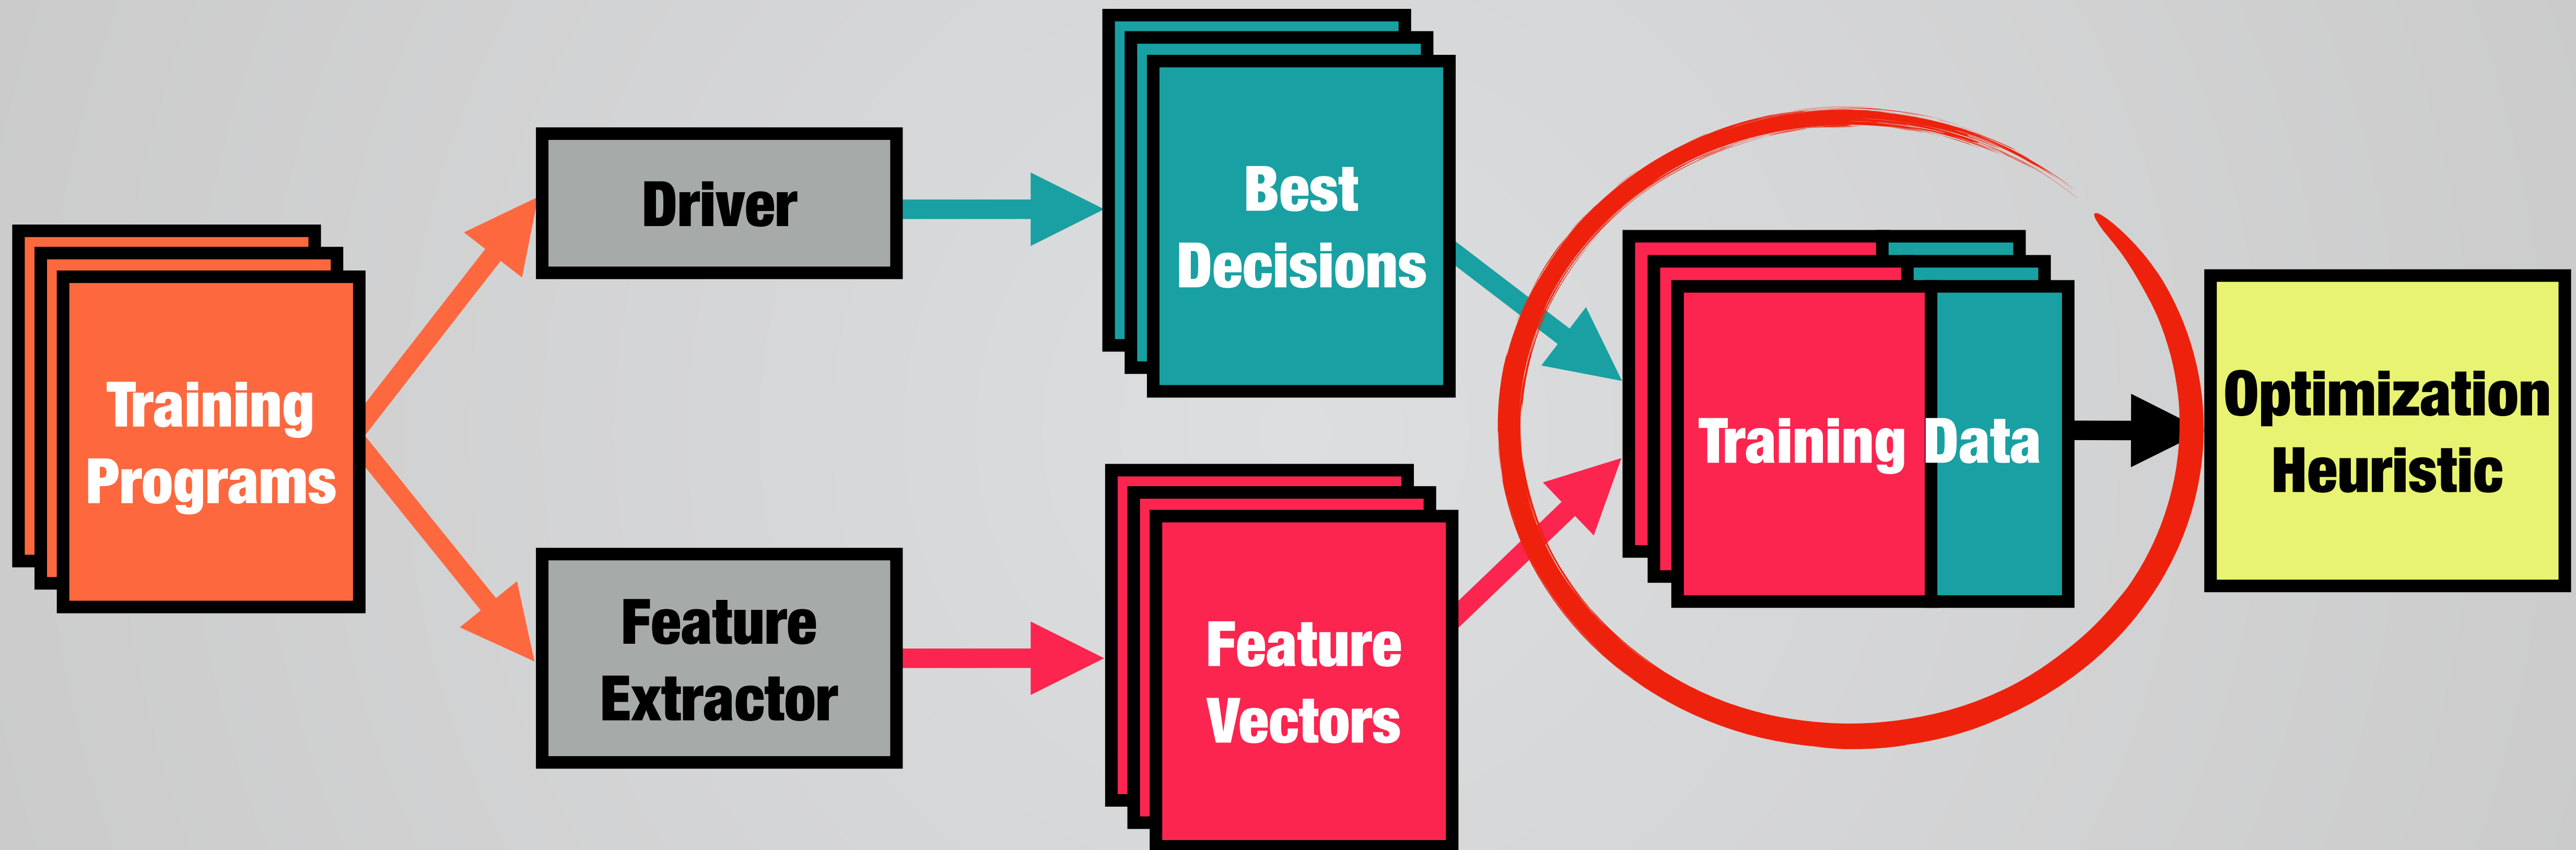
Features for Stencil Workgroup size


- 1. Device**
- 2. Program**
- 3. Dataset**

Learning an Optimization Heuristic



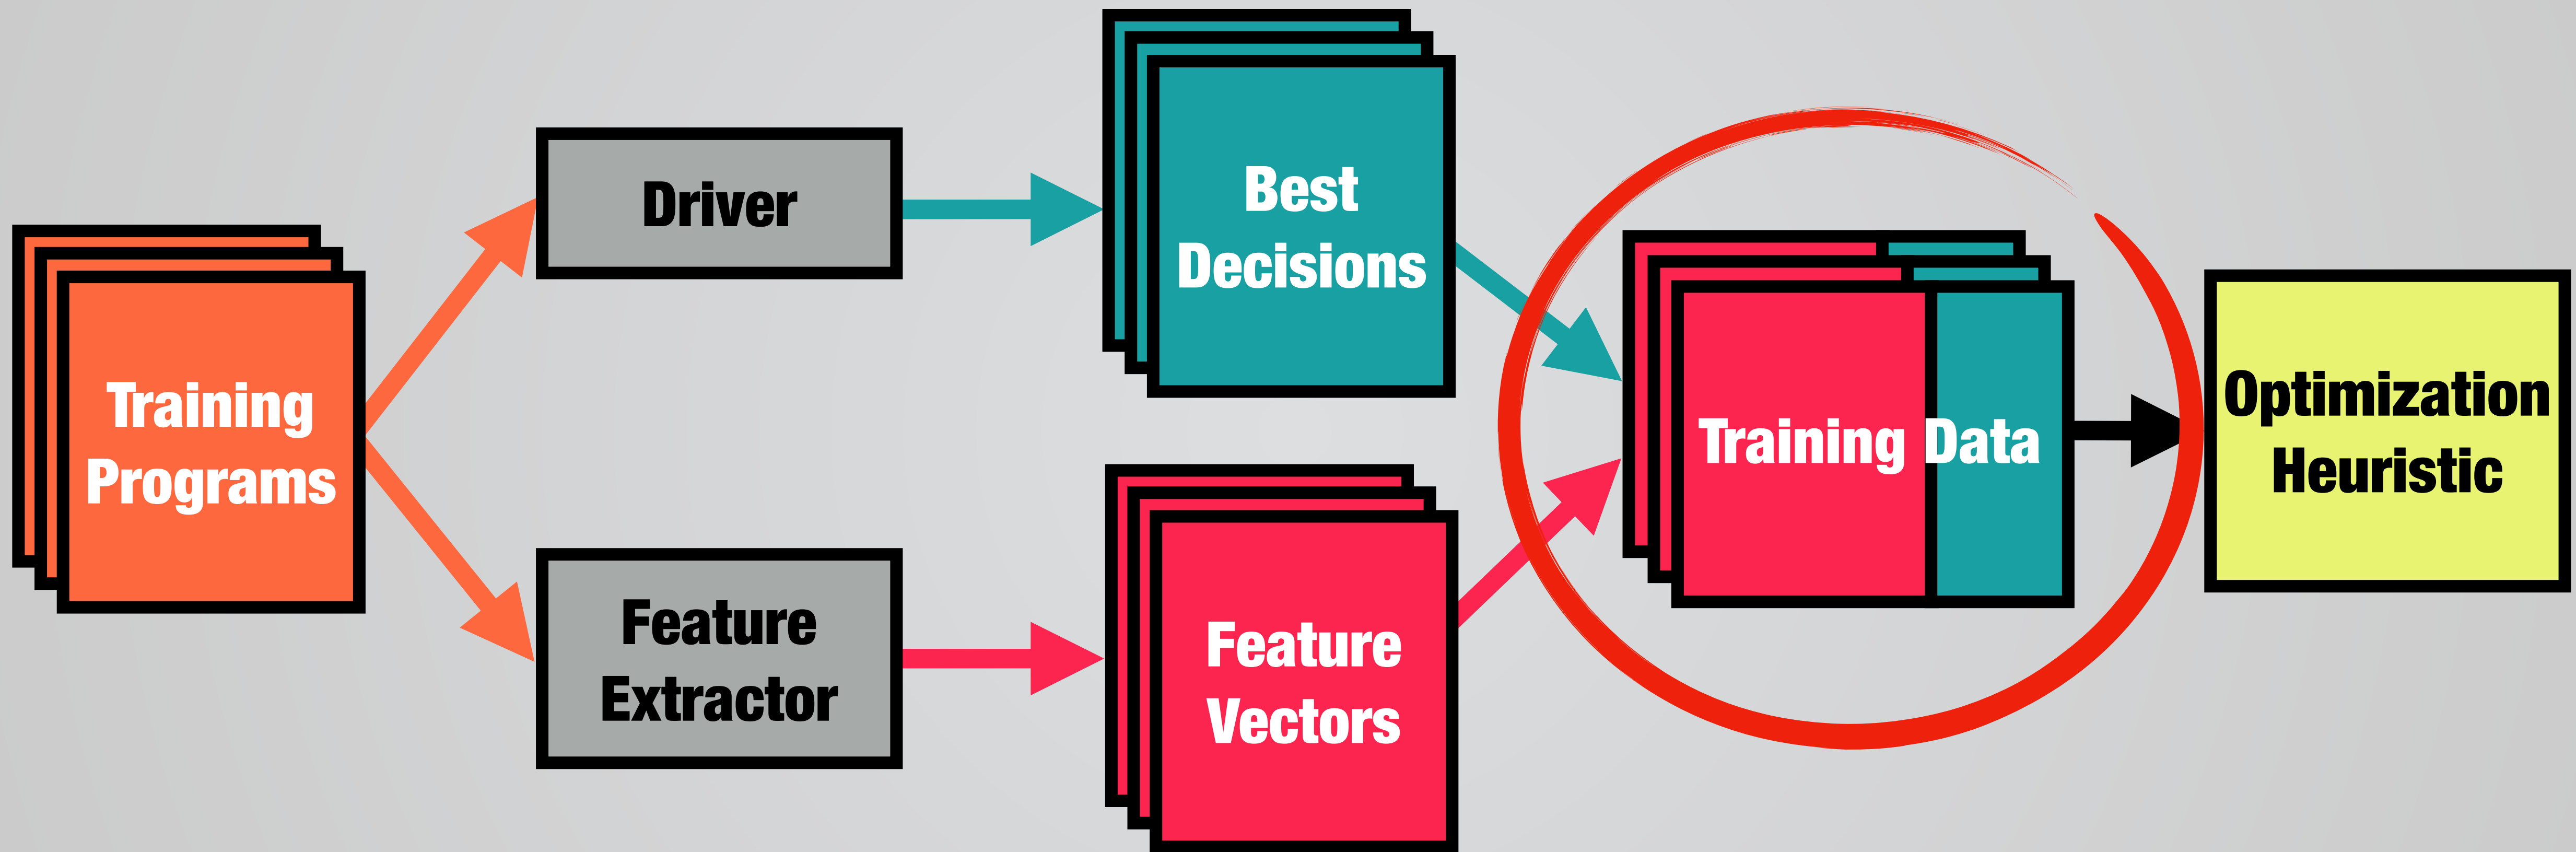
Learning an Optimization Heuristic



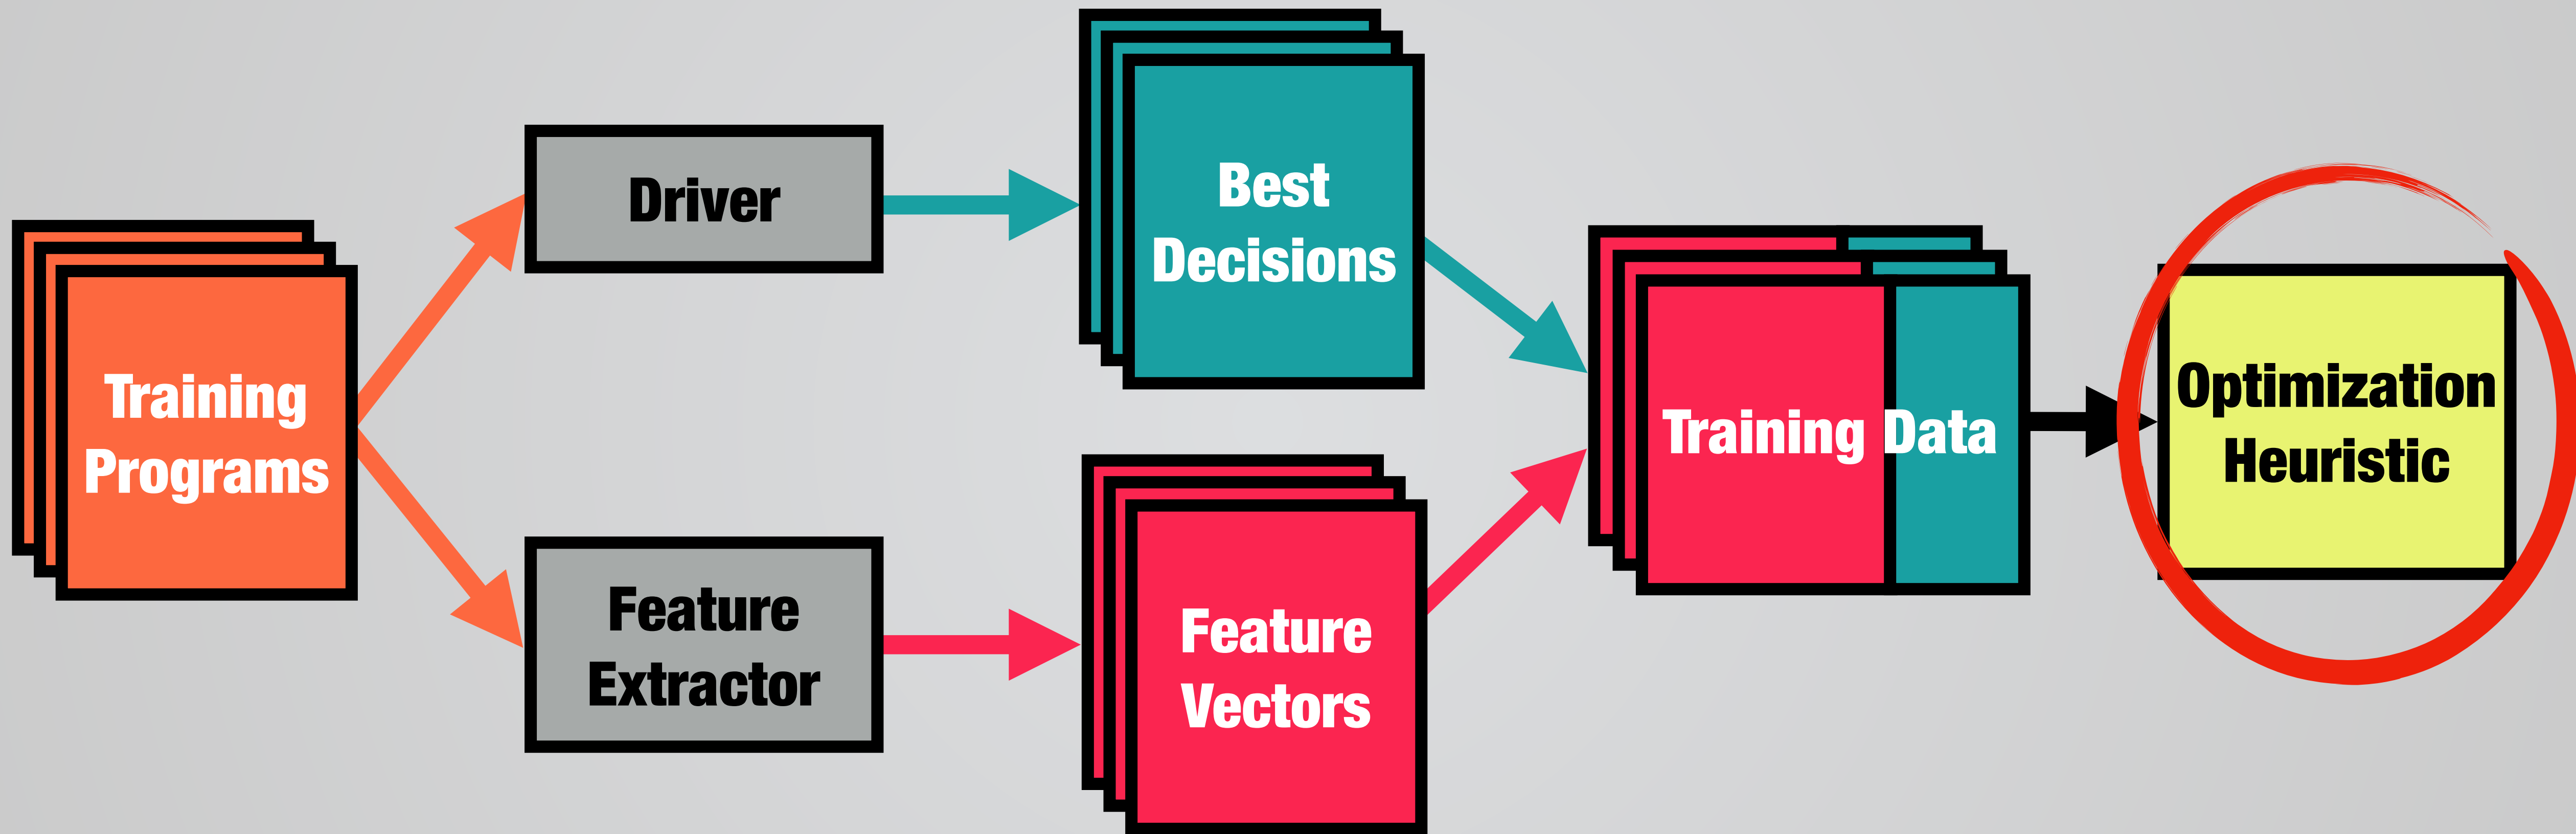




Learn a classification (*dotted line*)



Learning an Optimization Heuristic




Learning an Optimization Heuristic









26%
optimal





90% optimal



26% optimal

ML for optimization heuristics.

Advantages:

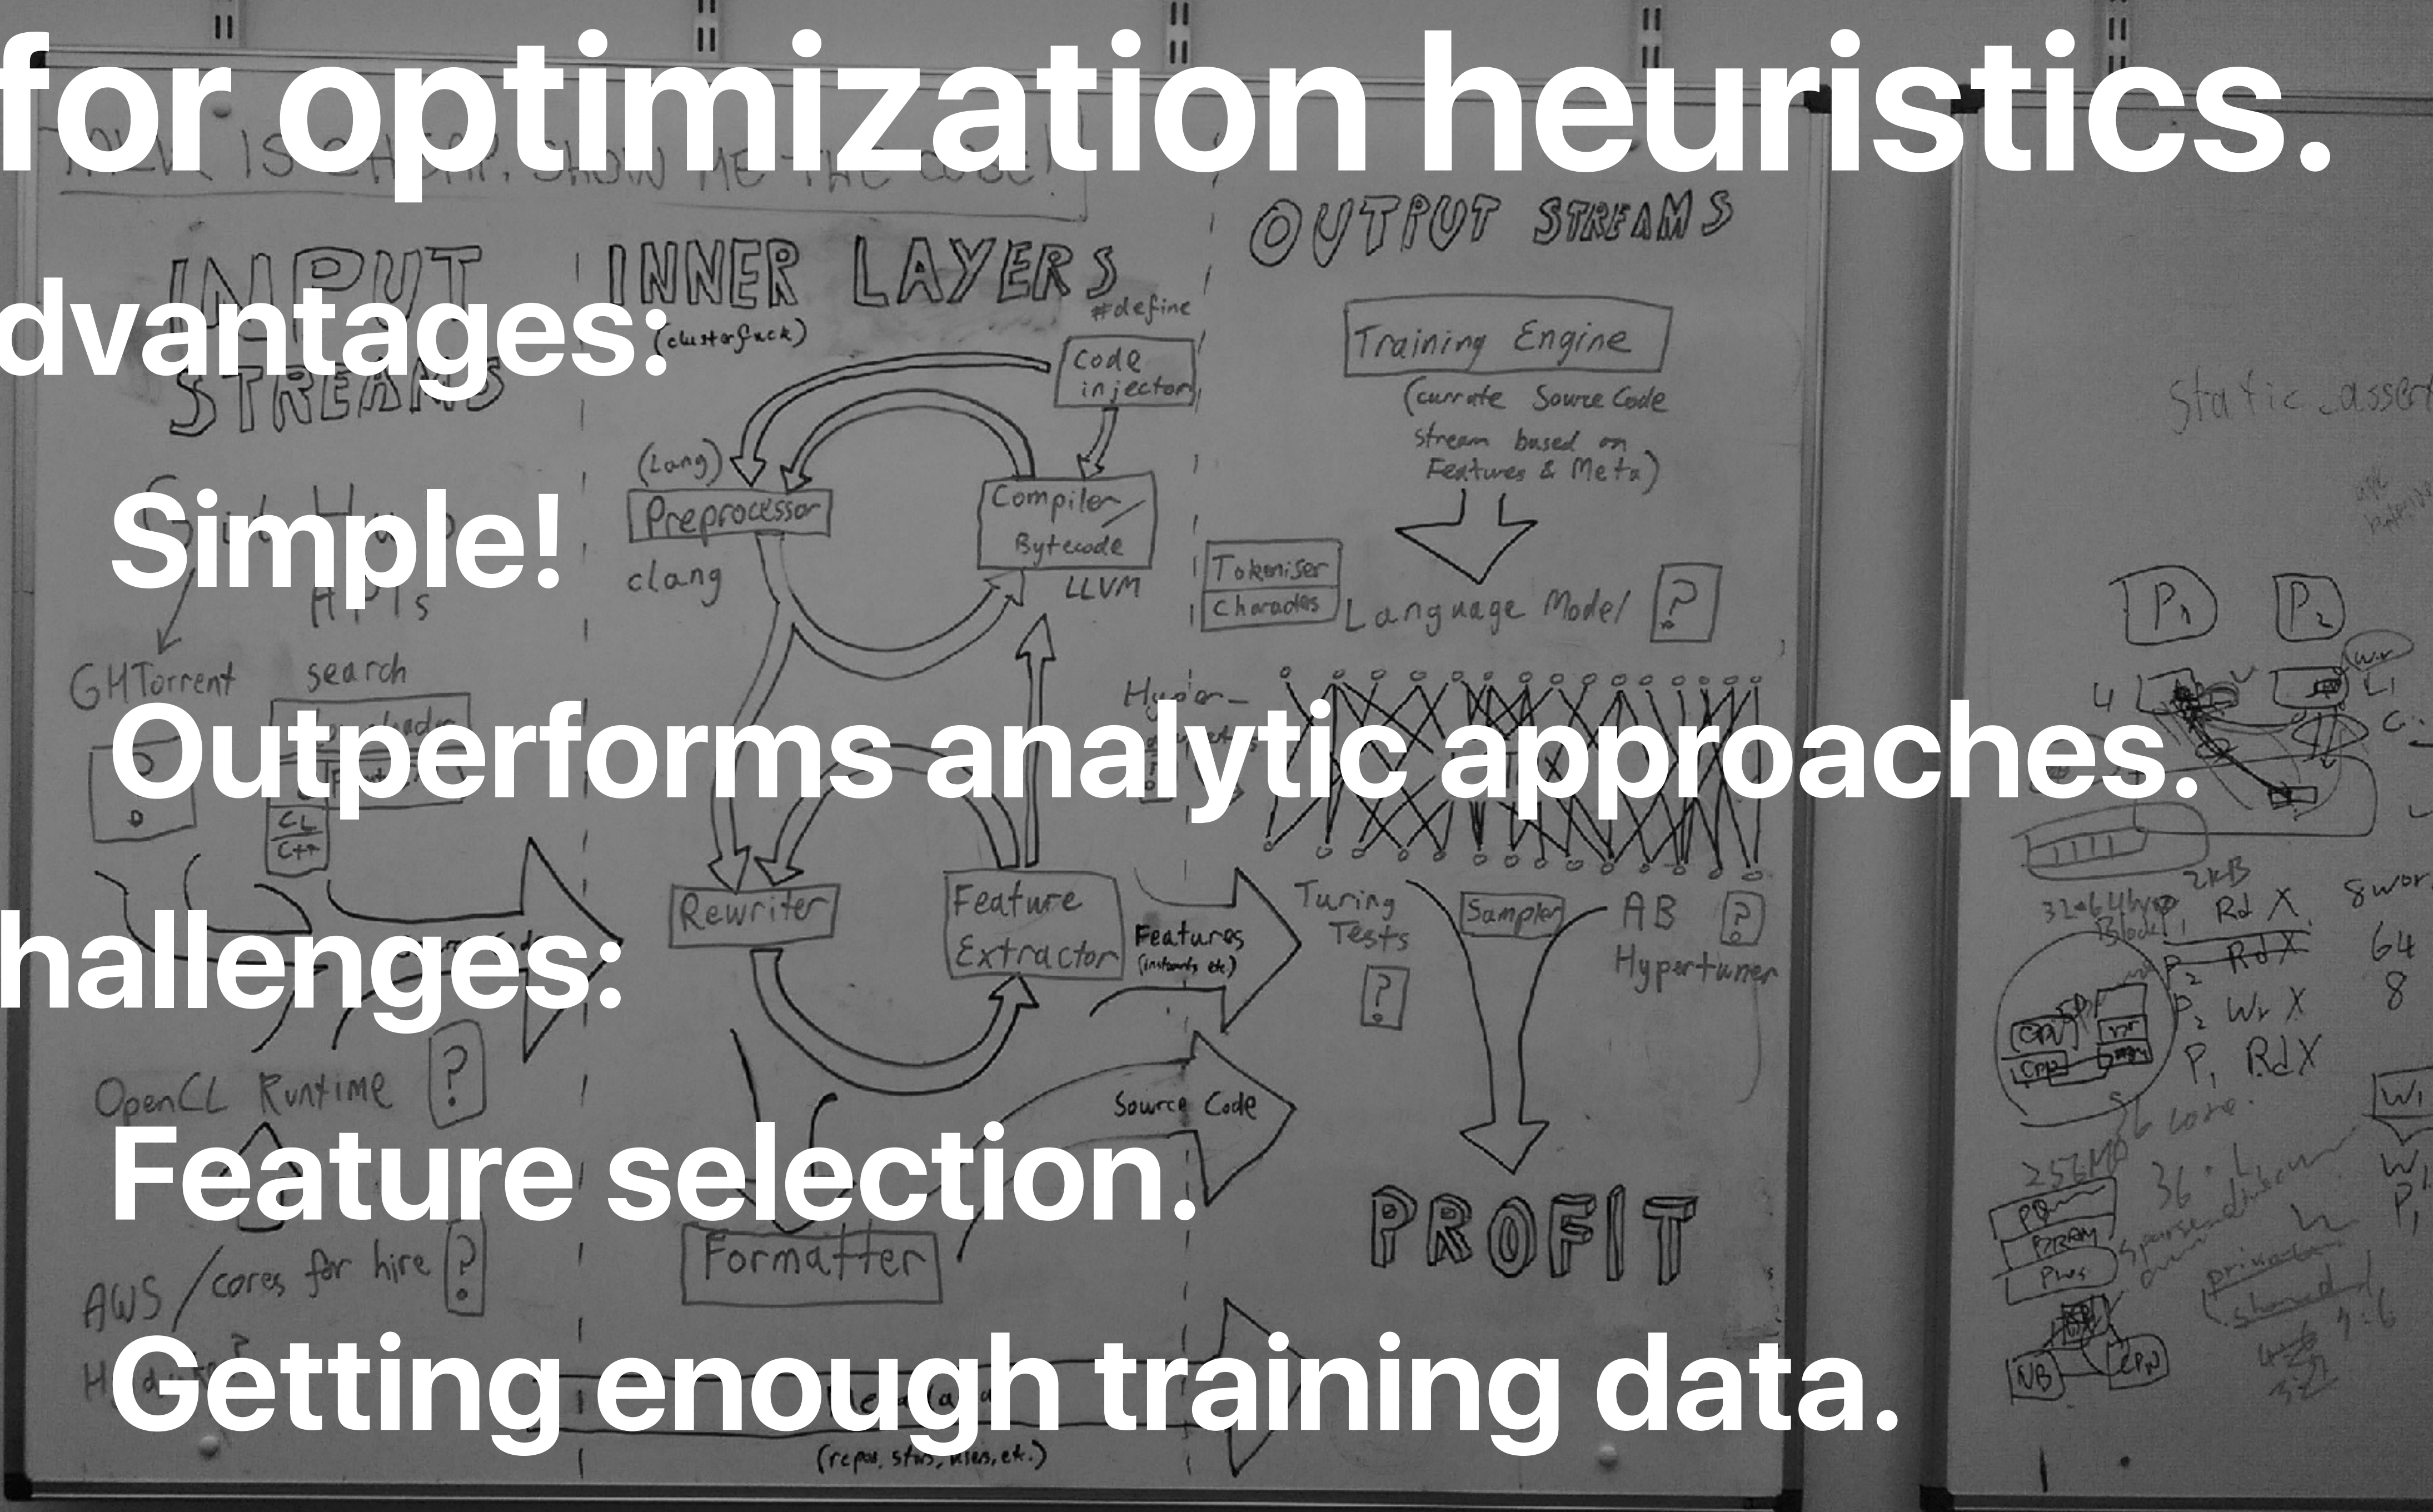
Simple!

Outperforms analytic approaches.

Challenges:

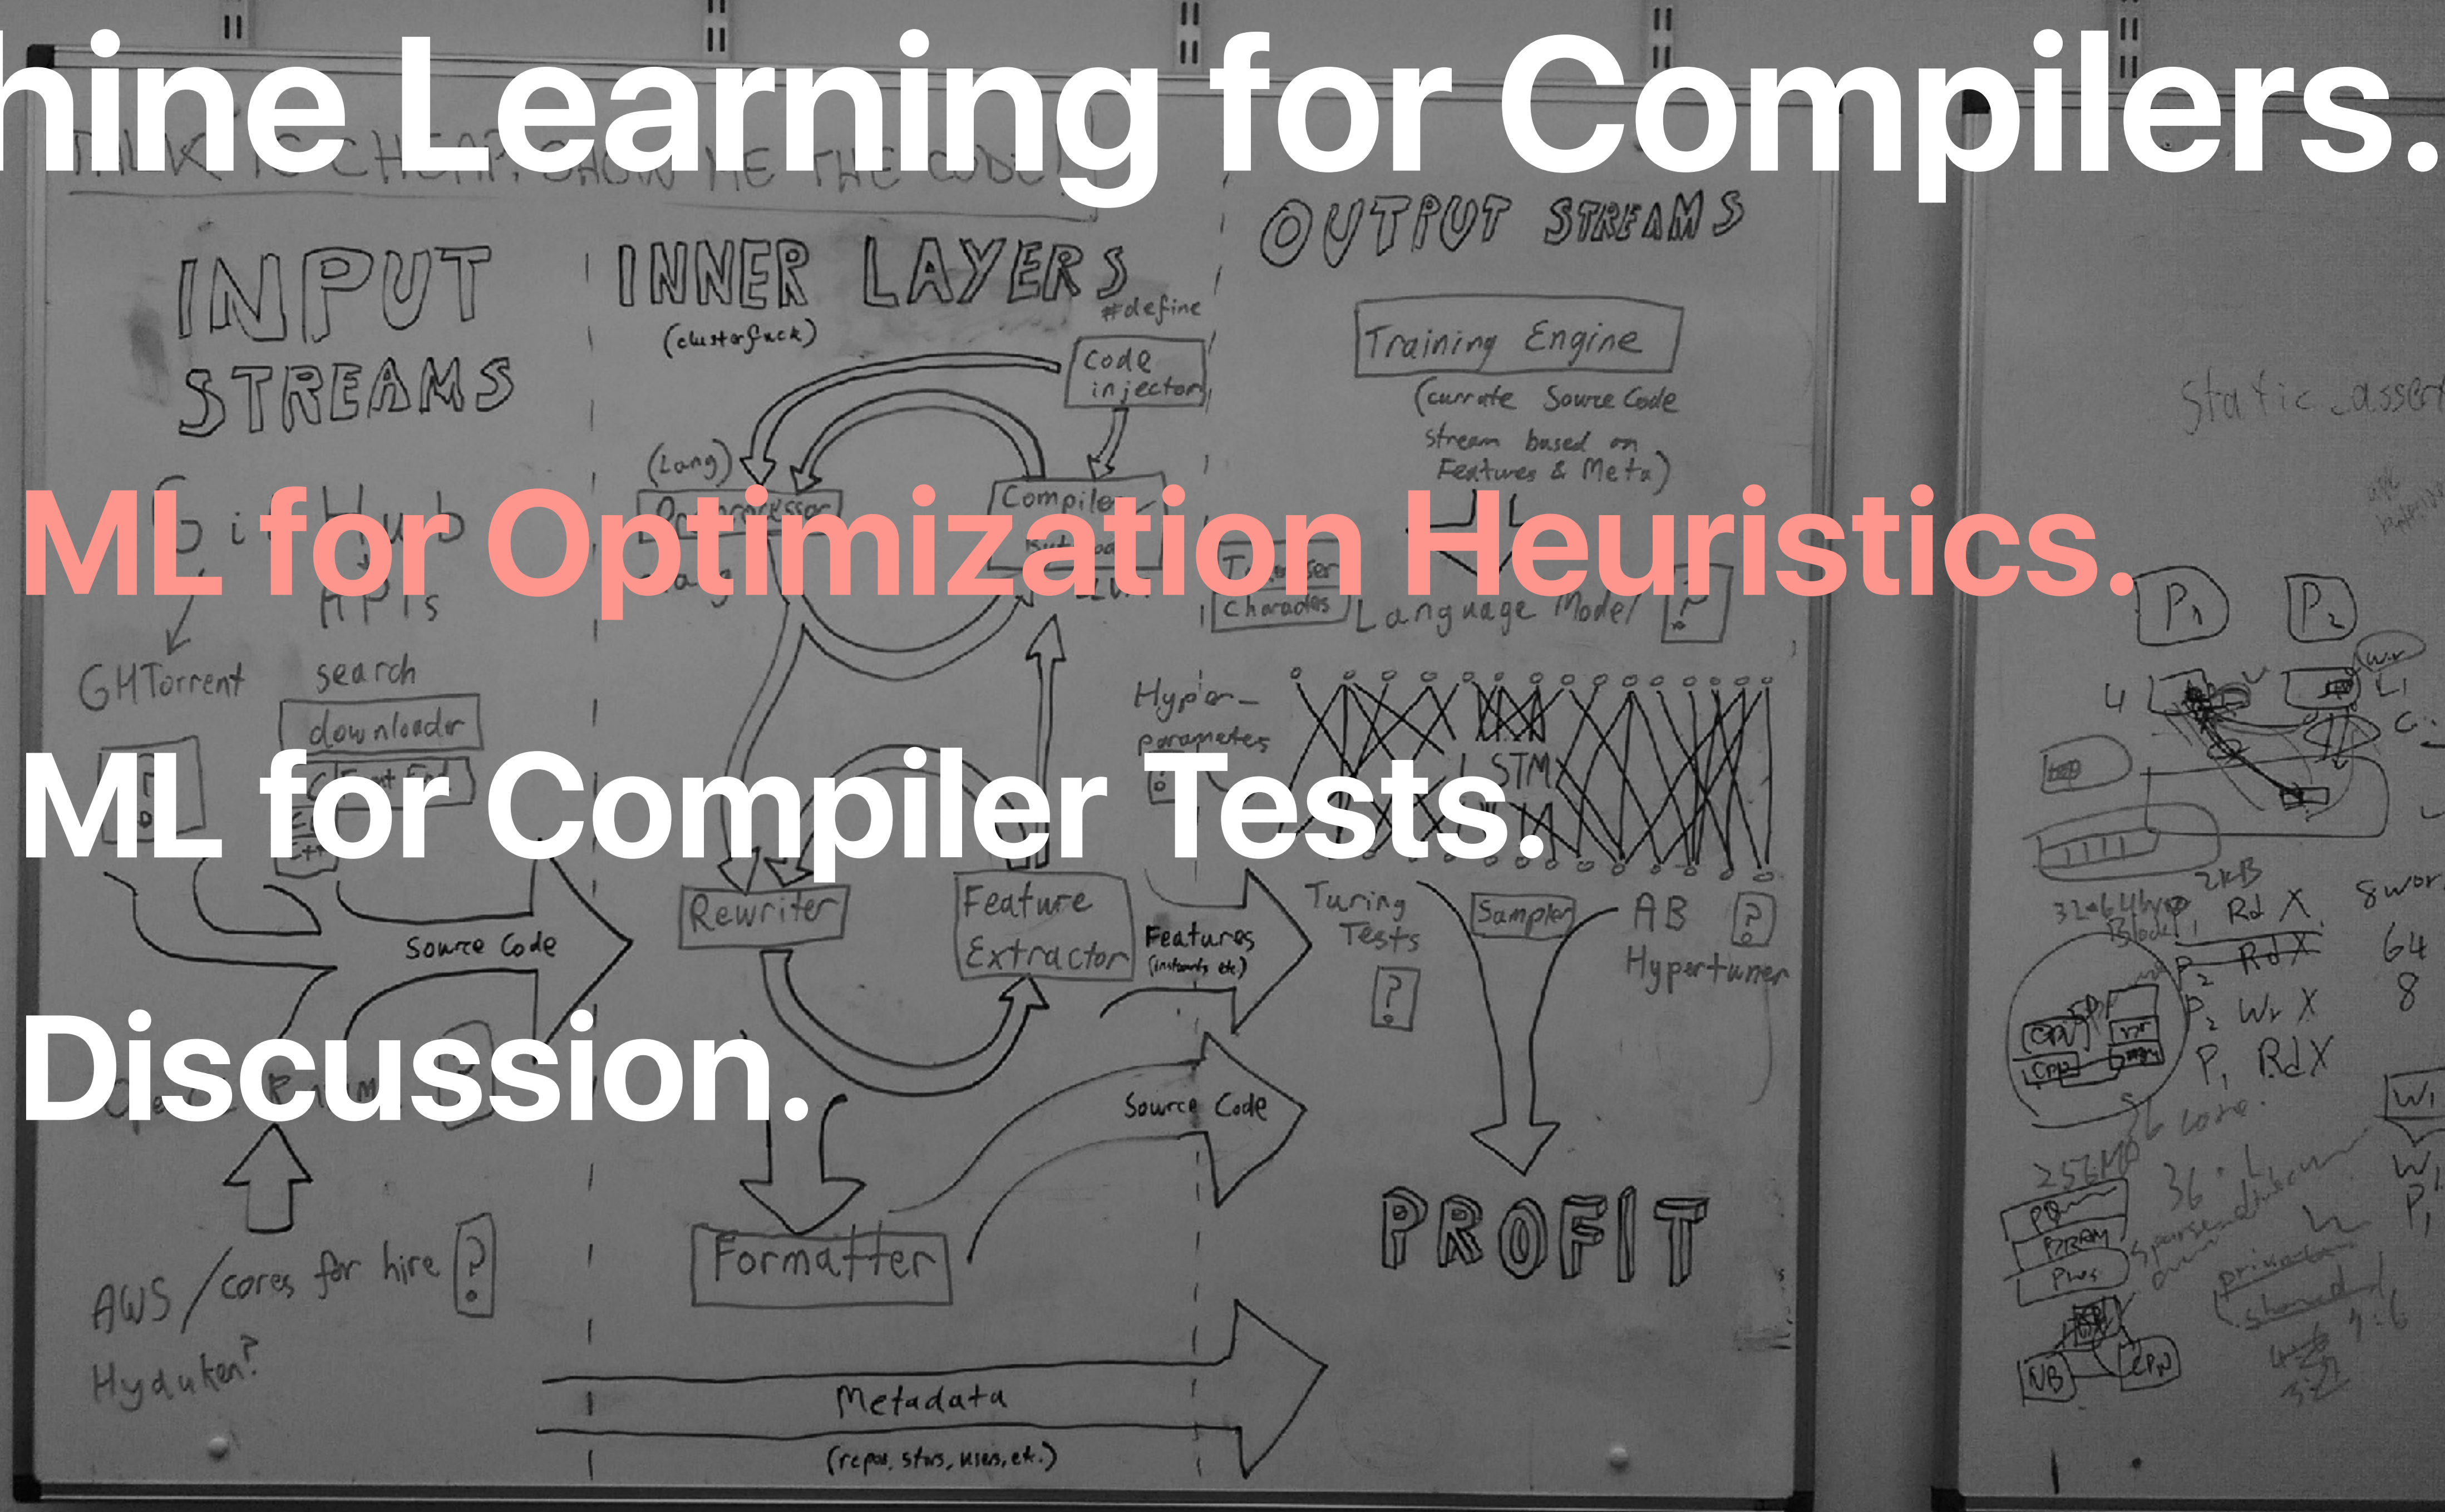
Feature selection.

Getting enough training data.



Machine Learning for Compilers.

1. ML for Optimization Heuristics.
2. ML for Compiler Tests.
3. Discussion.

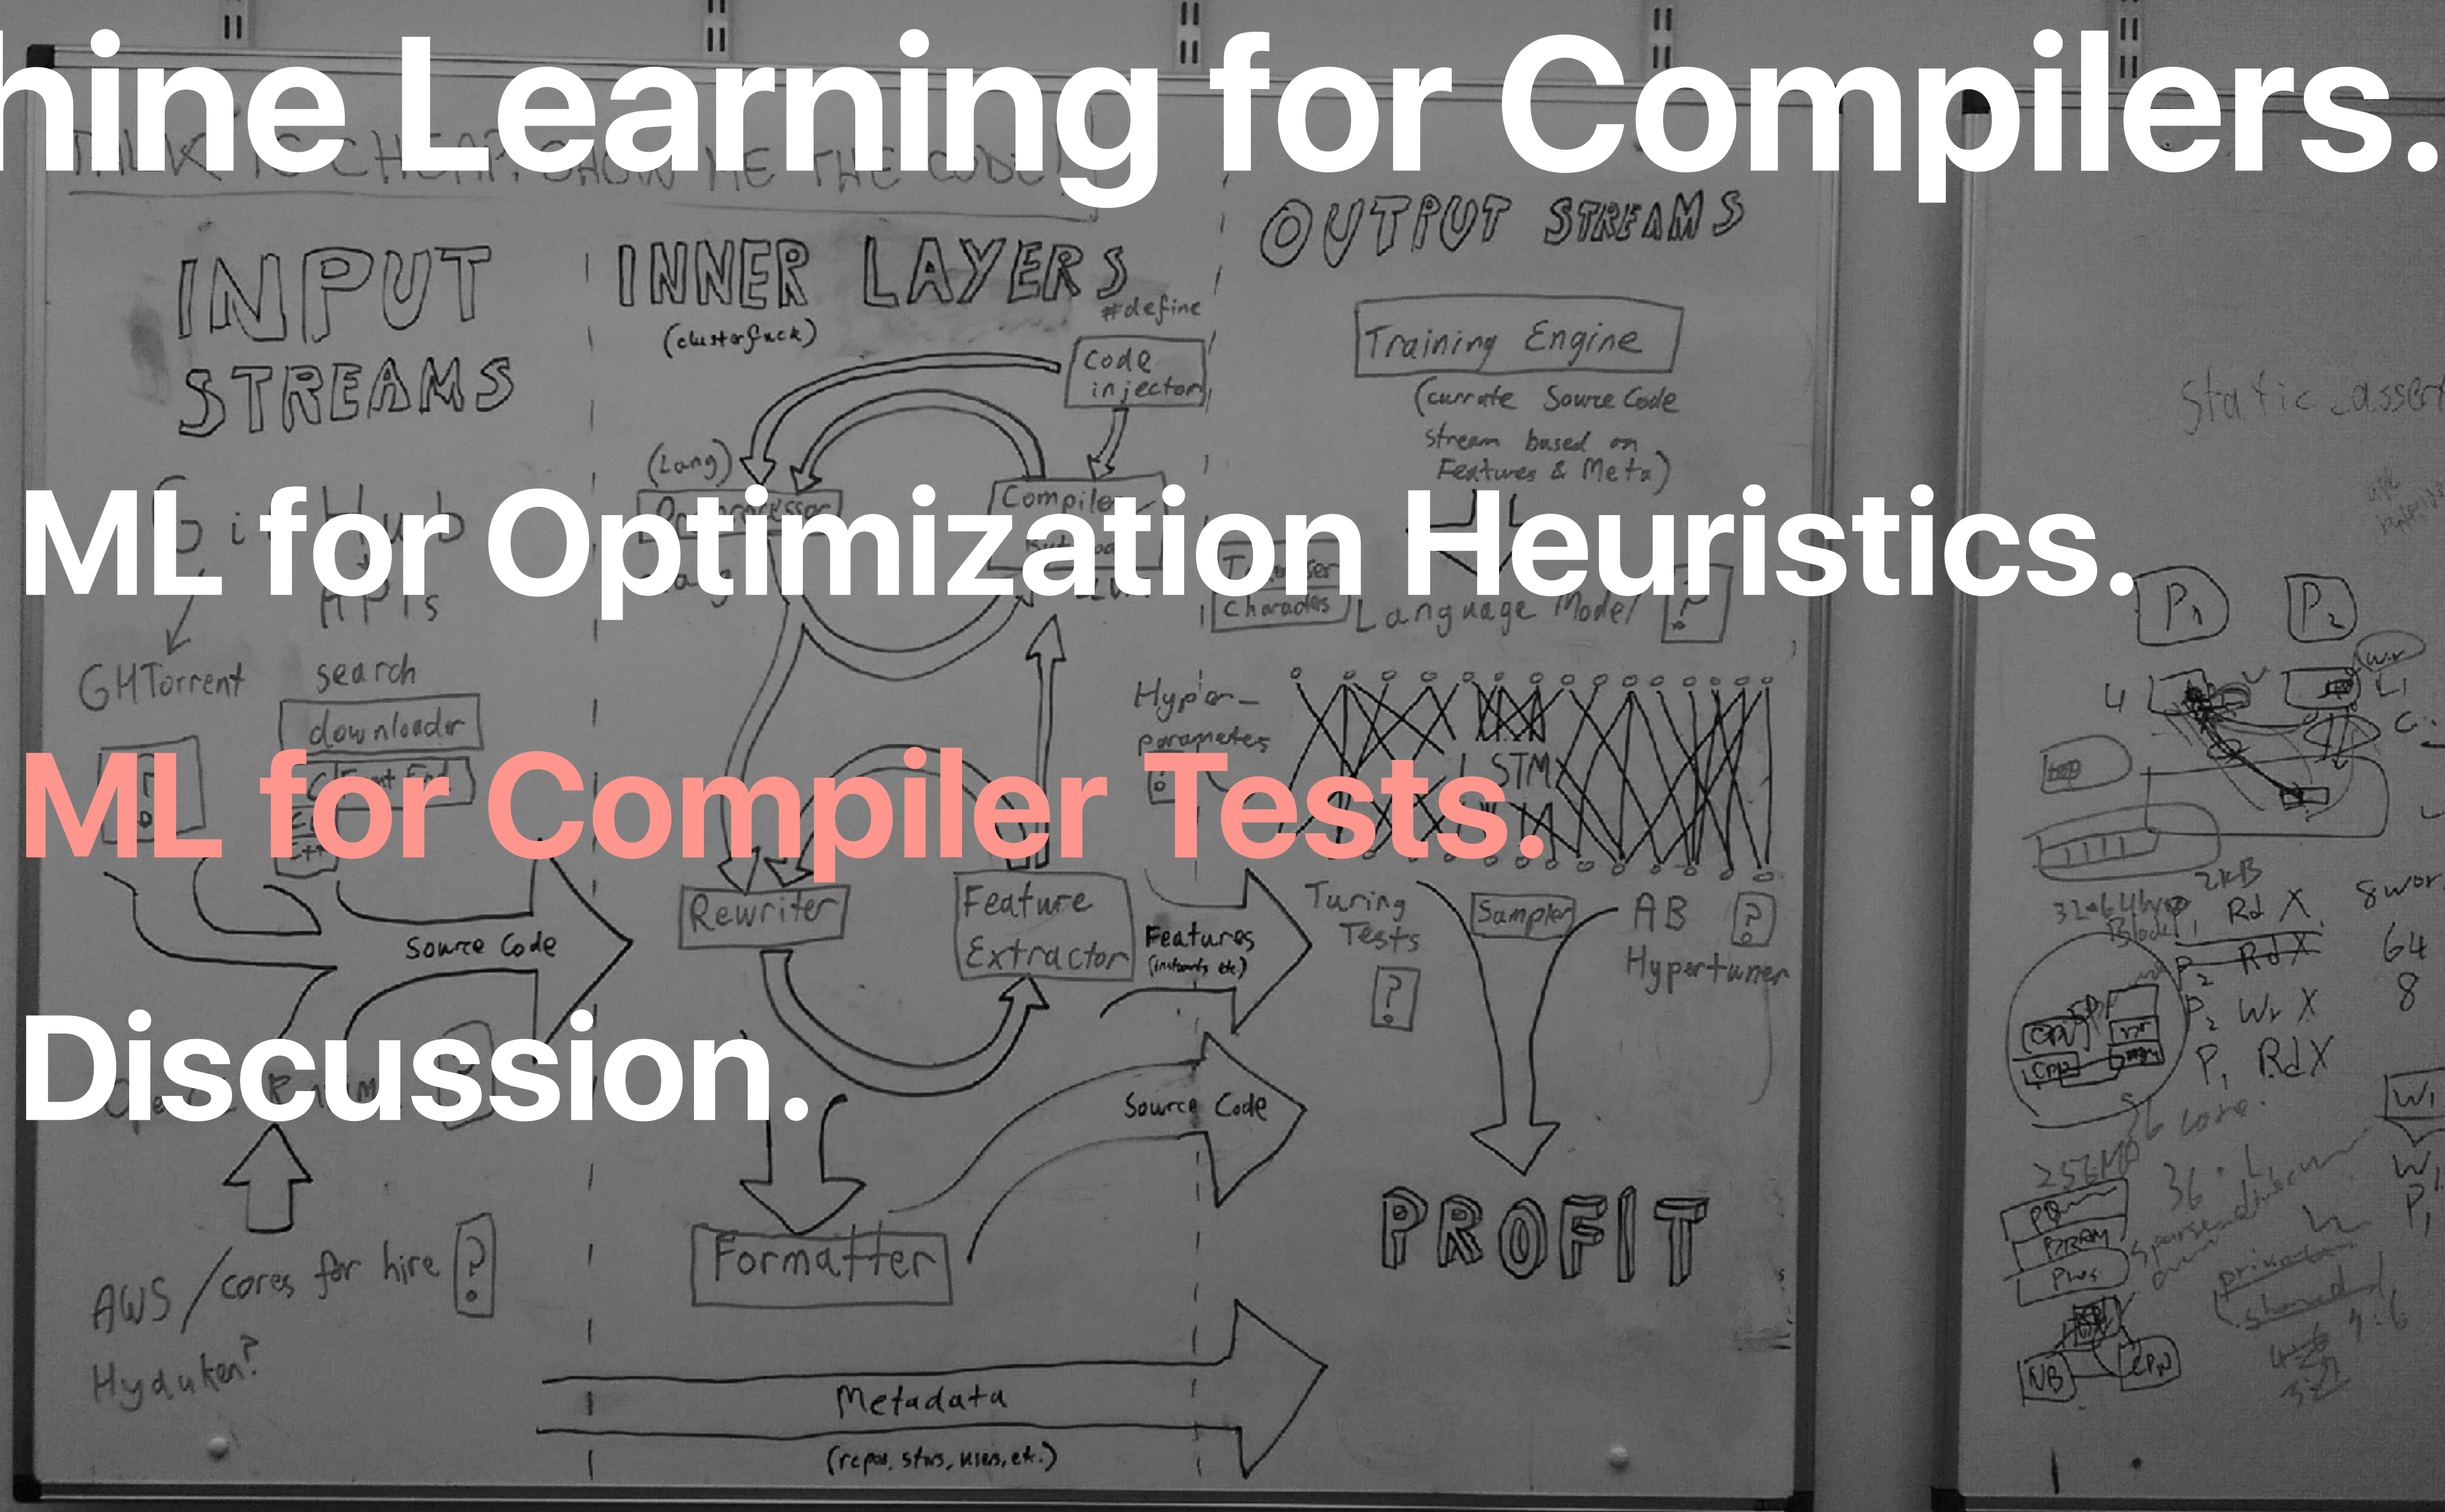


Machine Learning for Compilers.

1. ML for Optimization Heuristics.

2. ML for Compiler Tests.

3. Discussion.

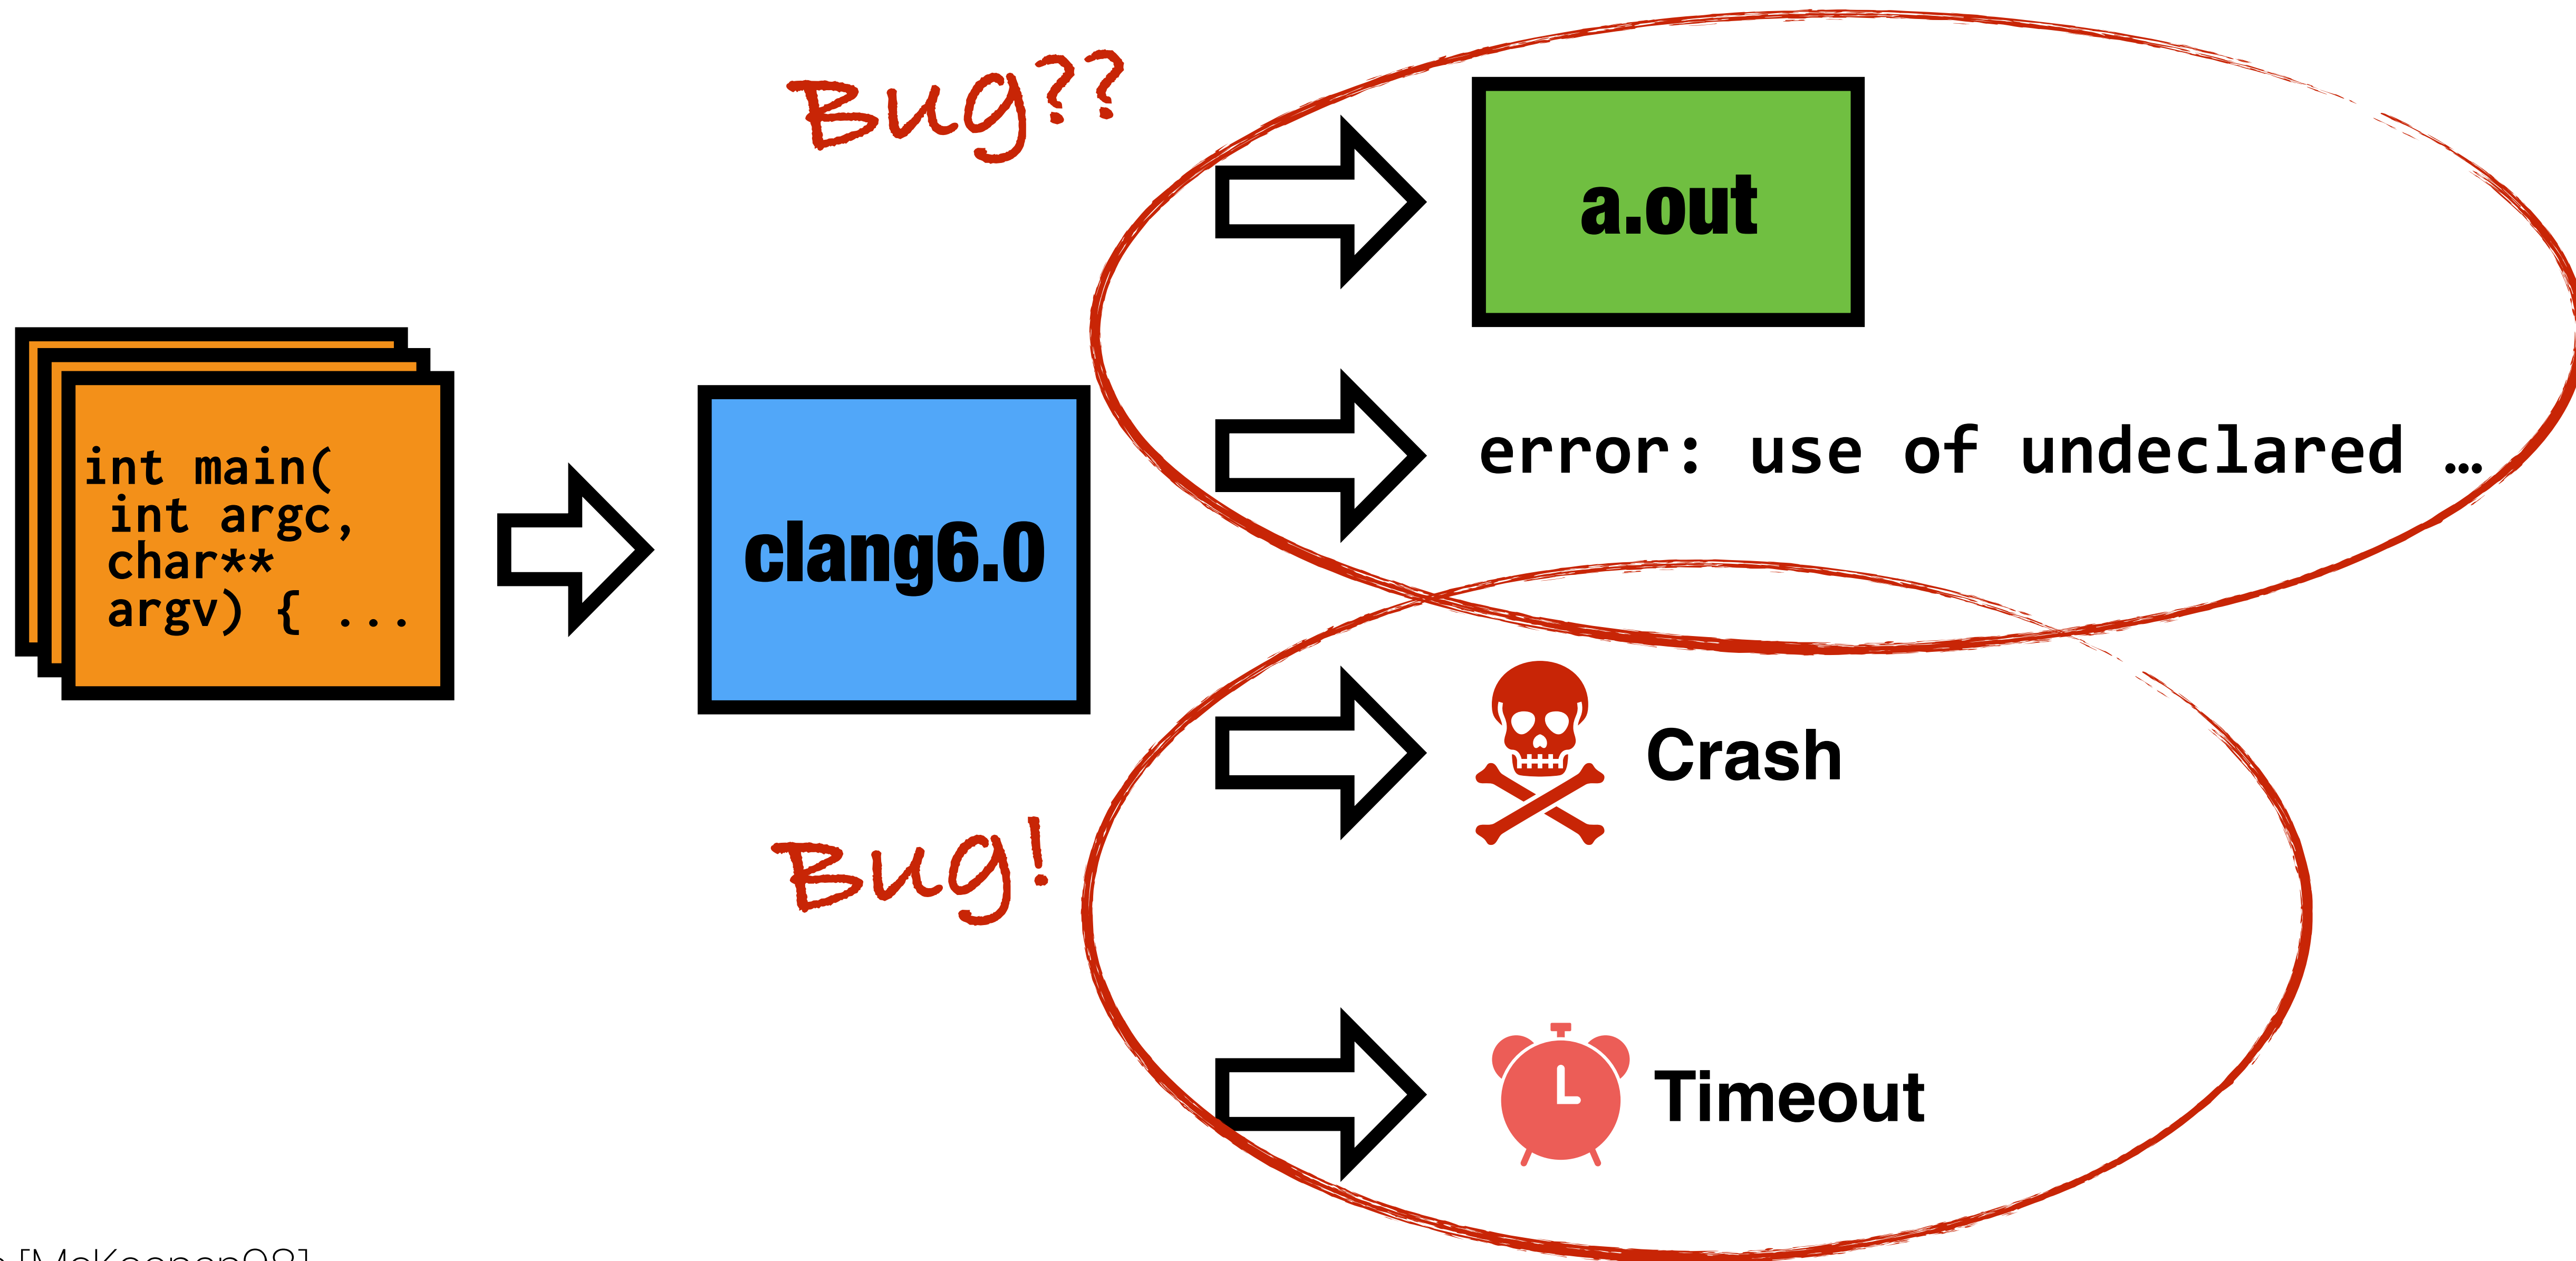


Problem:

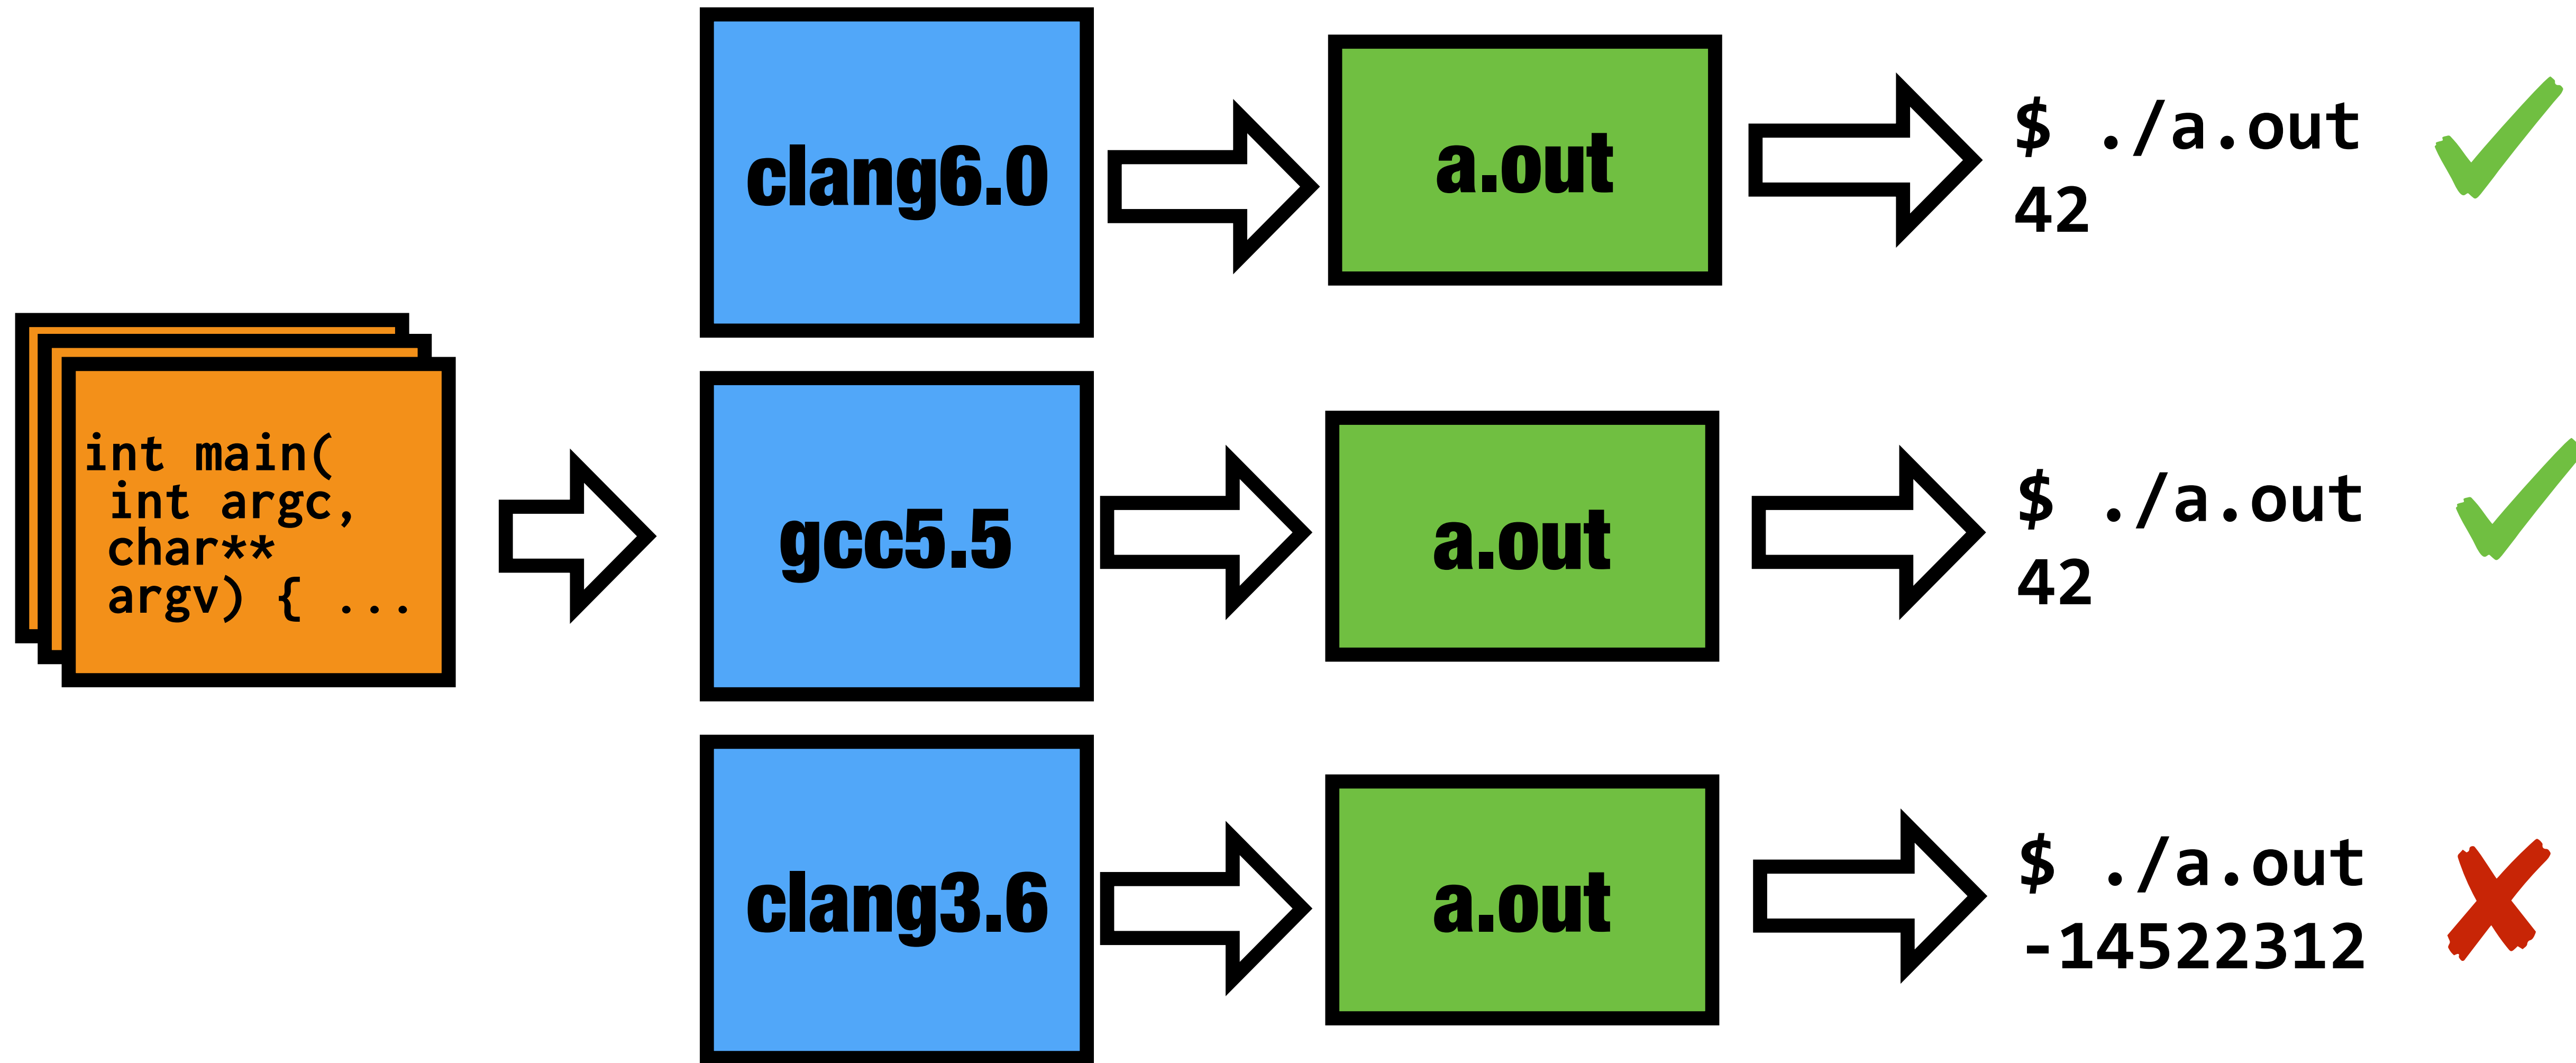
Generating inputs for compiler fuzzing is hard.

(requires complex grammar enumeration and analysis)

fuzzing a compiler

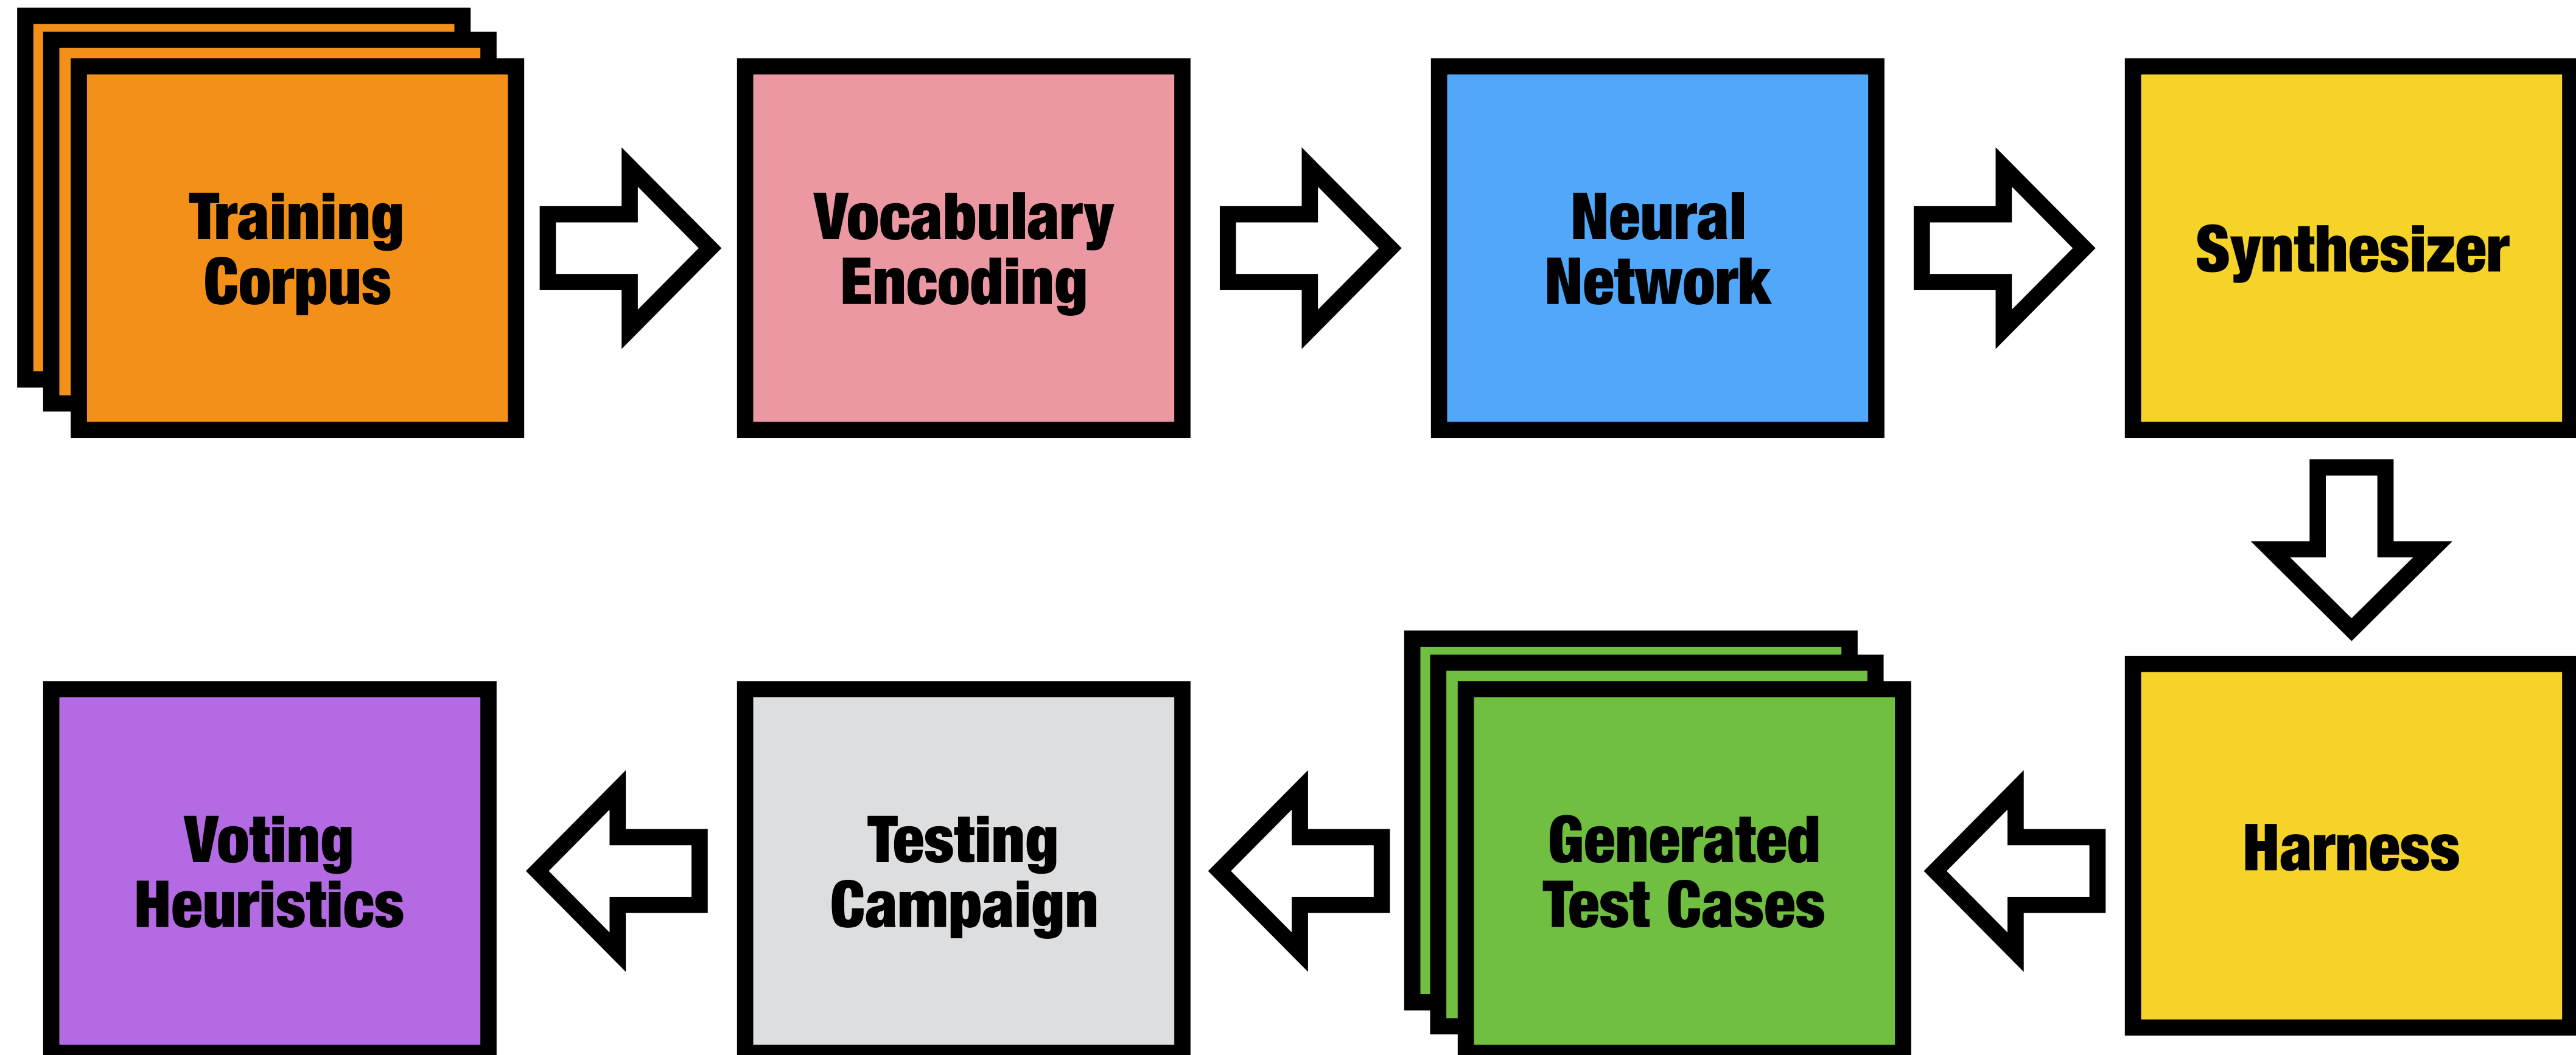


differential testing compilers



Majority rules

ML for compiler test generation



ML for compiler test generation

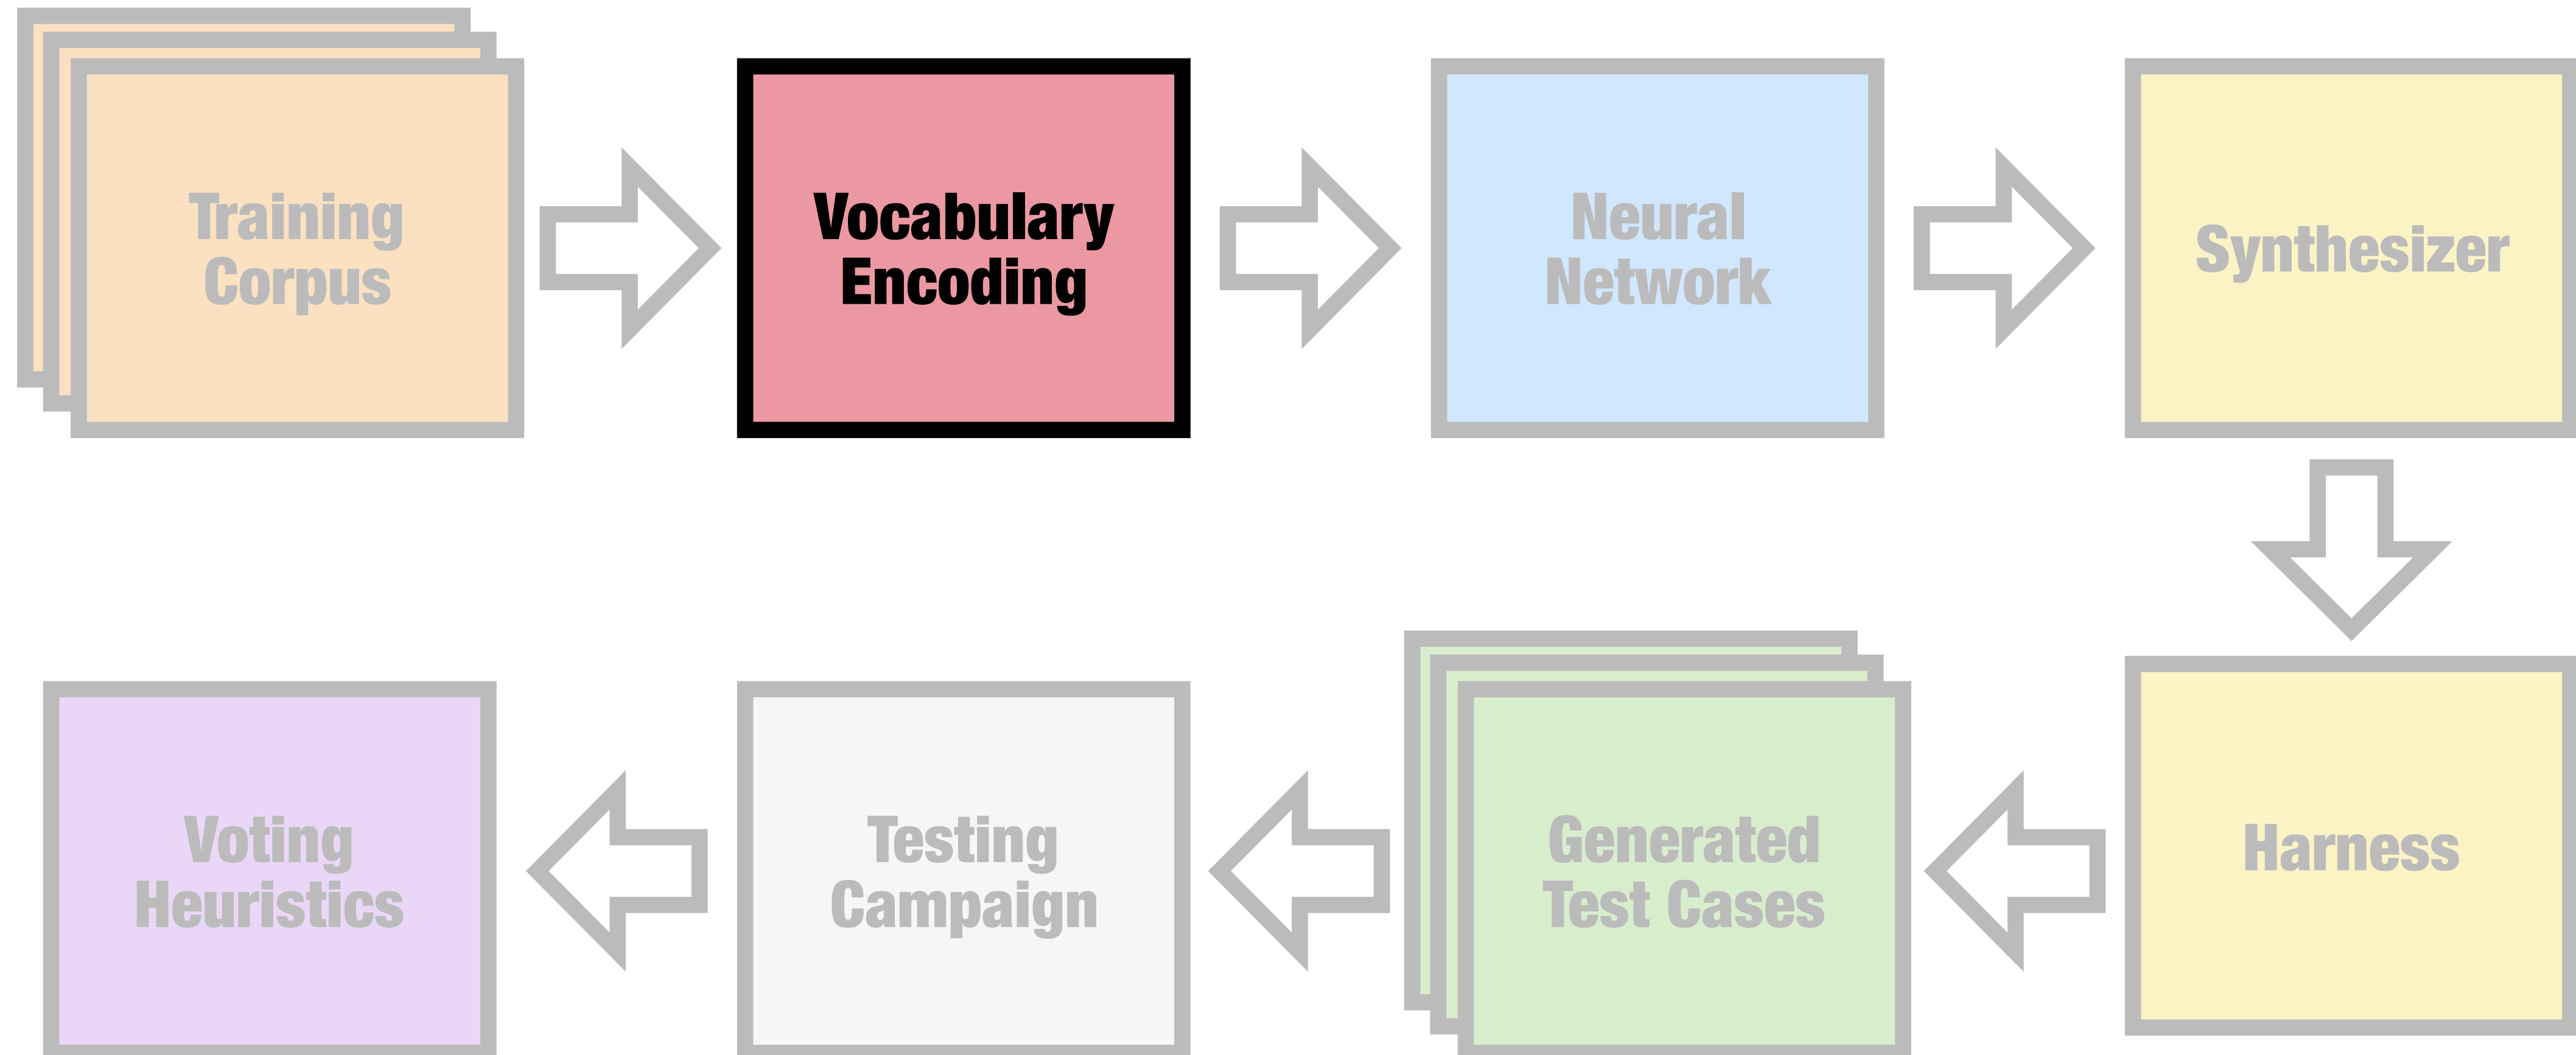
**Training
Corpus**

Mined from  **GitHub**
Filtered by oracle compiler.

**Voting
Heuristics**

- **1k repos**
- **10k files**
- **2.0M LOC**

ML for compiler test generation



vocabulary encoding

```
kernel void A(global float* a, const float b) {  
    a[get_global_id(0)] *= 3.14 + b;  
}
```

Vocab:

Token	Index
kernel	0
[space]	1
void	2
A	3
(4
global	5
float	6
*	7
a	8

Token	Index
,	9
const	10
b	11
)	12
{	13
\n	14
[15
get_global_id	16
0	17

Token	Index
]	18
=	19
3	20
.	21
1	22
4	23
+	24
;	25

Encoded:

0

1

2

1

3

4

5

1

6

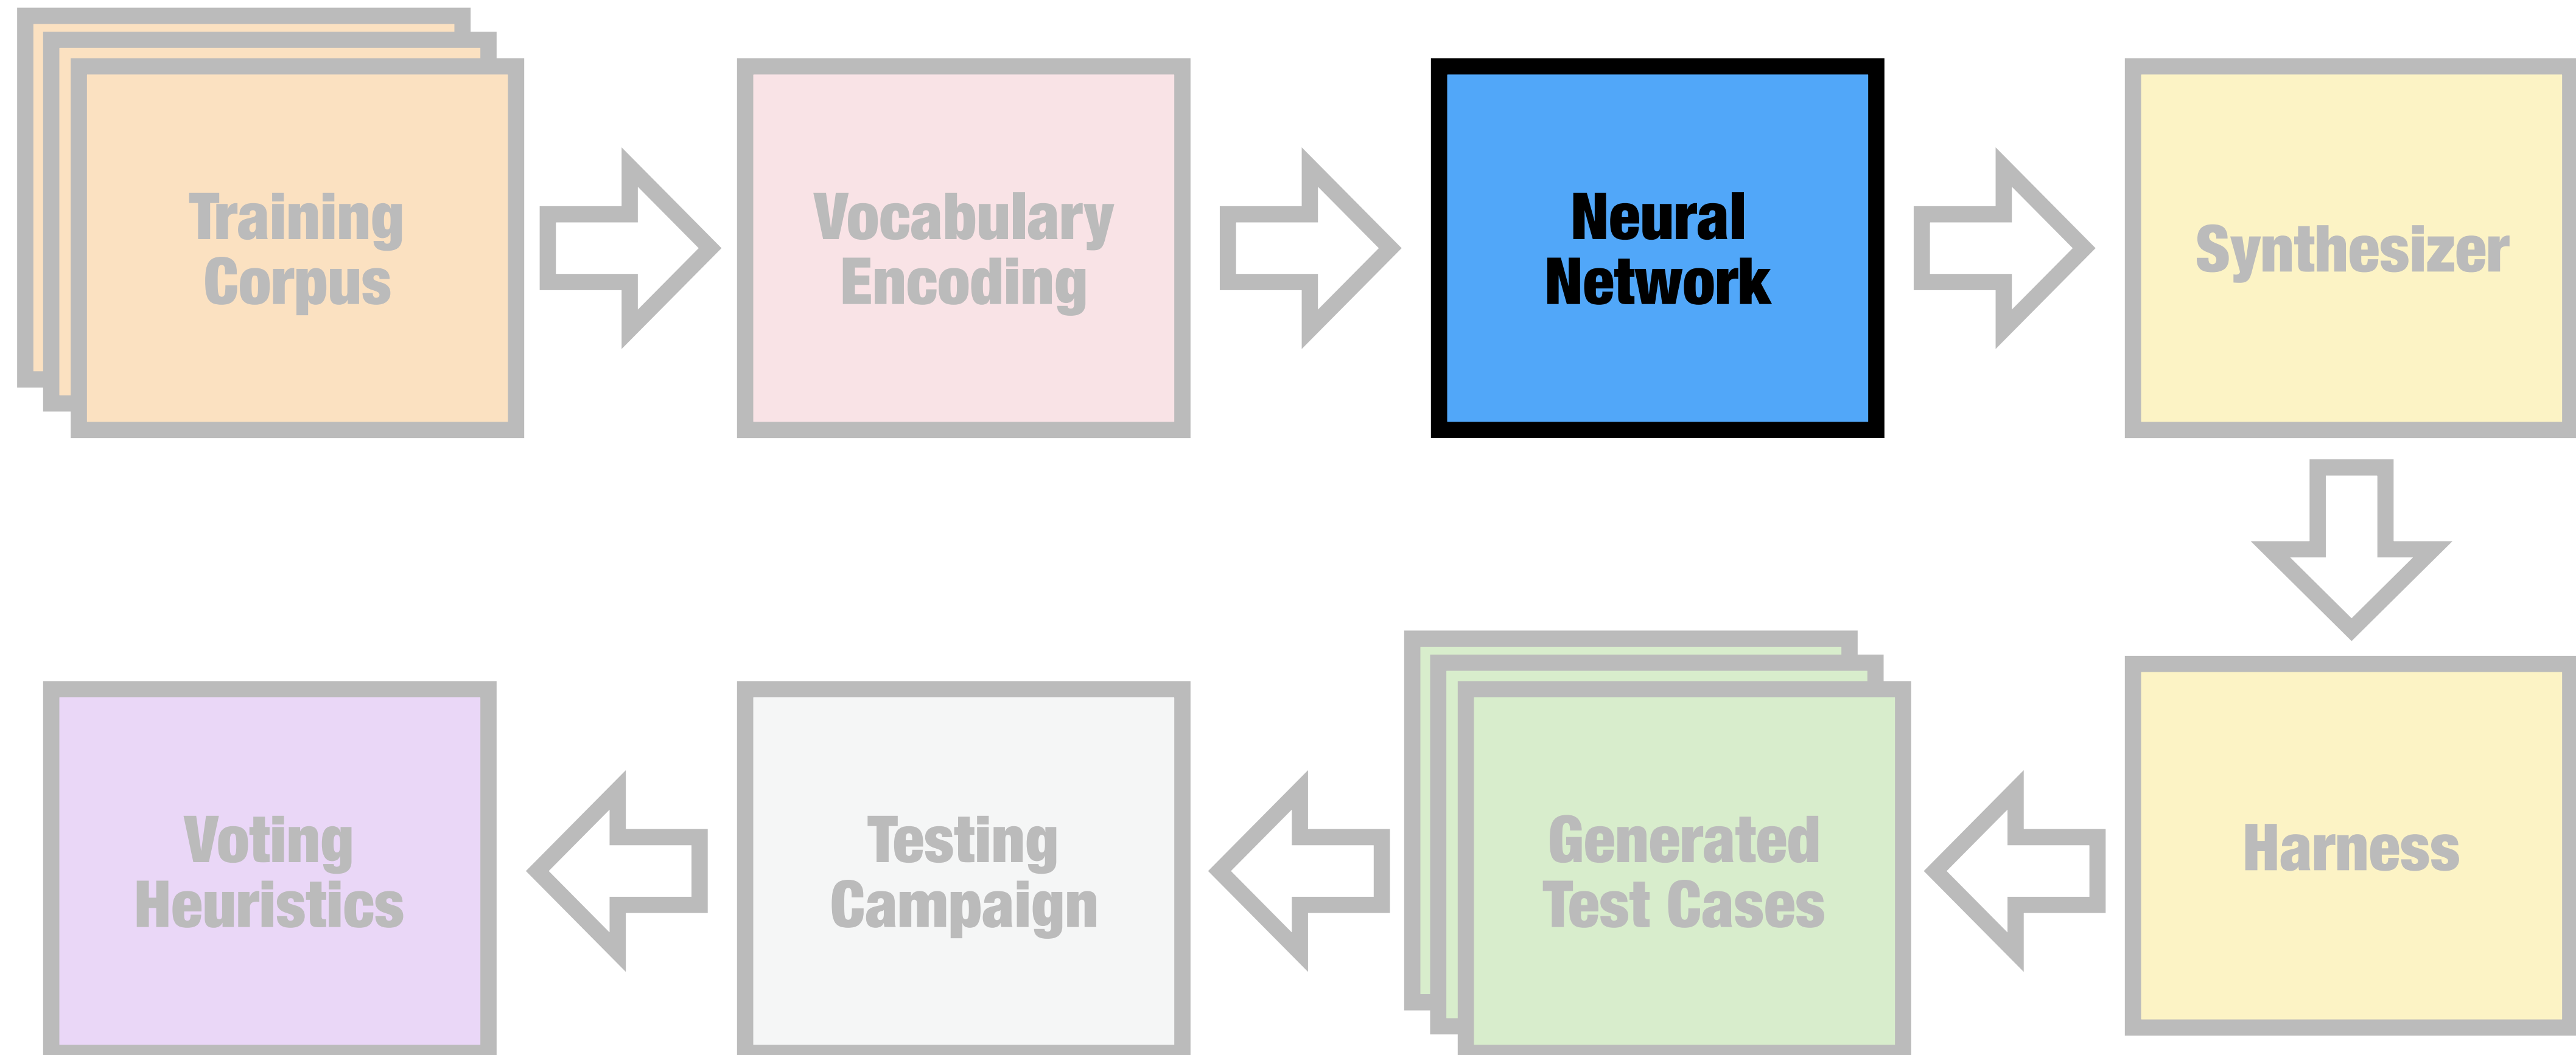
7

1


8

...

ML for compiler test generation



neural network

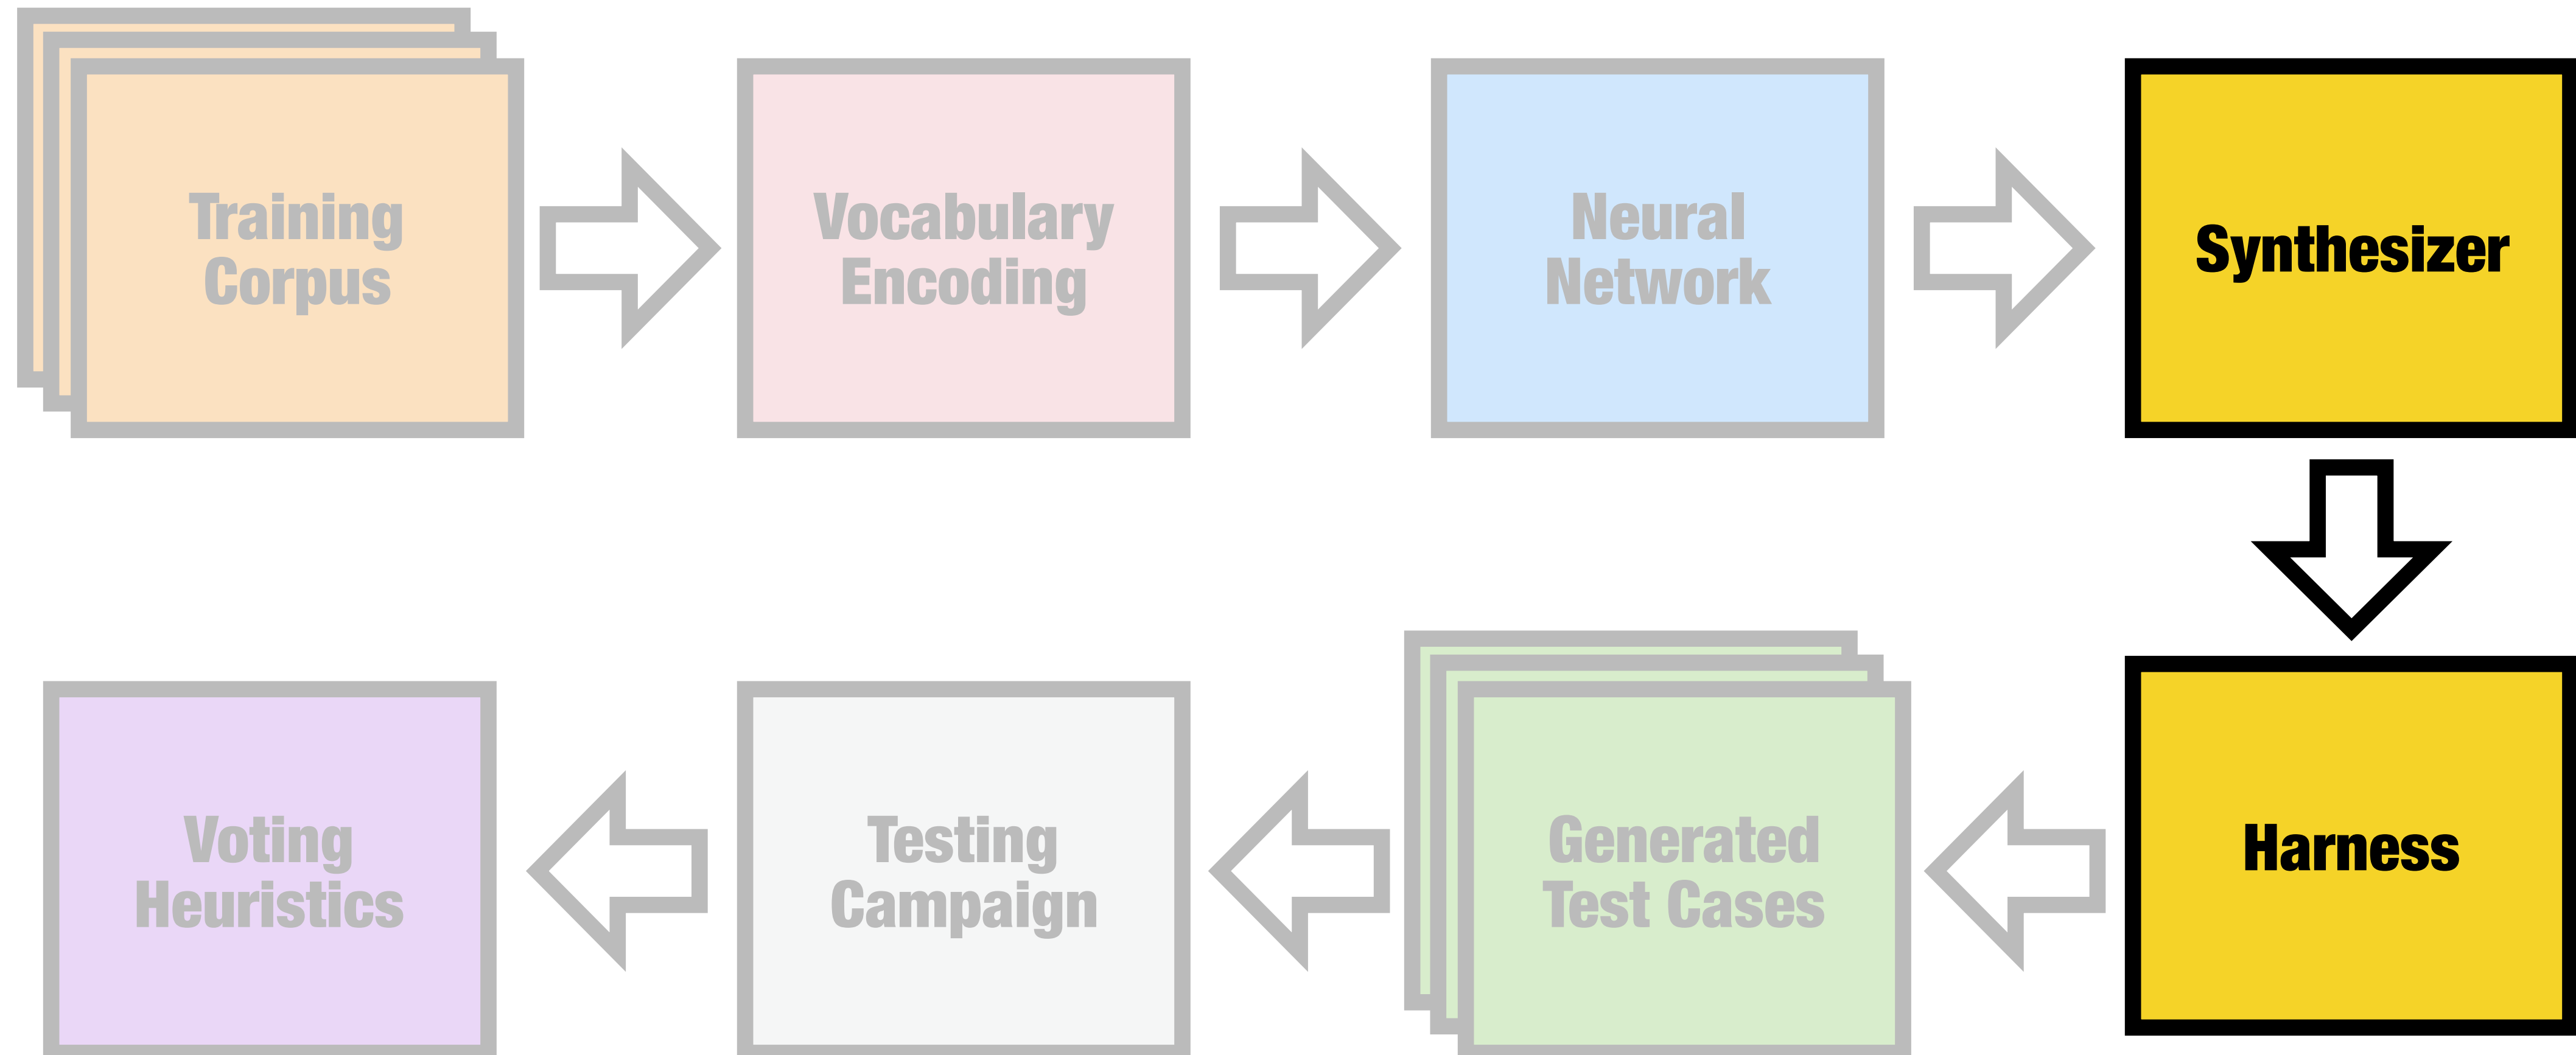


Input: 30M token corpus 0 1 2 ...

Learns probability distribution over corpus.

< 500 lines of code, 12 hours training on GPU.



ML for compiler test generation



synthesizer + harness

- 1. Seed the model with the start of a program.**
- 2. Predict tokens until { } brackets balance.**
- 3. Can we parse signature?**
 - Yes: Generate input data, compile and run it.**
 - No: Compile it but don't run it.**

ML for compiler test generation



Standard majority voting.

**False-positive filtering of
runtime behavior:**

**Combination of off-the-shelf tools
and ad-hoc filters.**

Took ~1 dev-day to develop.

Round 1

Player: 969, Robot: 1031

```
__kernel void A(__global uint *a, __global uint *b, int c, int d) {
    const int e = get_global_id(0);
    const int f = get_global_id(1);
    if (e >= c) return;
    if (f >= d) return;

    int g = e + f * c;
    int h = a[g];

    float3 i[4] = {
        {1.0, 1.0, 1.0}, {1.0, 0.7, 0.3}, {1.5, 0.8, 0.2}, {0.2, 0.8, 0.2}};

    int j = h & 3;
    float3 k = i[j];
    float l = convert_float((h & (255 - 7)));
    int m = l * k.x;
    if (m > 255) m = 255;
    int n = l * k.y;
    if (n > 255) n = 255;
    int o = l * k.z;
    if (o > 255) o = 255;
    b[g] = o + n * 256 + m * 65536;
}
```

```
__kernel void A(__global int* a, __global int* b, __global int* c, const int
d) {
    unsigned int e = get_global_id(0);
    unsigned int f = get_local_size(0);


    for (unsigned int g = e; g < e + 1; g++) {
        for (unsigned int h = 0; h < f; ++h) {
            a[g * d + h] = 0;
        }
    }
}
```

52%

⚙️ This side is the robot


⚙️ This side is the Robot

testing campaign




10 OpenCL compilers

3 GPUs, 5 CPUs, Xeon Phi, Emulator



Test with optimizations on / off


Treat as separate testbeds



48 hours per testbed

results overview

Errors in ever compiler!



ML for compiler tests.

Advantages:

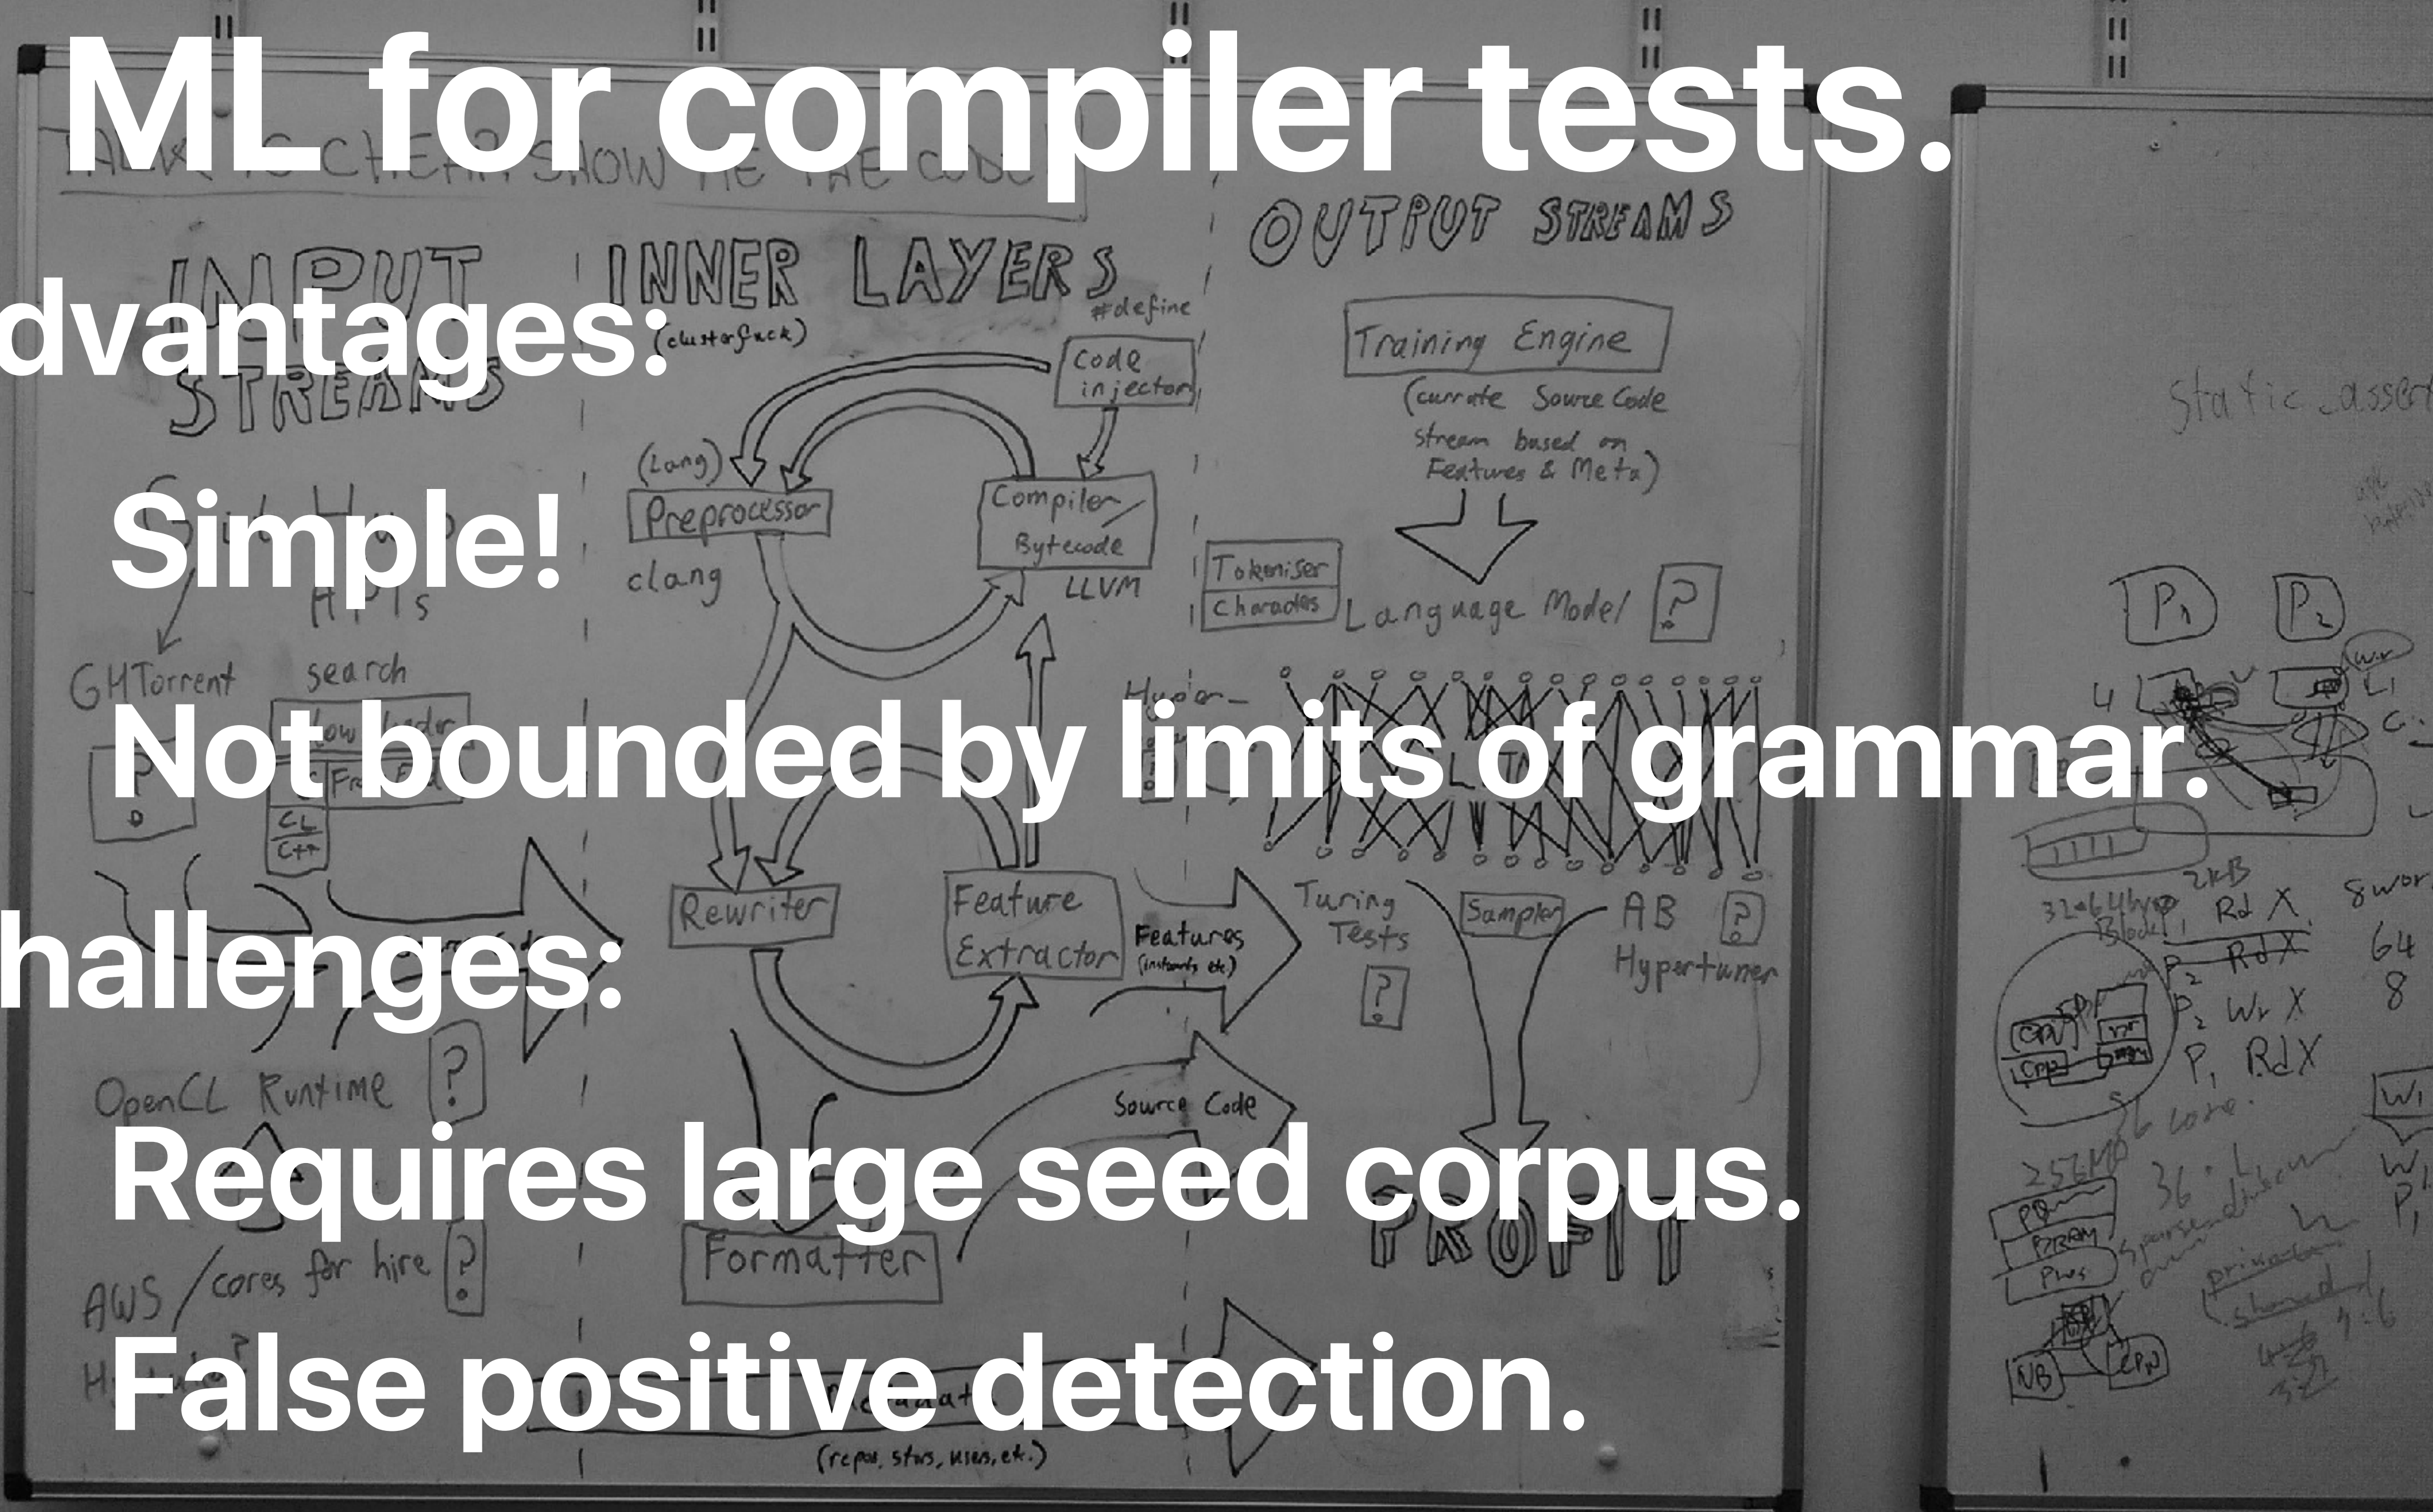
Simple!

Not bounded by limits of grammar.

Challenges:

Requires large seed corpus.

False positive detection.

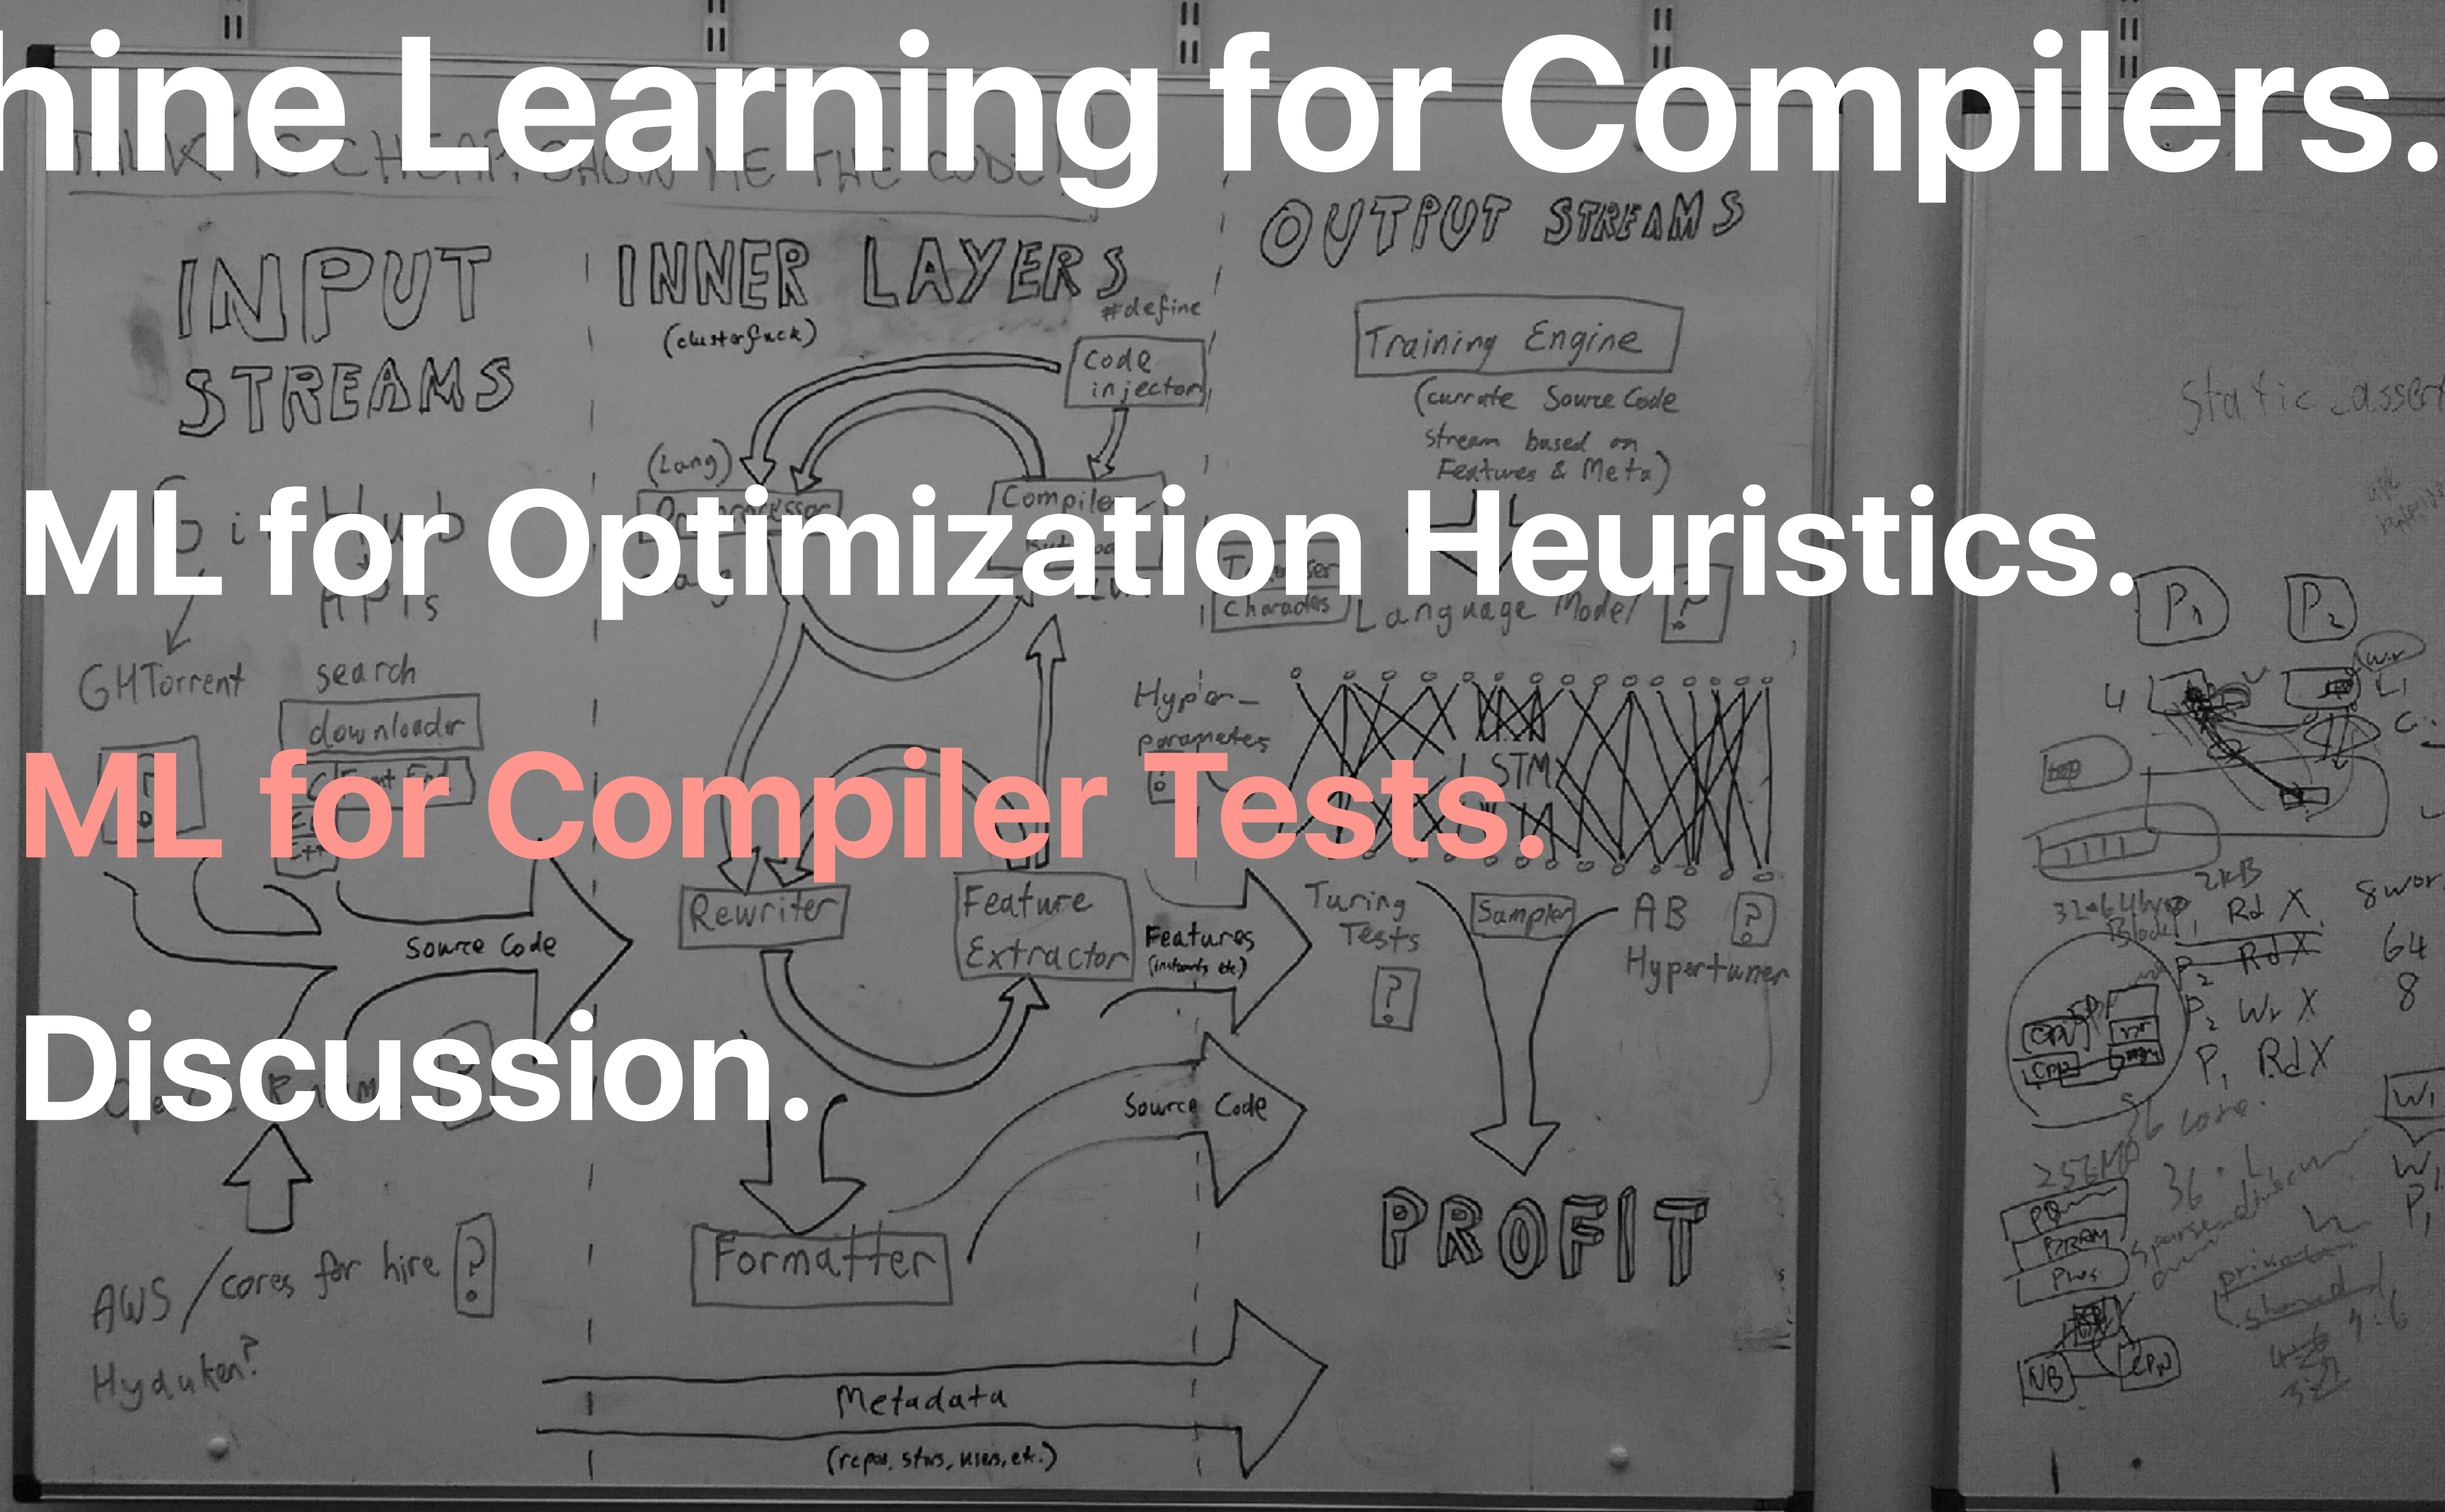


Machine Learning for Compilers.

1. ML for Optimization Heuristics.

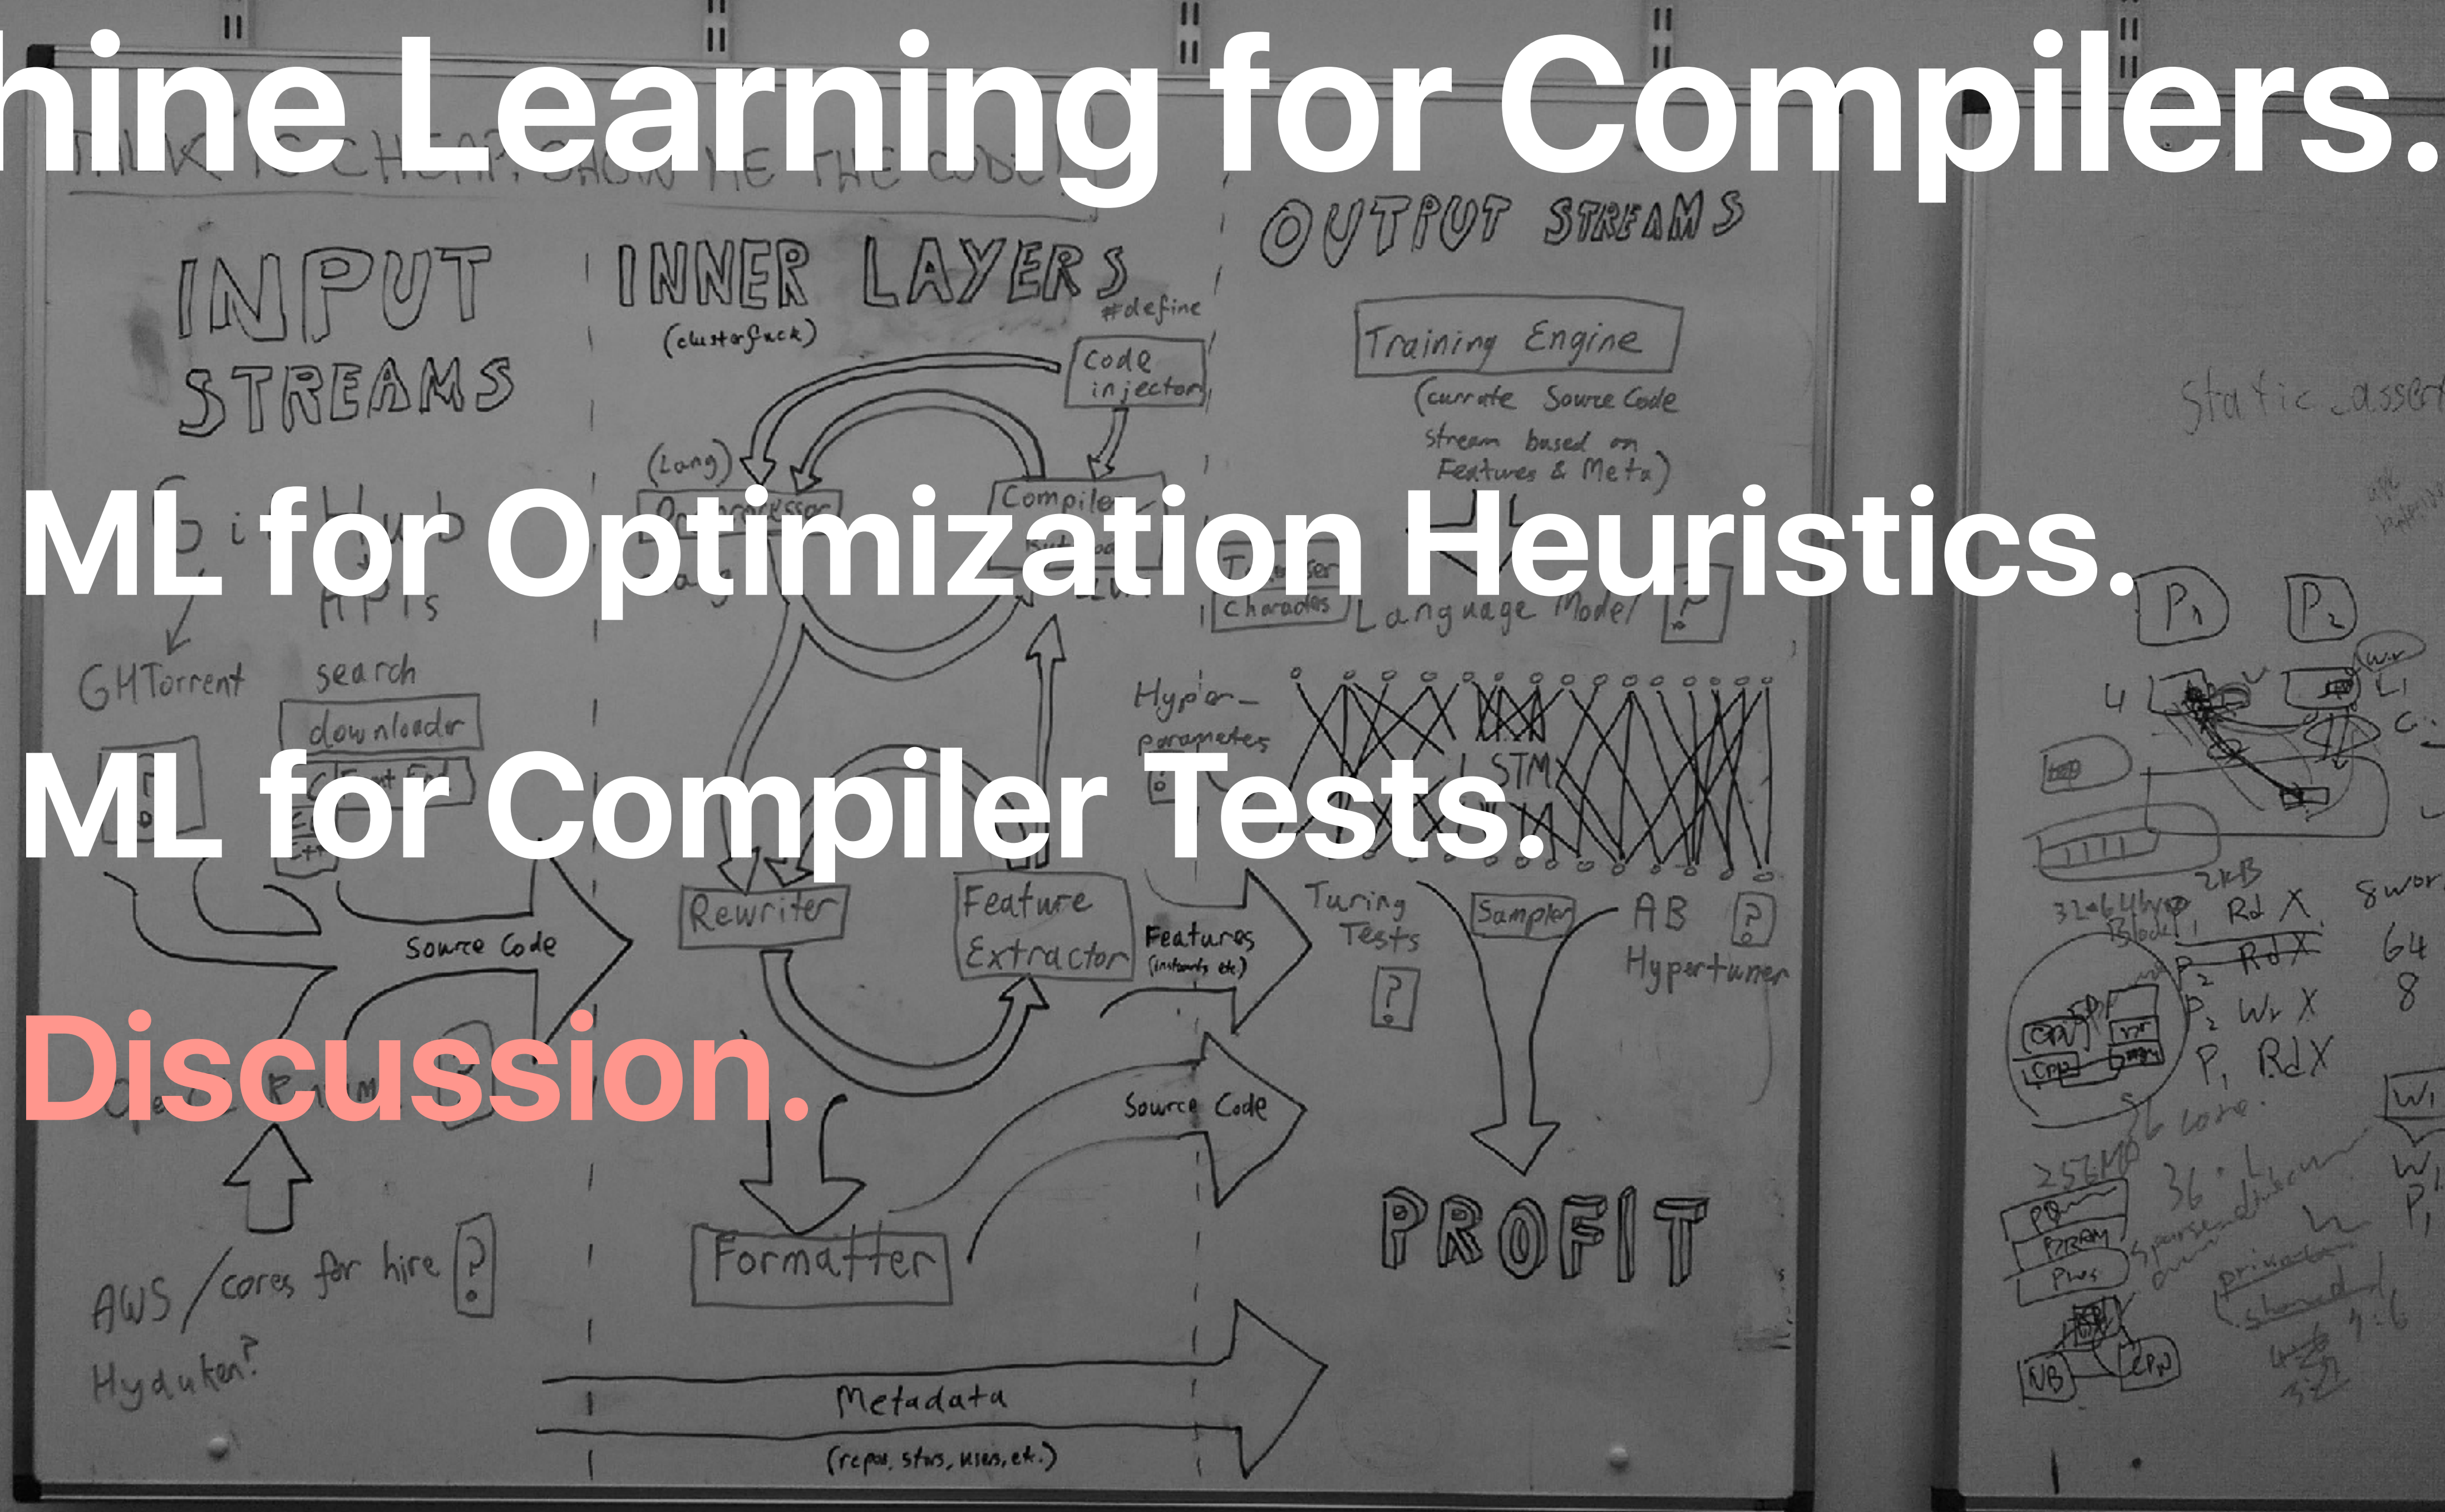
2. ML for Compiler Tests.

3. Discussion.



Machine Learning for Compilers.

1. ML for Optimization Heuristics.
2. ML for Compiler Tests.
3. Discussion.



Discussion

What problems do you know of where:

There is room for improvement?

Currently rely on domain expertise?

Challenges with ML:

Feature design / selection.

Correctness and data distribution.